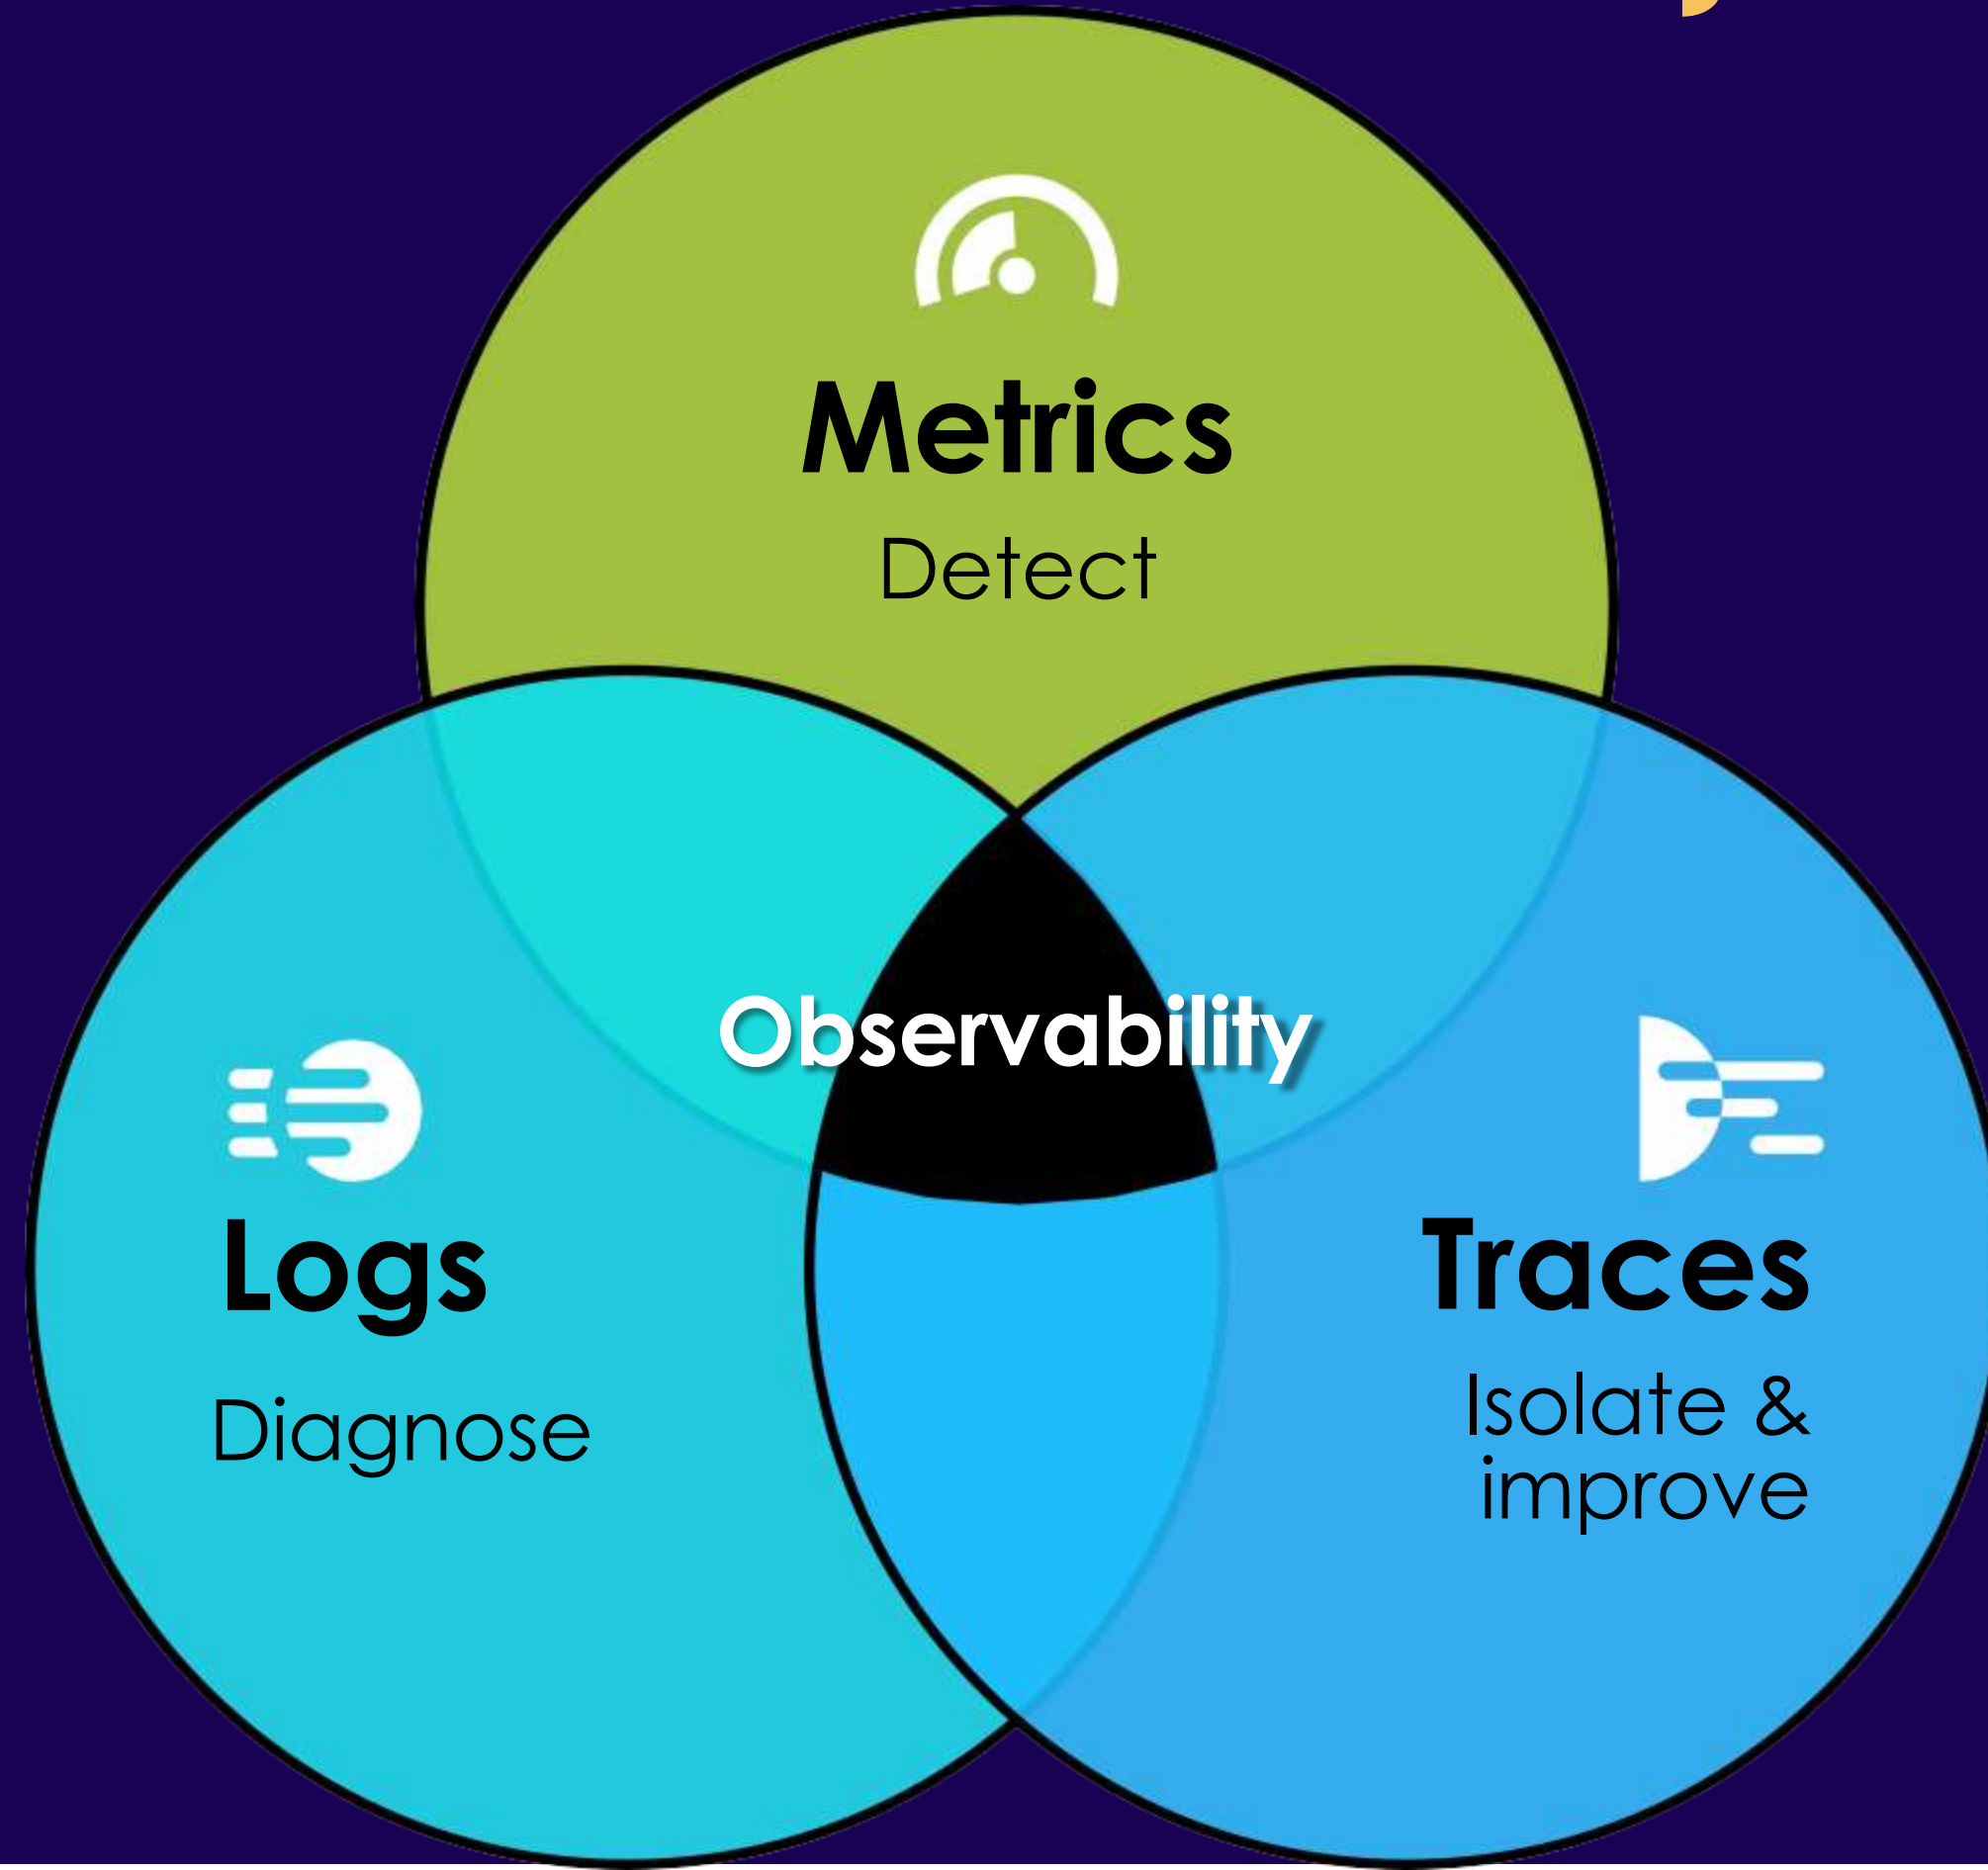


Getting started with Prometheus & Grafana

Seshagiri Sriram - Aug 26 2025

The vision: **unified observability**



Observability

- Being able to measure “things” or witness state changes.
- Not useful if doing so alters behavior (significantly).
- **Measurement:** a single measurement of something
- **Metric:** something that you are measuring The version of deployed code
 - Total cost on Amazon services total bugs filed, bug backlog Total queries executed

Observability

- **Measurement Velocity** - The rate @ which measurement is taken
- **Perspective**
- **Visualization**
- **Trends**
- **Alerting**
- **MONITORING IS ALL OF THIS 😊**

A new Perspective

Monitoring		Observability
1	Says whether the System is Working or Not	Why its not working
2	Collects Metrics and Logs from a System	Actionable Insights gained from the Metrics
3	Failure Centric	Overall Behavior of the System
4	Is “the How” of something you do	Is “The Process” of something you have
5	I monitor you	You make yourself observable

Pillars of Observability

Logs/events



Immutable records of discrete events that happen over time

Metrics



Numbers describing a particular process or activity measured over intervals of time

Traces



Data that shows, for each invocation of each downstream service, which instance was called, which method within that instance was invoked, how the request performed, and what the results were

Prometheus



- A monitoring & alerting system, Inspired by Google's BorgMon
- Originally built by SoundCloud in 2012
- Open Source, now part of the CNCF
- Simple text-based metrics format
- Multidimensional data model
- Rich, concise query language

Prometheus



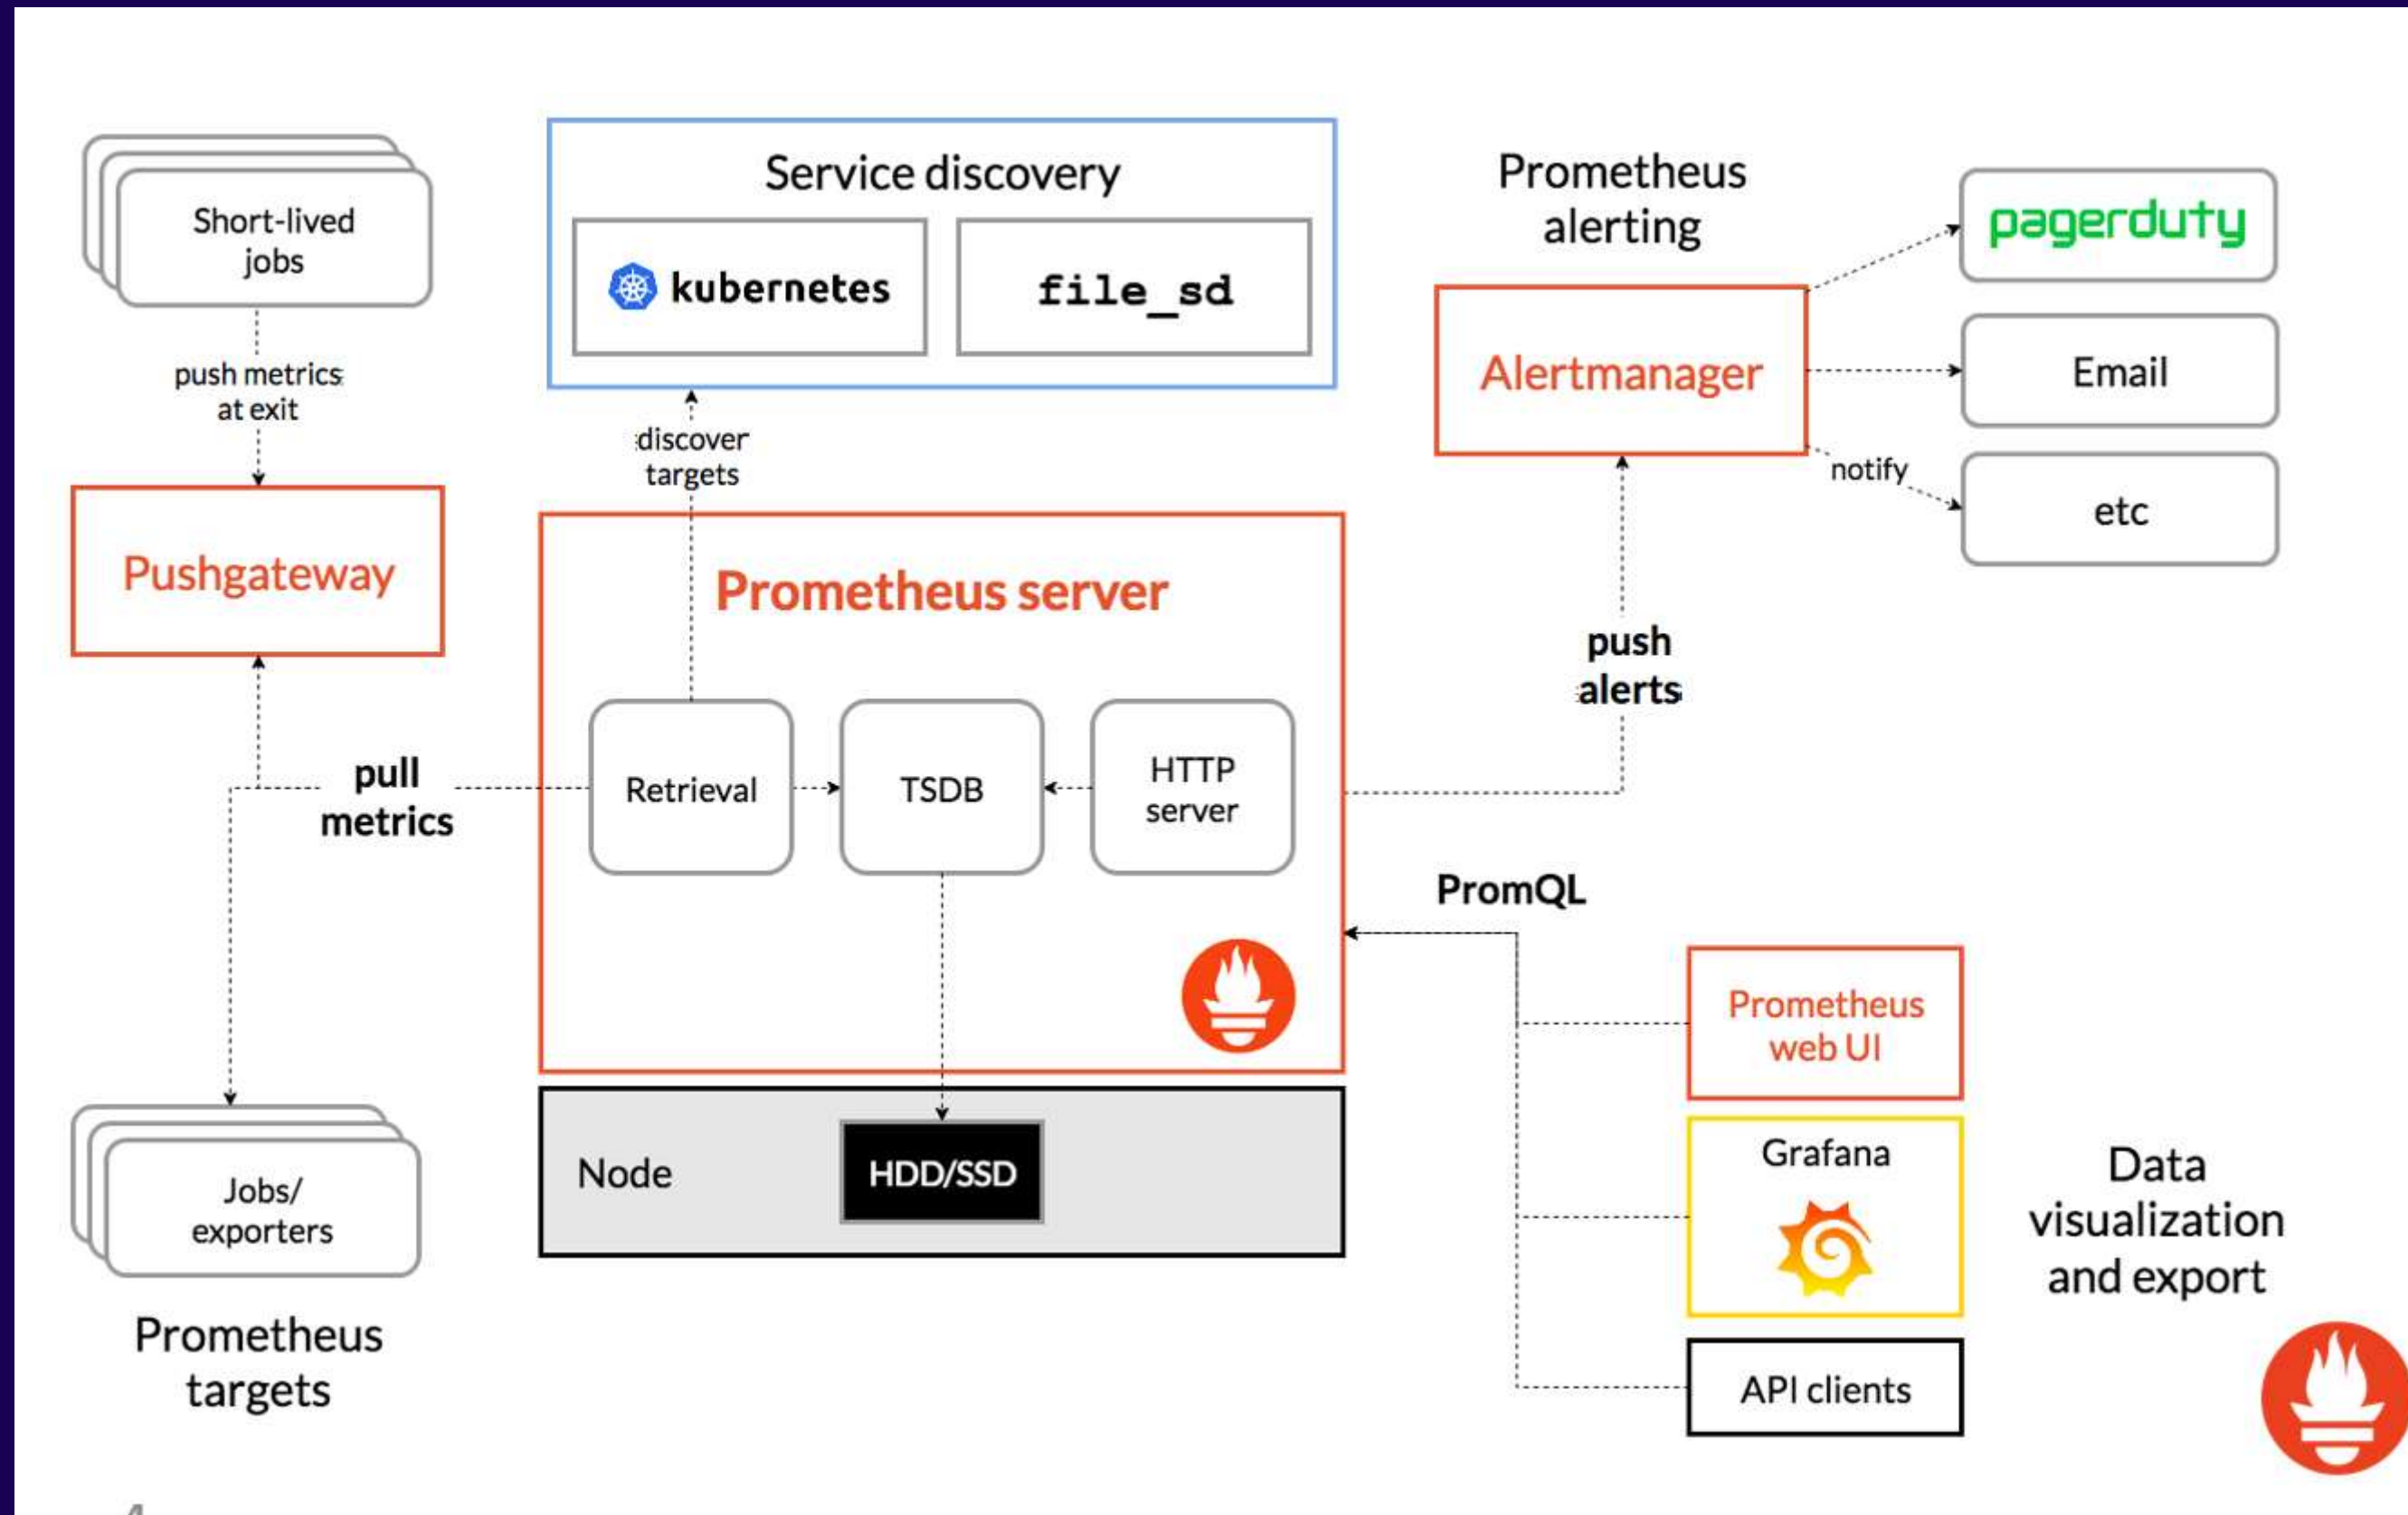
- A data scraper that pulls metrics data over HTTP periodically at a configured interval.
- A time-series database to store all the metrics data.
- A simple user interface where you can visualize, query, and monitor all the metrics.
- Written in Go, fully published in 2015.

Prometheus

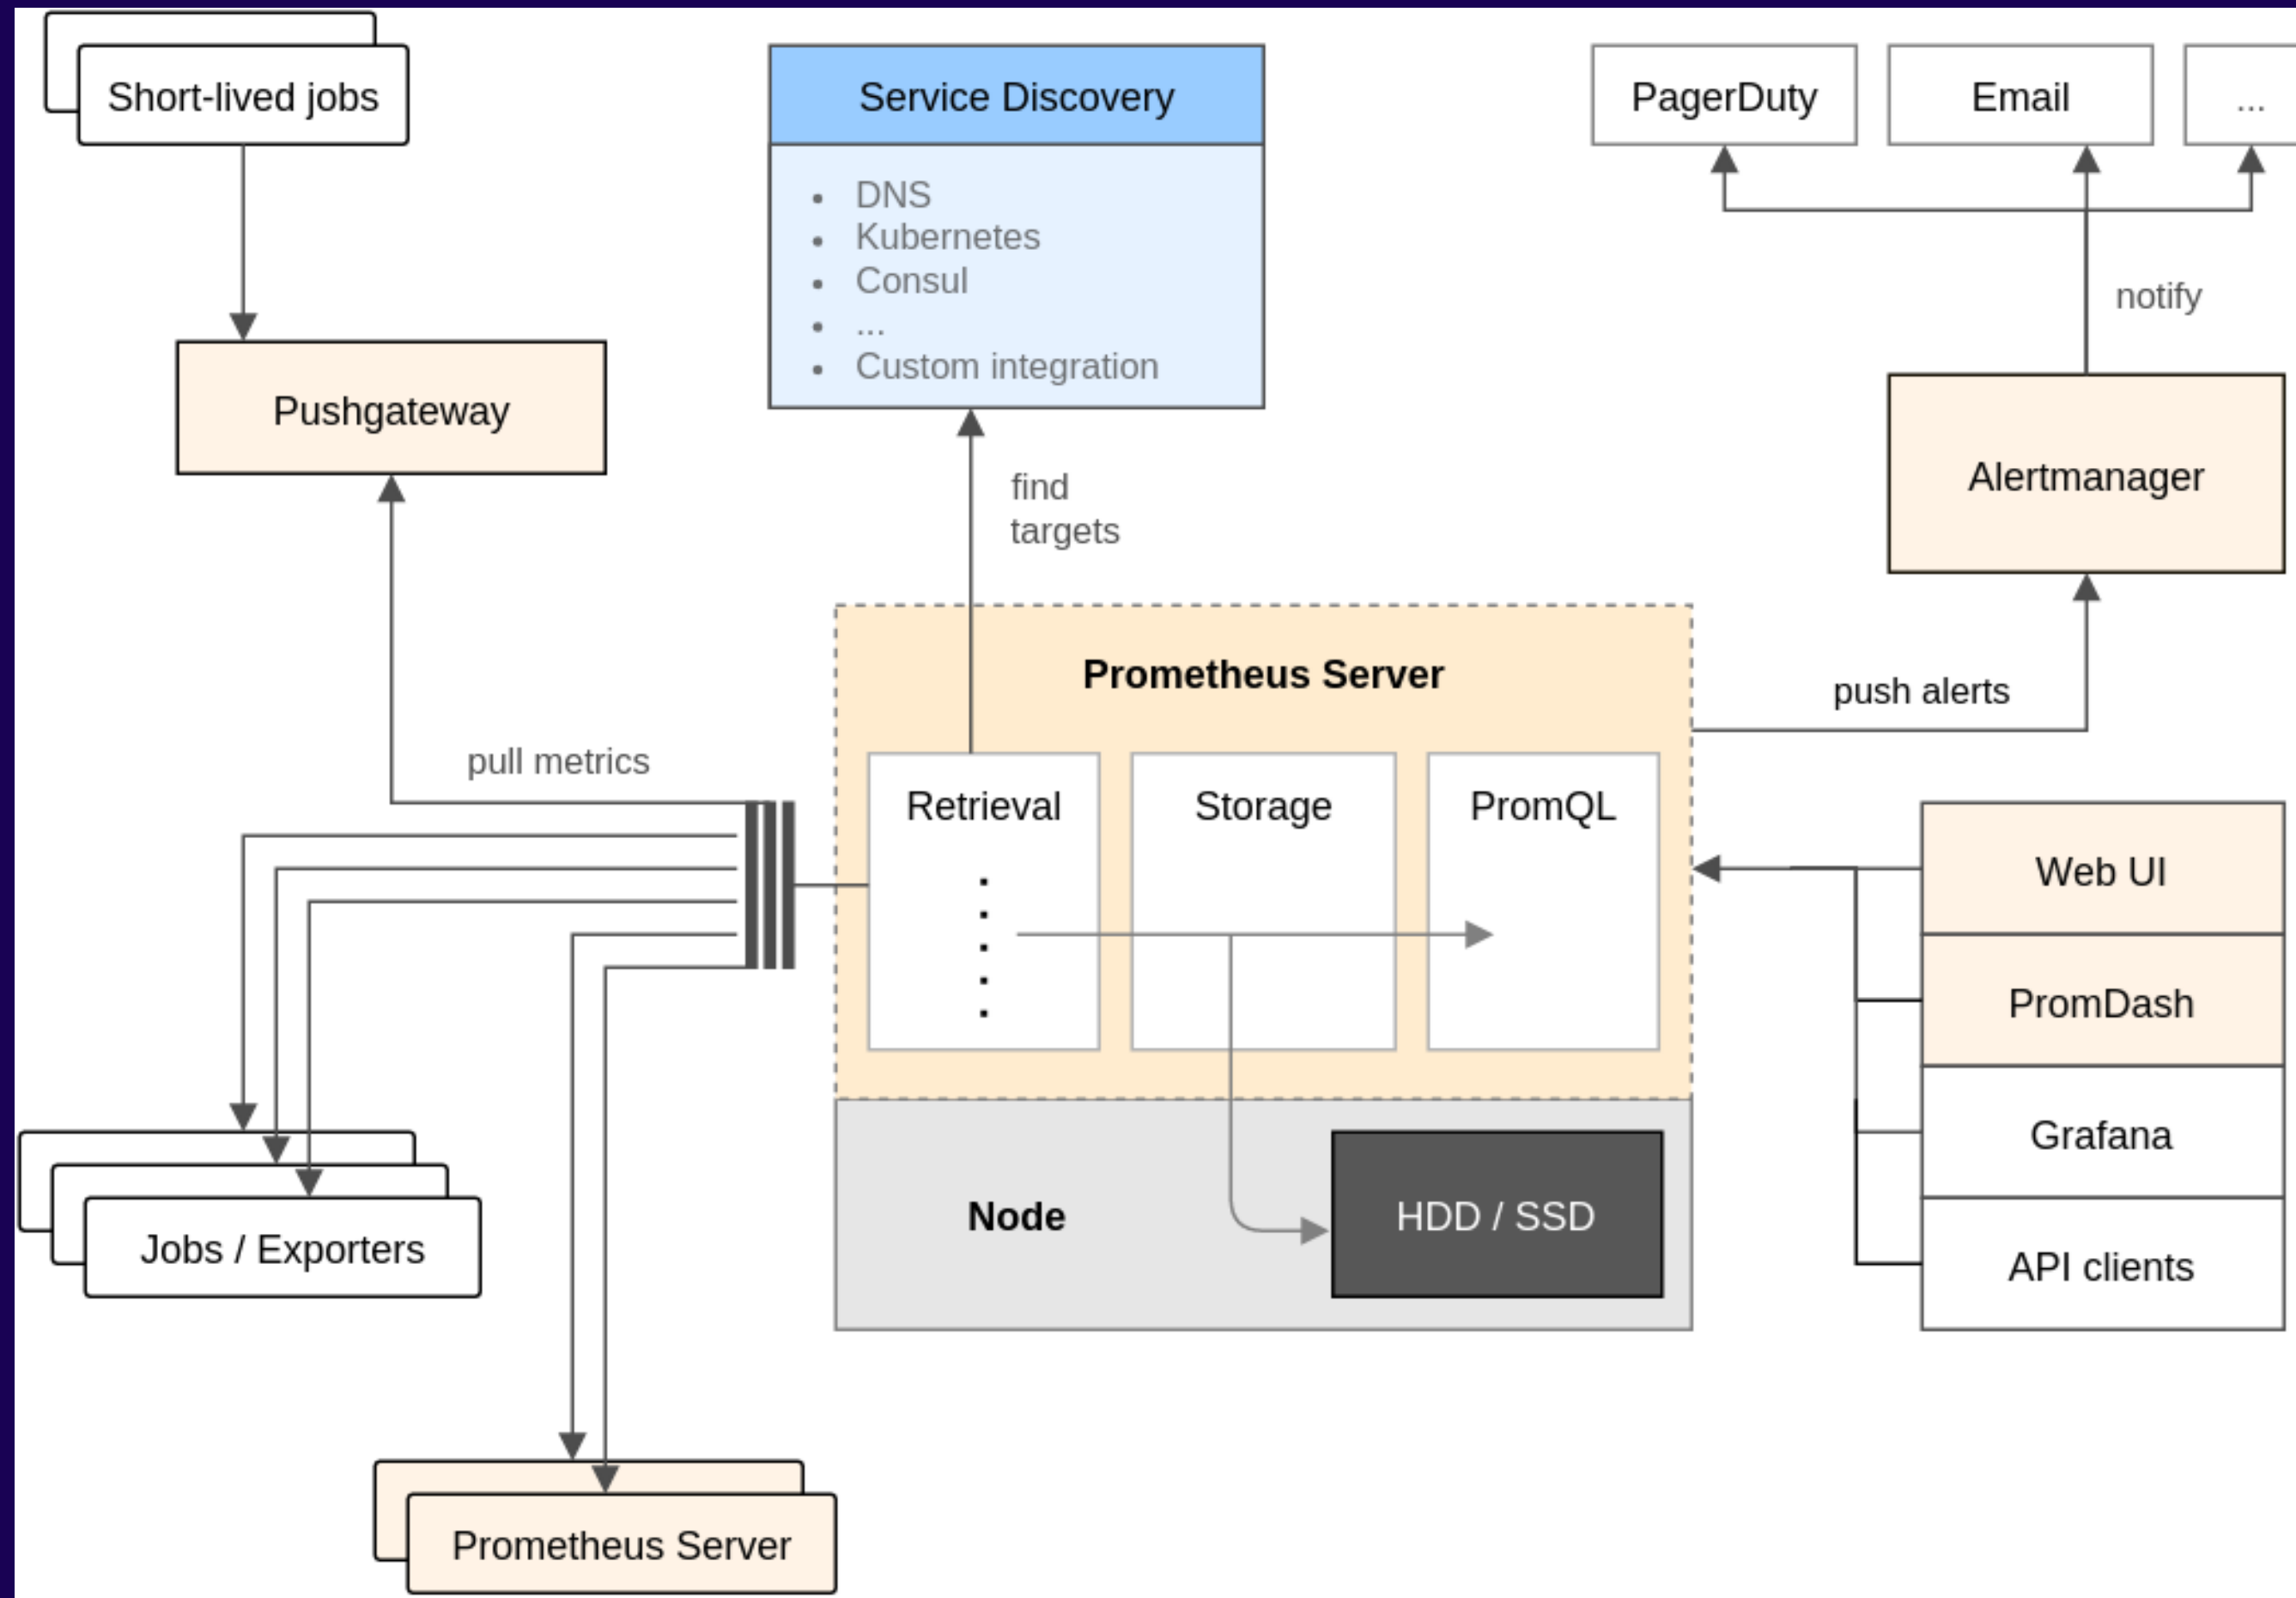


- Monitoring systems and TSDB
 - Instrumentation
 - Metrics collection and Storage
 - Querying
 - Alerting
 - Dashboarding / Graphing / Trending
- Focus on
 - Dynamic Cloud Environments
 - Operational Systems Monitoring

Prometheus



Prometheus



Prometheus



What I can do

Dimensional Data Model

Powerful Query Language

Efficiency

Operational Simplicity

What it cannot do

Raw Log/event Collection

Request Tracing

Anomaly Detection

Automatic horizontal scaling

User Management and authentication has to be handled separately

Expression Browser


Prometheus Alerts Graph Status Help

sort_desc(sum(bazooka_instance_memory_limit_bytes - bazooka_instance_memory_usage_bytes) by (app, proc)) / 1024 / 1024 / 1024 Execute

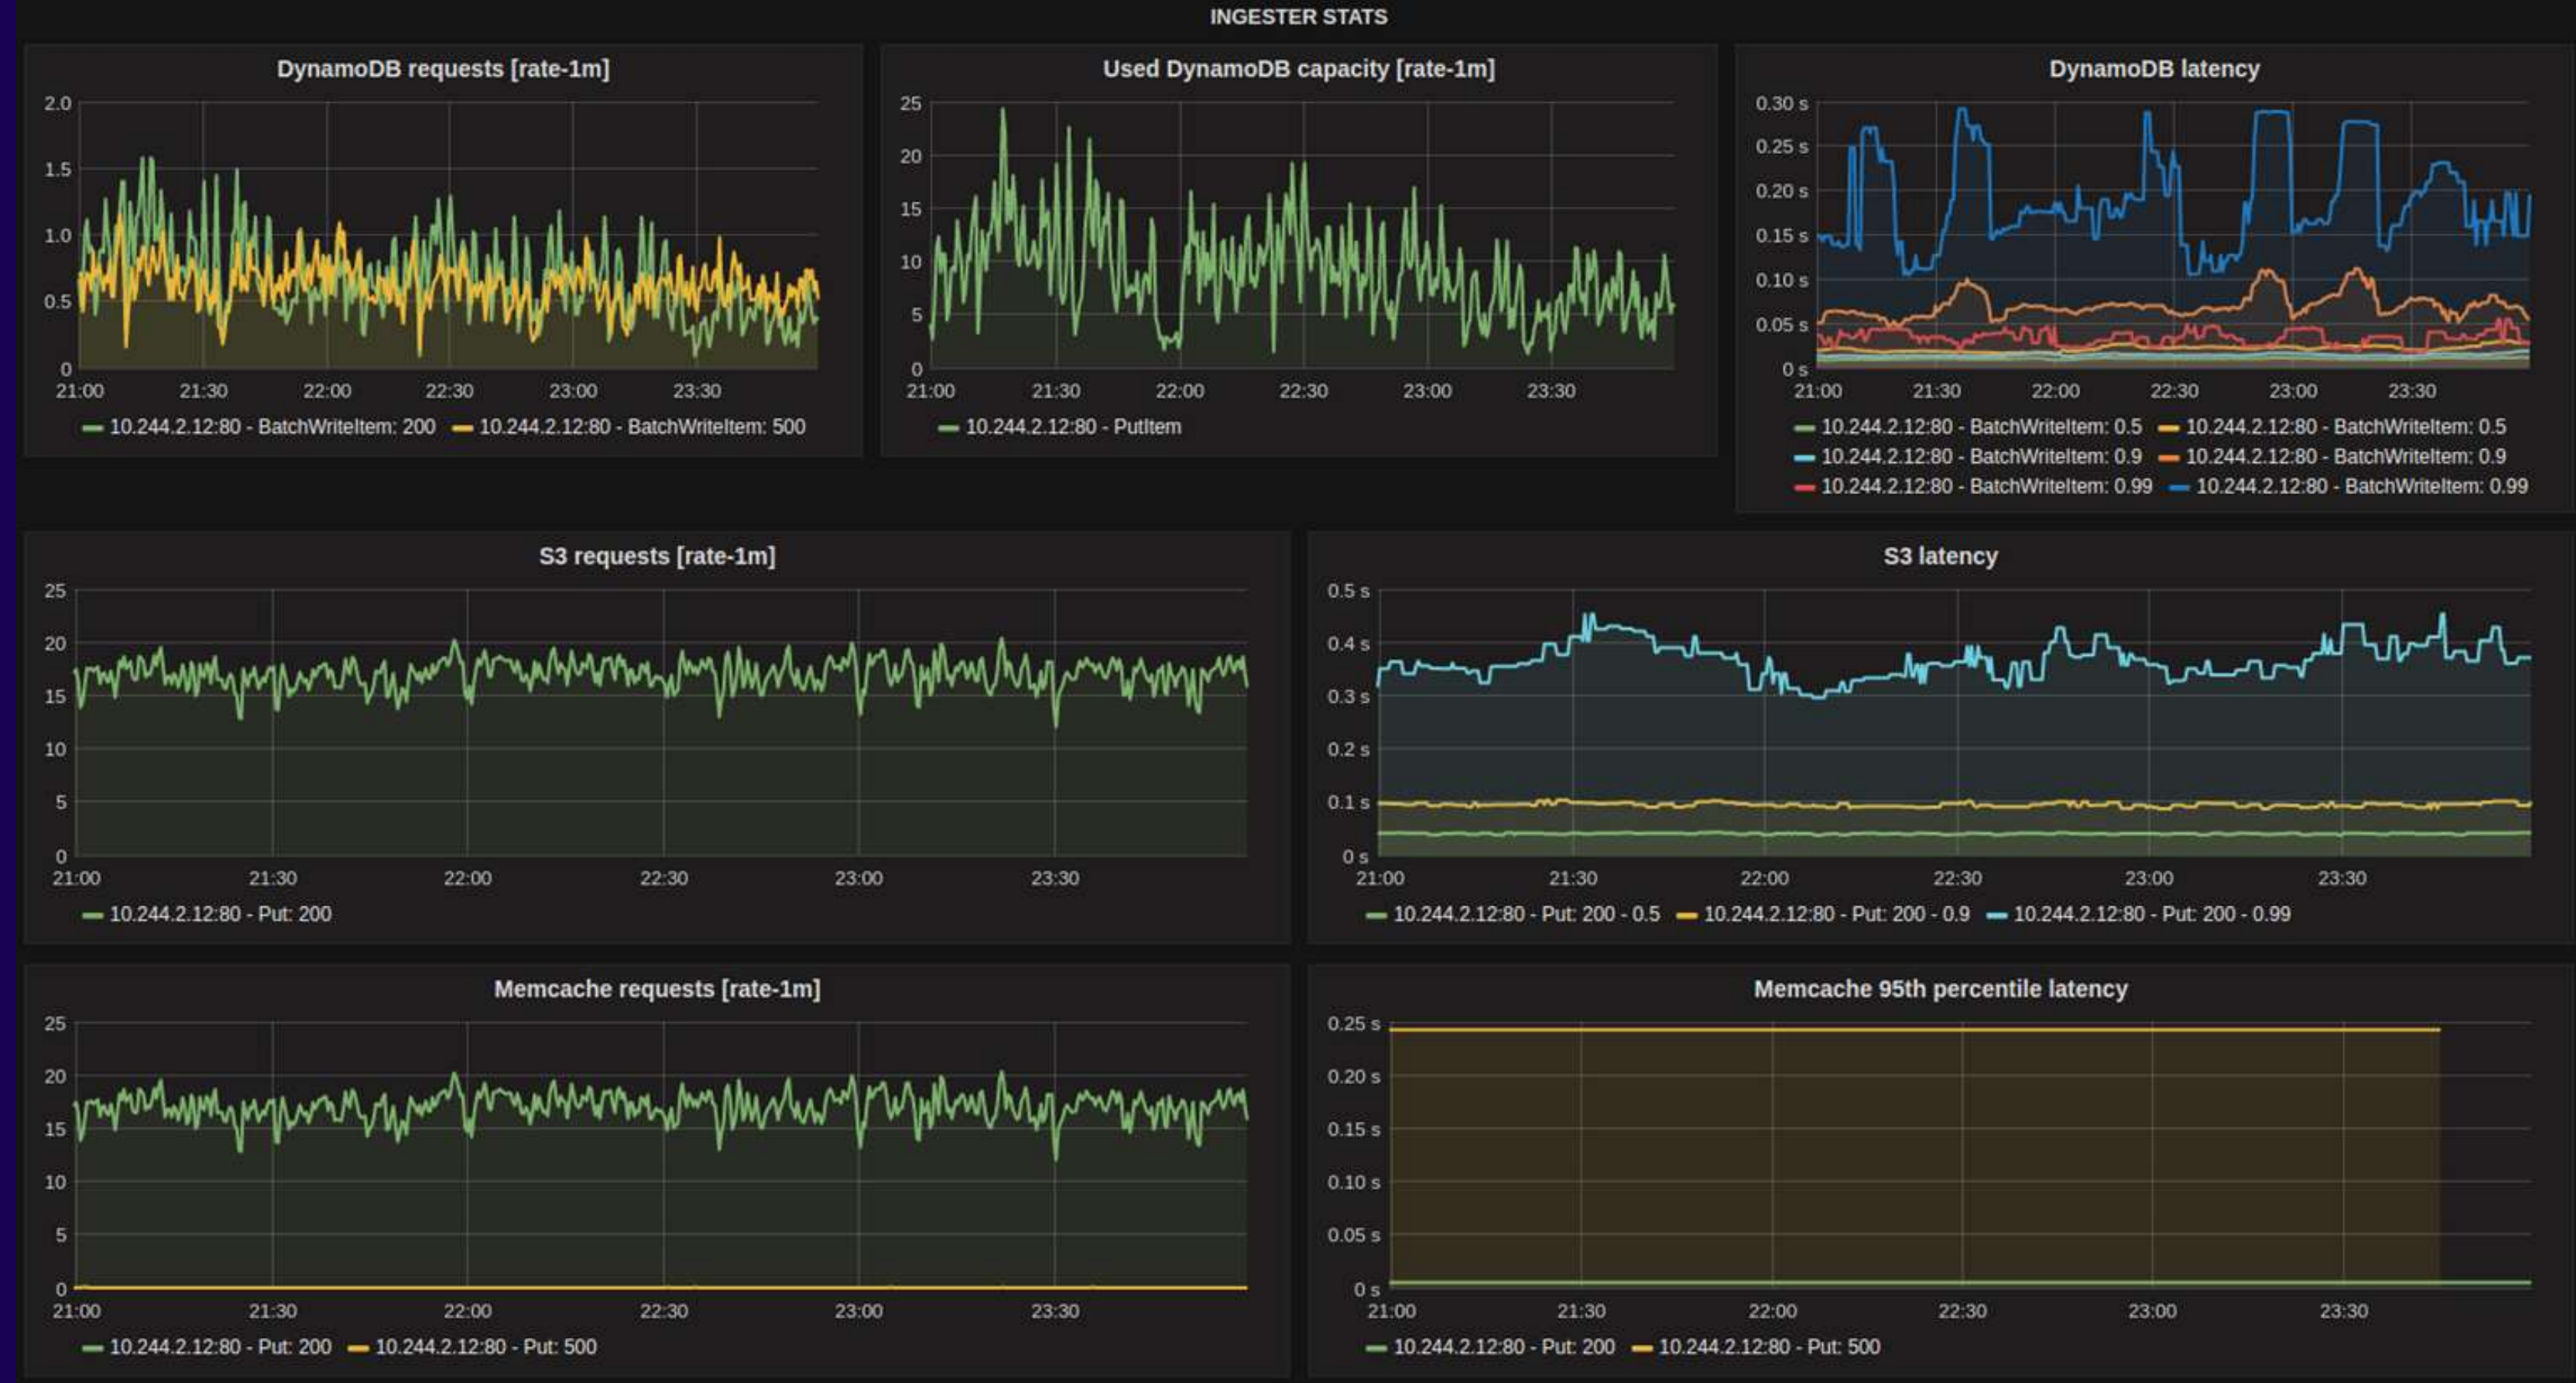
Graph Console

Element	Value
{app="harsh-dagger",proc="api"}	132.720802
{app="quality-locomotive",proc="web"}	89.547081
{app="husky-long-oyster",proc="web"}	68.982738
{app="vital-albatross",proc="api"}	48.033772
{app="autopsy-gutsy",proc="widget"}	47.410583
{app="western-python",proc="cruncher"}	40.126926
{app="harsh-dagger",proc="api"}	28.527714
{app="outstanding-dagger",proc="api"}	26.119423
{app="gruesome-waterbird",proc="web"}	17.666714
{app="gutsy-square",proc="public"}	15.296242
{app="harsh-dagger",proc="web"}	14.738327
{app="northern-electron",proc="api"}	13.349815

Built in Graphing



Grafana Support



Data Model

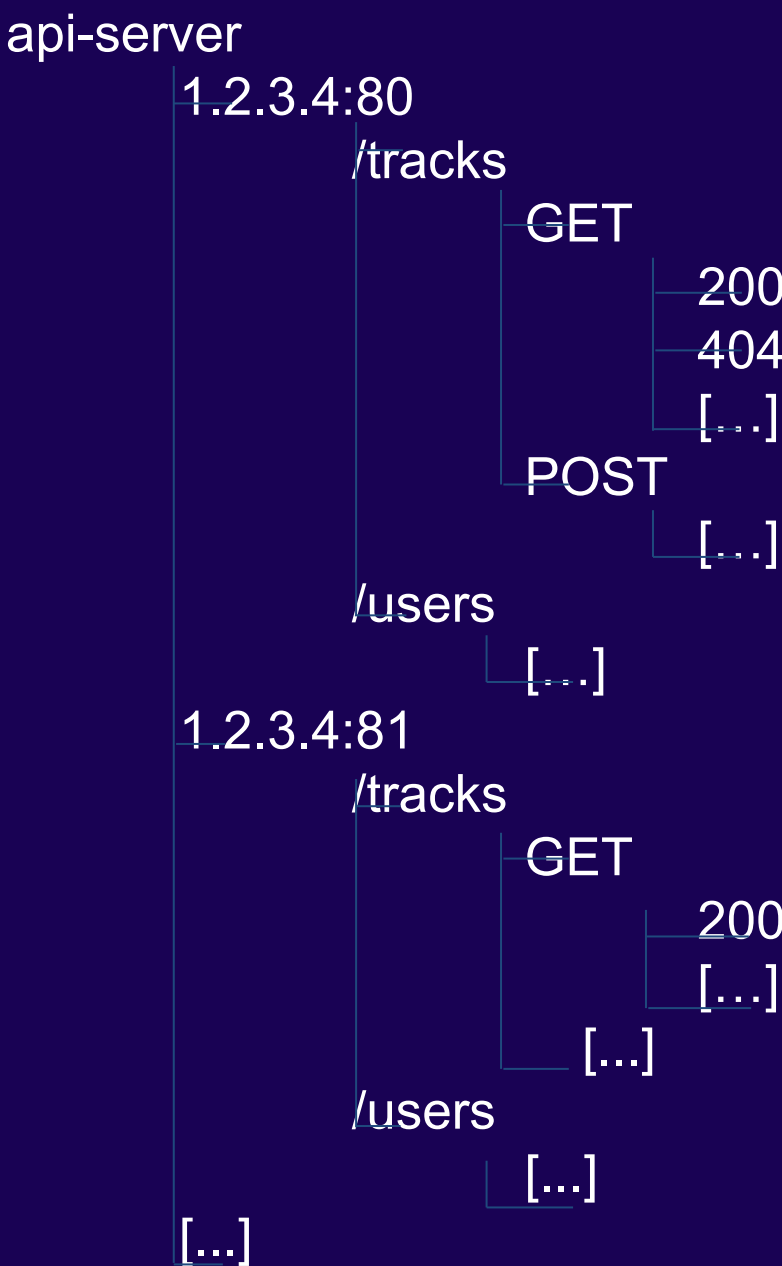
Labels

>

Hierarchy



api_http_requests_total{method="post"}
api-server.*.post.*



Prometheus

- Simple Data Model
 $\text{<identifier>} \rightarrow (t_0, v_0), (t_1, v_1), \dots (t_n, v_n)$
- Essentially a time series data
- Timestamps are in milliseconds
- Examples of these include: (See the series selectors below)

```
http_requests_total{job="nginx", instances="1.2.3.4:80", path="/home", status="200"}  
http_requests_total{job="nginx", instances="1.2.3.4:80", path="/home", status="500"}  
http_requests_total{job="nginx", instances="1.2.3.4:80", path="/settings", status="200"}  
http_requests_total{job="nginx", instances="1.2.3.4:80", path="/settings", status="502"}
```

Prometheus Metrics

Prometheus Metrics

Key-Value store (with BigTable semantics) seems suitable.

KEY			VALUE
Metric name	Labels	Timestamp	Sample Value
...			
http_requests_total	{status="200",method="GET"}	@1434317560938	94355
http_requests_total	{status="200",method="GET"}	@1434317561287	94934
http_requests_total	{status="200",method="GET"}	@1434317562344	96483
http_requests_total	{status="404",method="GET"}	@1434317560938	38473
http_requests_total	{status="404",method="GET"}	@1434317561249	38544
http_requests_total	{status="404",method="GET"}	@1434317562588	38663
http_requests_total	{status="200",method="POST"}	@1434317560885	4748
http_requests_total	{status="200",method="POST"}	@1434317561483	4795
http_requests_total	{status="200",method="POST"}	@1434317562589	4833
http_requests_total	{status="404",method="POST"}	@1434317560939	122
...			

PROMQL Query Language

PromQL: `rate(api_http_requests_total[5m])`

SQL: `SELECT job, instance, method, status, path, rate(value, 5m) FROM api_http_requests_total`

PromQL: `avg by(city) (temperature_celsius{country="germany"})`

SQL: `SELECT city, AVG(value) FROM temperature_celsius WHERE country="germany" GROUP BY city`

PromQL: `rate(errors{job="foo"}[5m]) / rate(total{job="foo"}[5m])`

SQL:

`SELECT errors.job, errors.instance, [...more labels...], rate(errors.value, 5m) / rate(total.value, 5m)
FROM errors JOIN total ON [...all the label equalities...] WHERE errors.job="foo" AND total.job="foo"`

PROMQL Query Language

- PromQL has a number of features.
- It can select a vector of values, use functions **and**
- Aggregate by dimension e.g.
 - *sum by (path) (rate(http_requests_total{job="nginx",status =~ "5.."}[1m]))*
- **And** do binary operations e.g.
 - *sum by (path) (rate(http_requests_total{job="nginx",status =~ "5.."}[1m])) / sum by (path) (rate(http_requests_total{job="nginx"}[1m]))*

Metrics

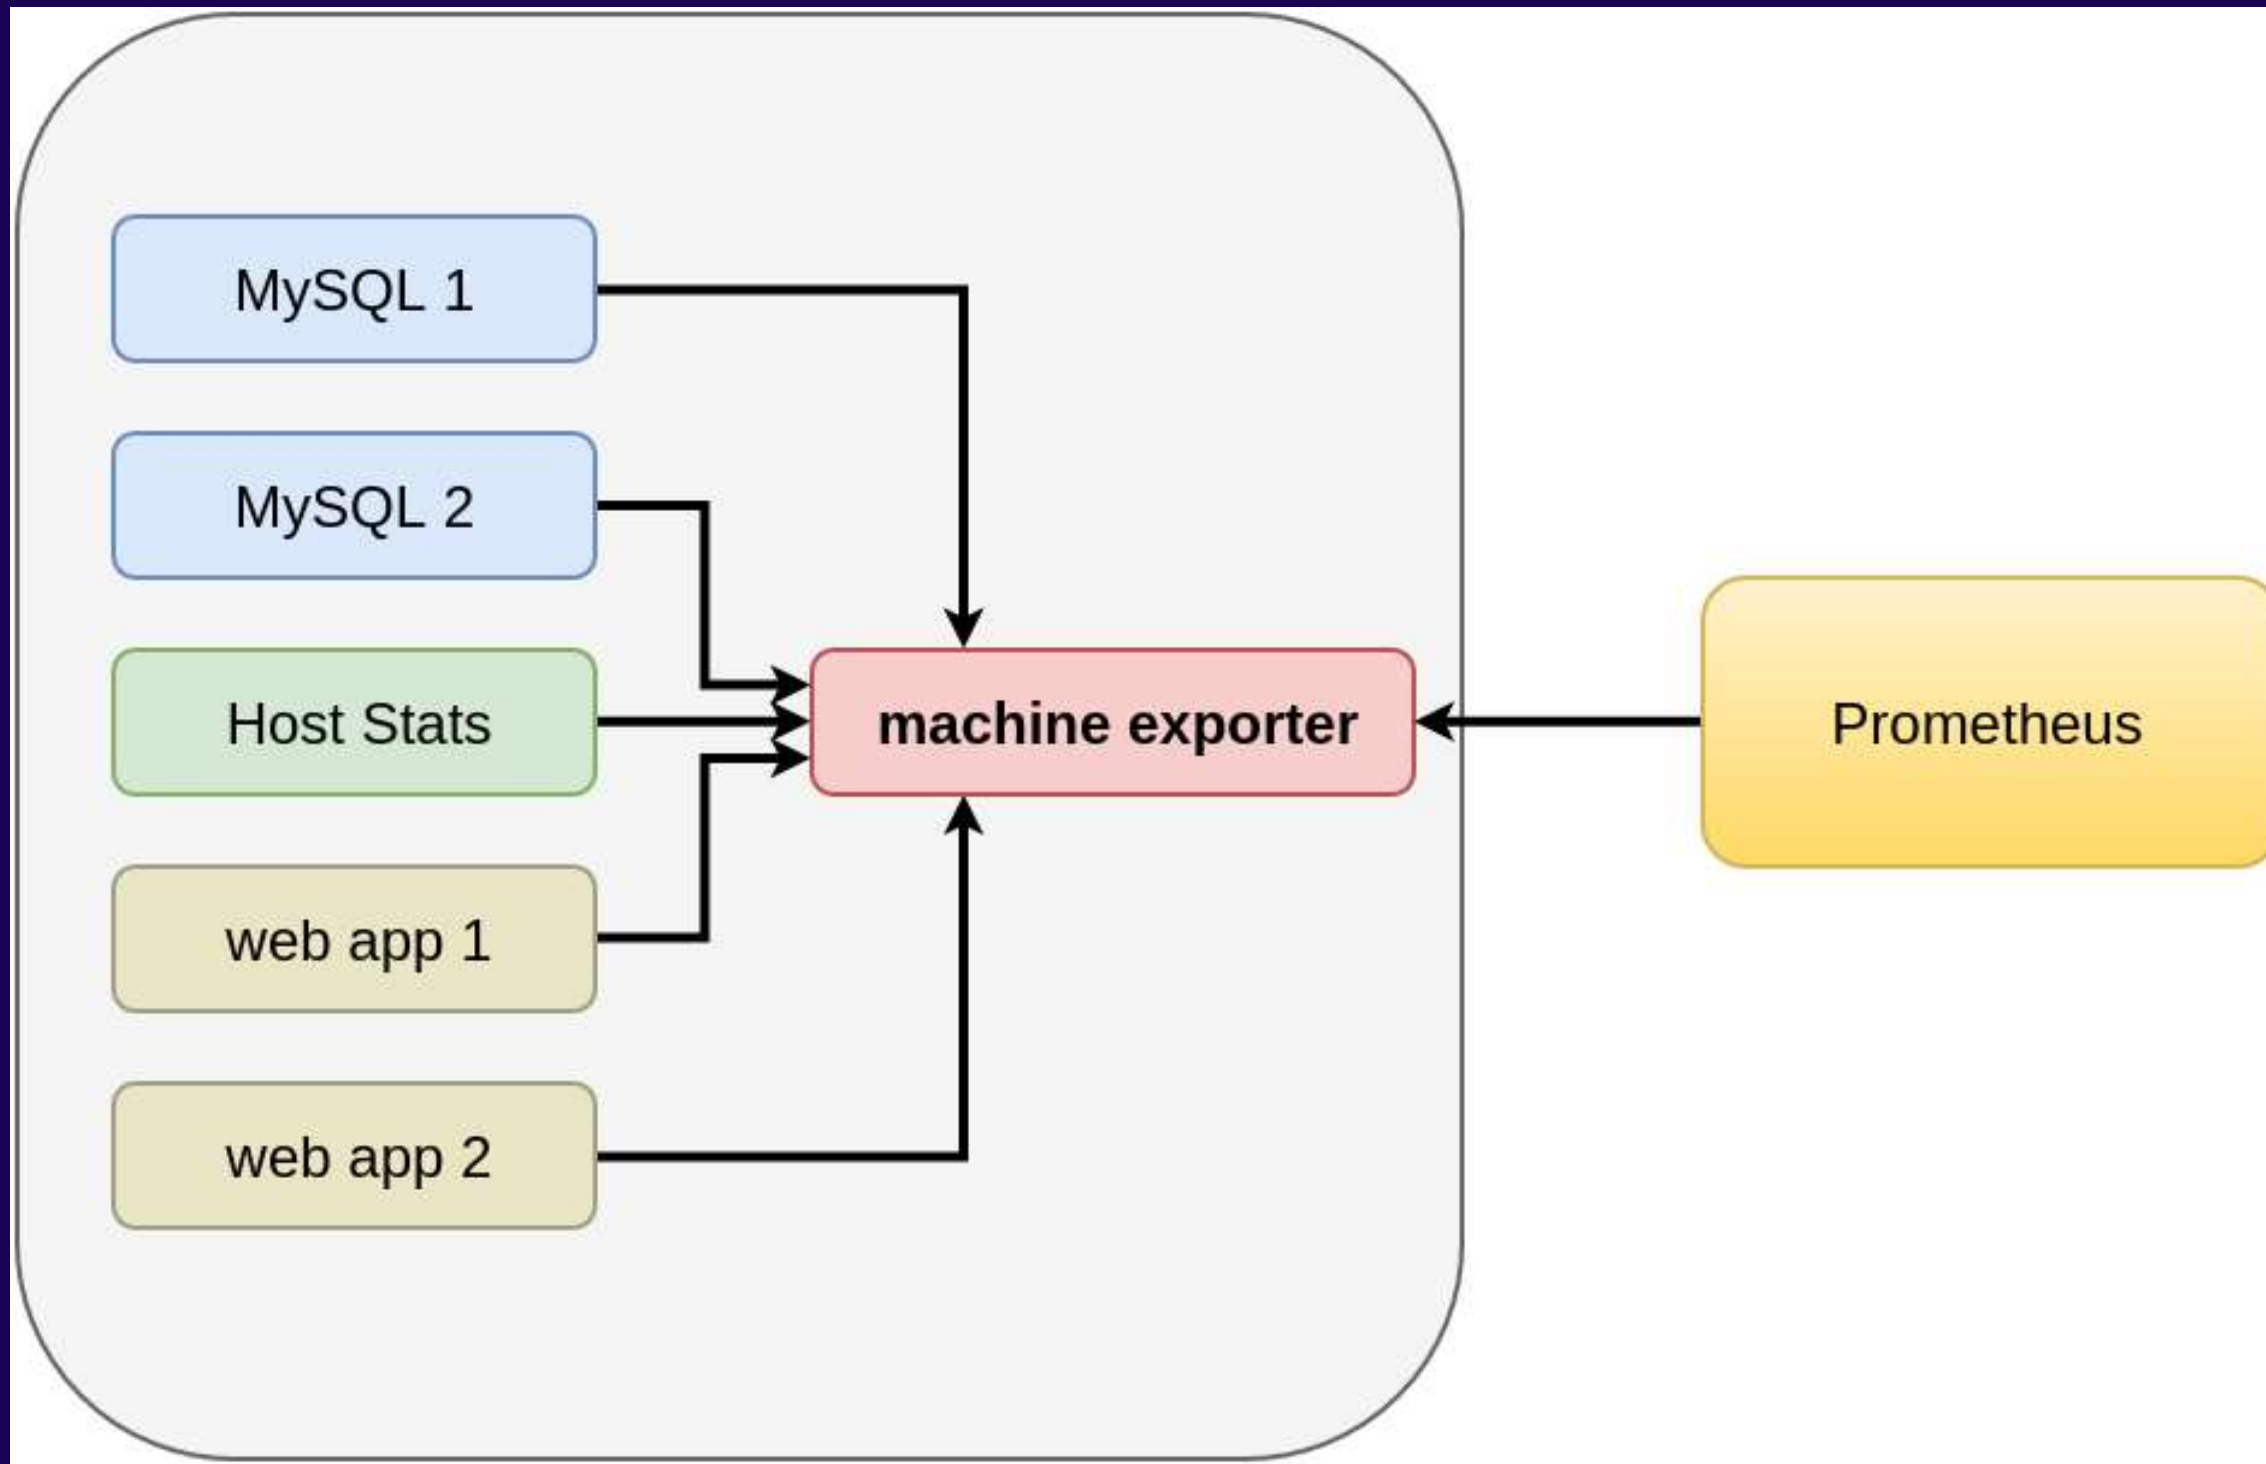
Category of Metrics

- USE
- RED
- AD-HOC

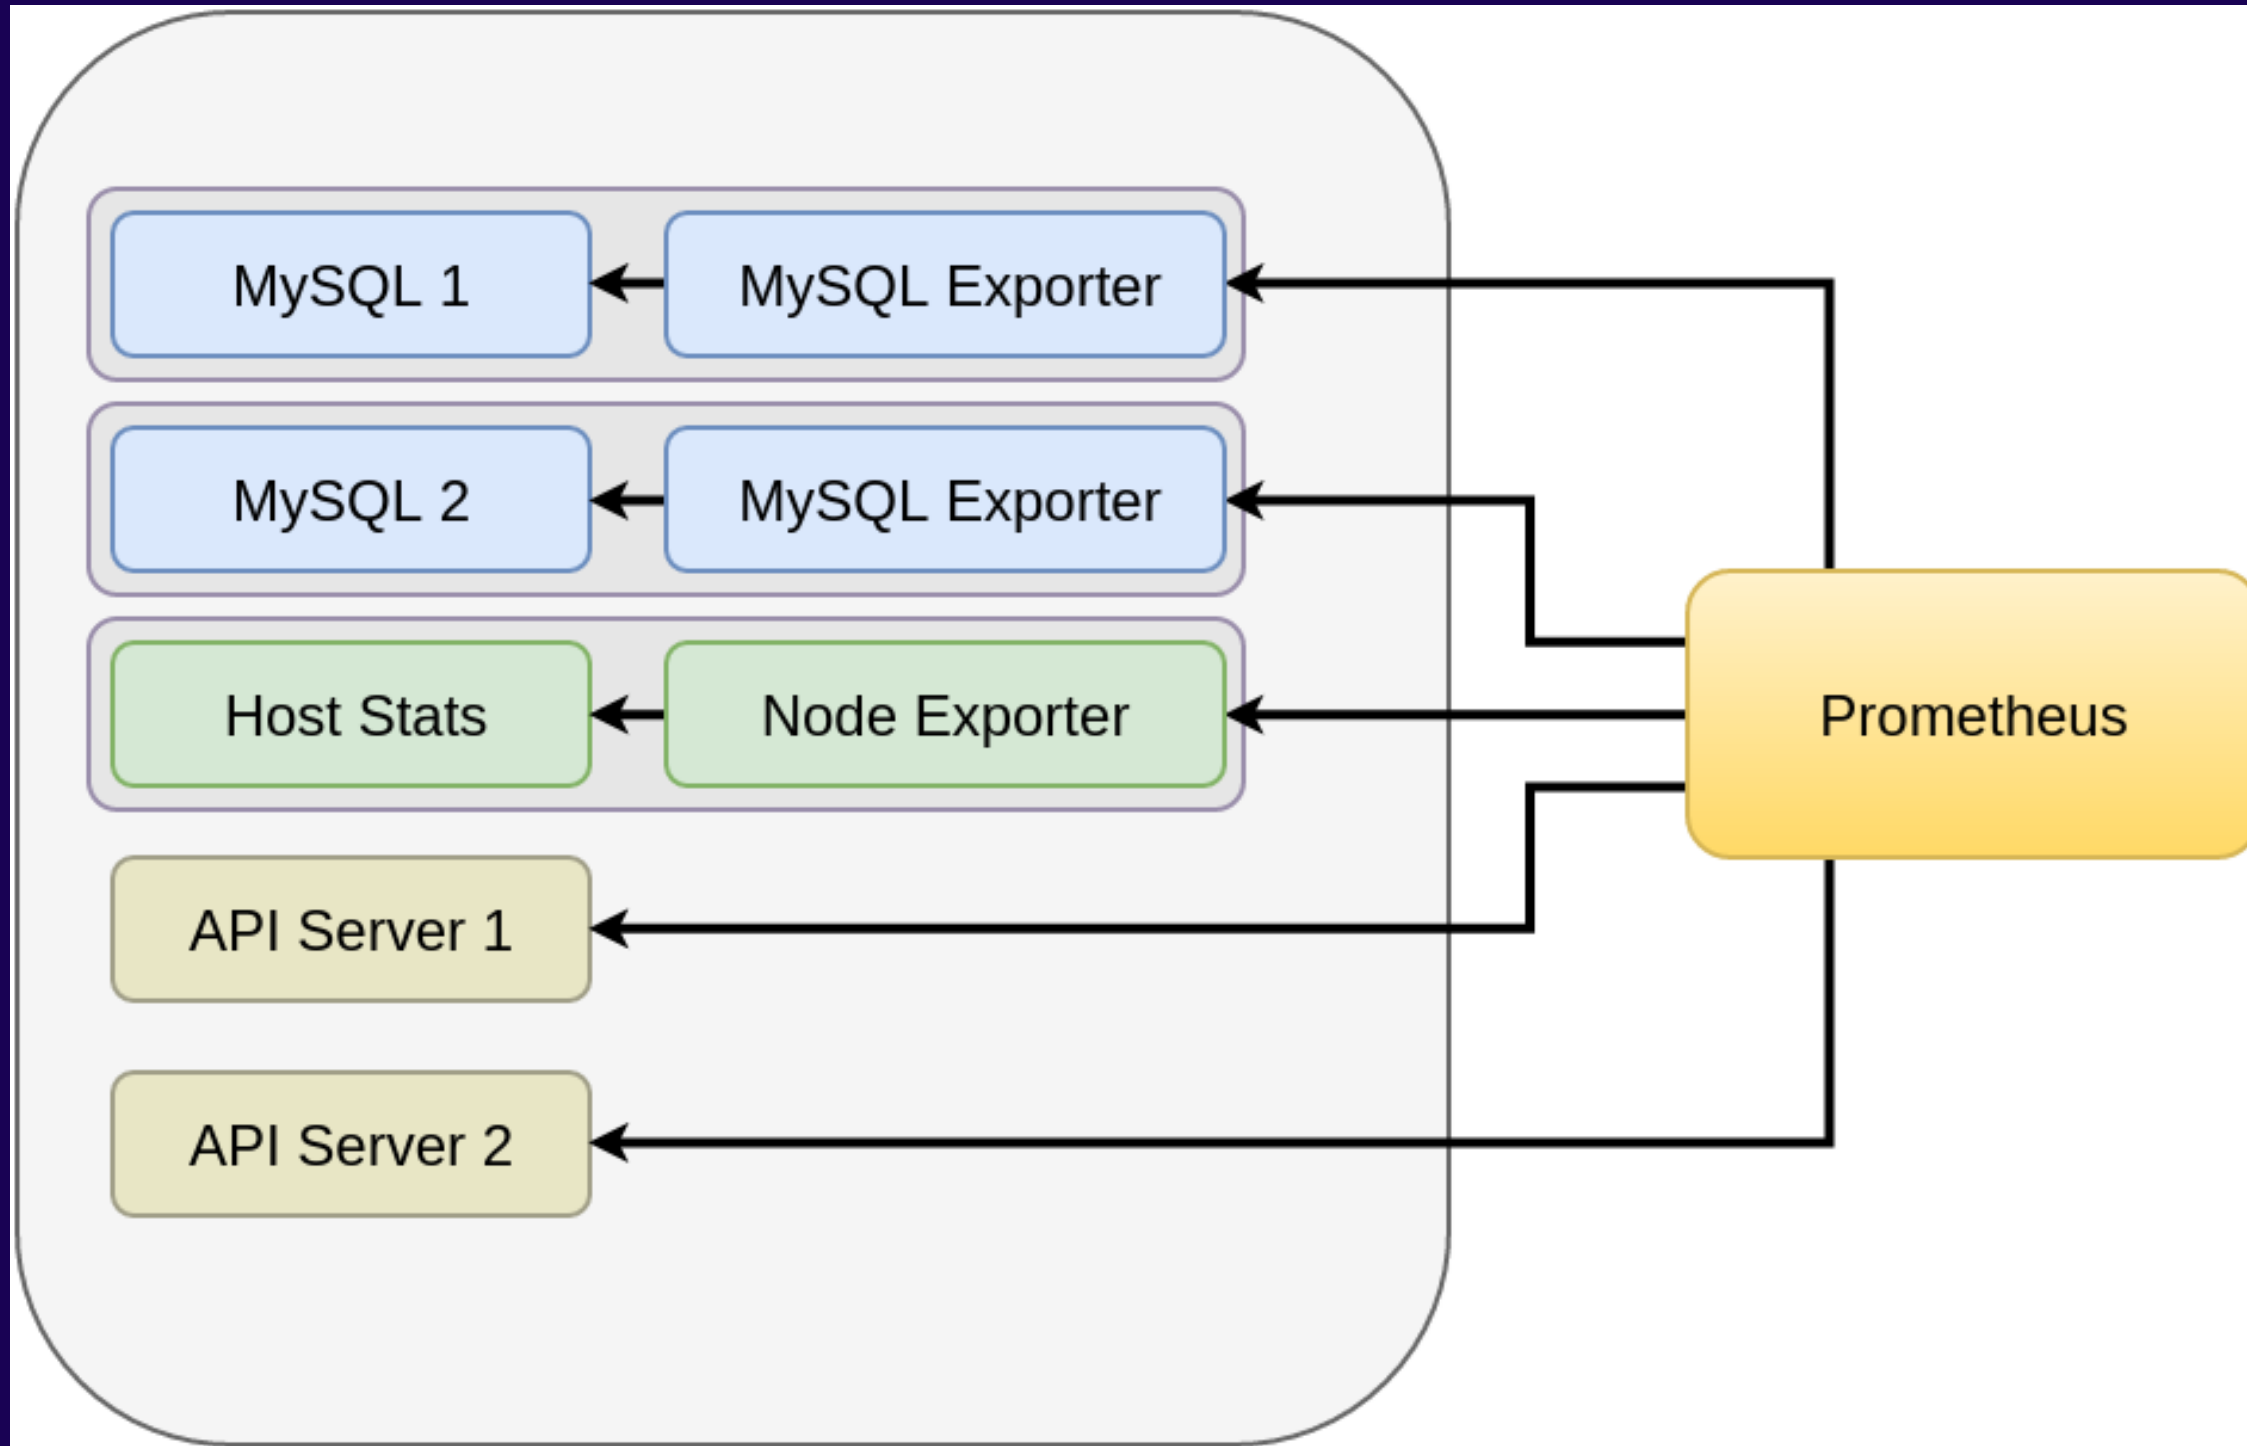
Type of Metrics

- Count
- Gauge
- Histogram

Exporters



Exporters



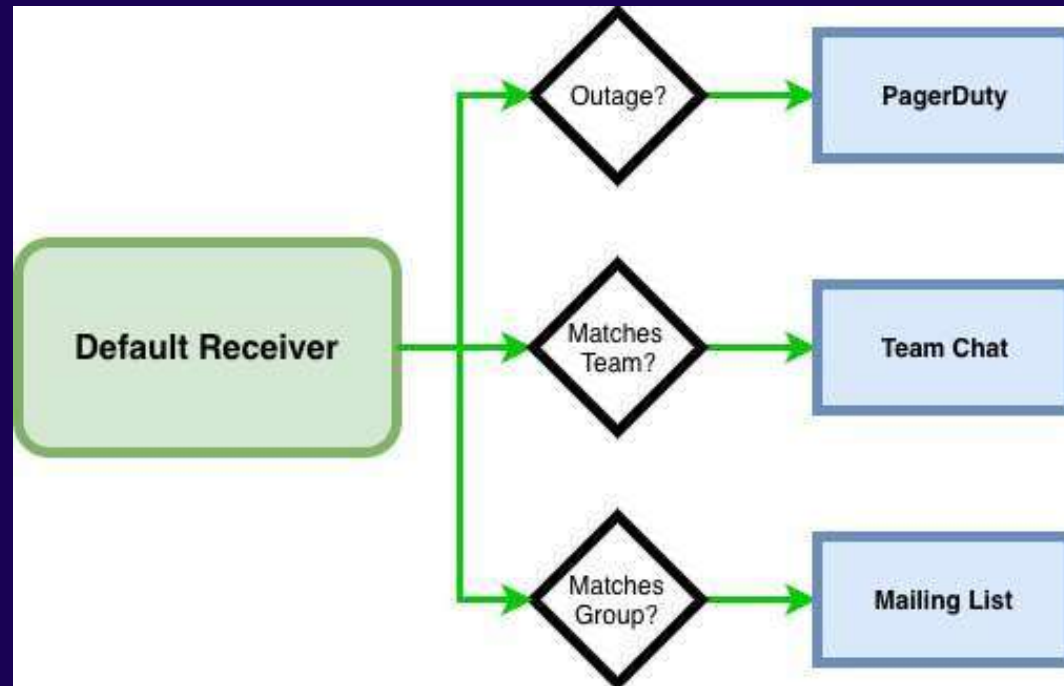
Exporters

A process that exposes Metrics for Prometheus to harvest



The available exporters can be find from here: <https://prometheus.io/docs/instrumenting/exporters/>

Alert Manager Rules



- AlertManager rules are conceptualized as routes, giving you the ability to write sophisticated sets of rules to determine where notifications should end up
- A default receiver should be configured for every notification, and then additional services can be configured through child routes which will match certain conditions

A full configuration reference is available here:

<https://prometheus.io/docs/alerting/configuration>

Alert Manager Rules

- Our config YAML file will be responsible for **setting up routing rules** that will determine how events are triaged
- As mentioned before, all events should **start with a default receiver**, called default-receiver, which will be the starting point for any route
- From there, any number of **sub-receivers can be configured**
- **Sample Configuration one called 'slack'** which will be invoked when the "service" tag of the event that has been triggered matches "activemq"
- Next, **configure our receivers**
- Sample Slack receiver config will **contain WebHook** into Slack

```
global:
  smtp_smarthost: 'localhost:25'
  smtp_from: 'alertmanager@monitoring.com'

route:
  receiver: 'default-receiver'
  group_wait: 30s
  group_interval: 5m
  repeat_interval: 4h
  group_by: [cluster, alertname]
  routes:
    - receiver: 'slack'
      group_wait: 10s
      match_re:
        service: activemq

receivers:
  - name: 'default-receiver'
    email_configs:
      - to: 'justin.reock@roguewave.com'

  - name: 'slack'
    slack_configs:
      - api_url: https://hooks.slack.com/services/
        channel: '#general'
```

Alert Manager Rules

Configure Sample Rules

- **two simple events**, but, events can be created out of a **huge range of possible query configurations**

Alert /

```
groups:
- name: activemq
  rules:
  - alert: DLQ
    expr: org_apache_activemq_Broker_DLQ > 1
    for: 1m
    labels:
      severity: minor
      service: activemq
    annotations:
      summary: A message has gone into the DLQ
      dashboard: http://192.168.40.120:3000/dashboard/db/activemq-broker
      impact: A message has been misfired
      runbook: http://activemq.apache.org
  - alert: Broker Down
    expr: up{job="activemq"} == 0
    labels:
      severity: major
      service: activemq
    annotations:
      summary: The broker has crashed
      dashboard: http://192.168.40.120:3000/dashboard/db/activemq-broker
      impact: Broker is down
      runbook: http://activemq.apache.org
```

Integrating with Prometheus

15s

- **configure Prometheus to push** alert events into AlertManager
 - Add an alerting section to the Prometheus YAML file
- **Update prom- amq.yml** configuration file from earlier to integrate with our newly configured AlertManager instance
- Upon restarting Prometheus, we **should see our alerts** in the Prometheus dashboard