

Getting started with Prometheus & Grafana

Seshagiri Sriram - Aug 26 2028

Prometheus



- A monitoring & alerting system, Inspired by Google's BorgMon
- Originally built by SoundCloud in 2012
- Open Source, now part of the CNCF
- Simple text-based metrics format
- Multidimensional data model
- Rich, concise query language

Prometheus



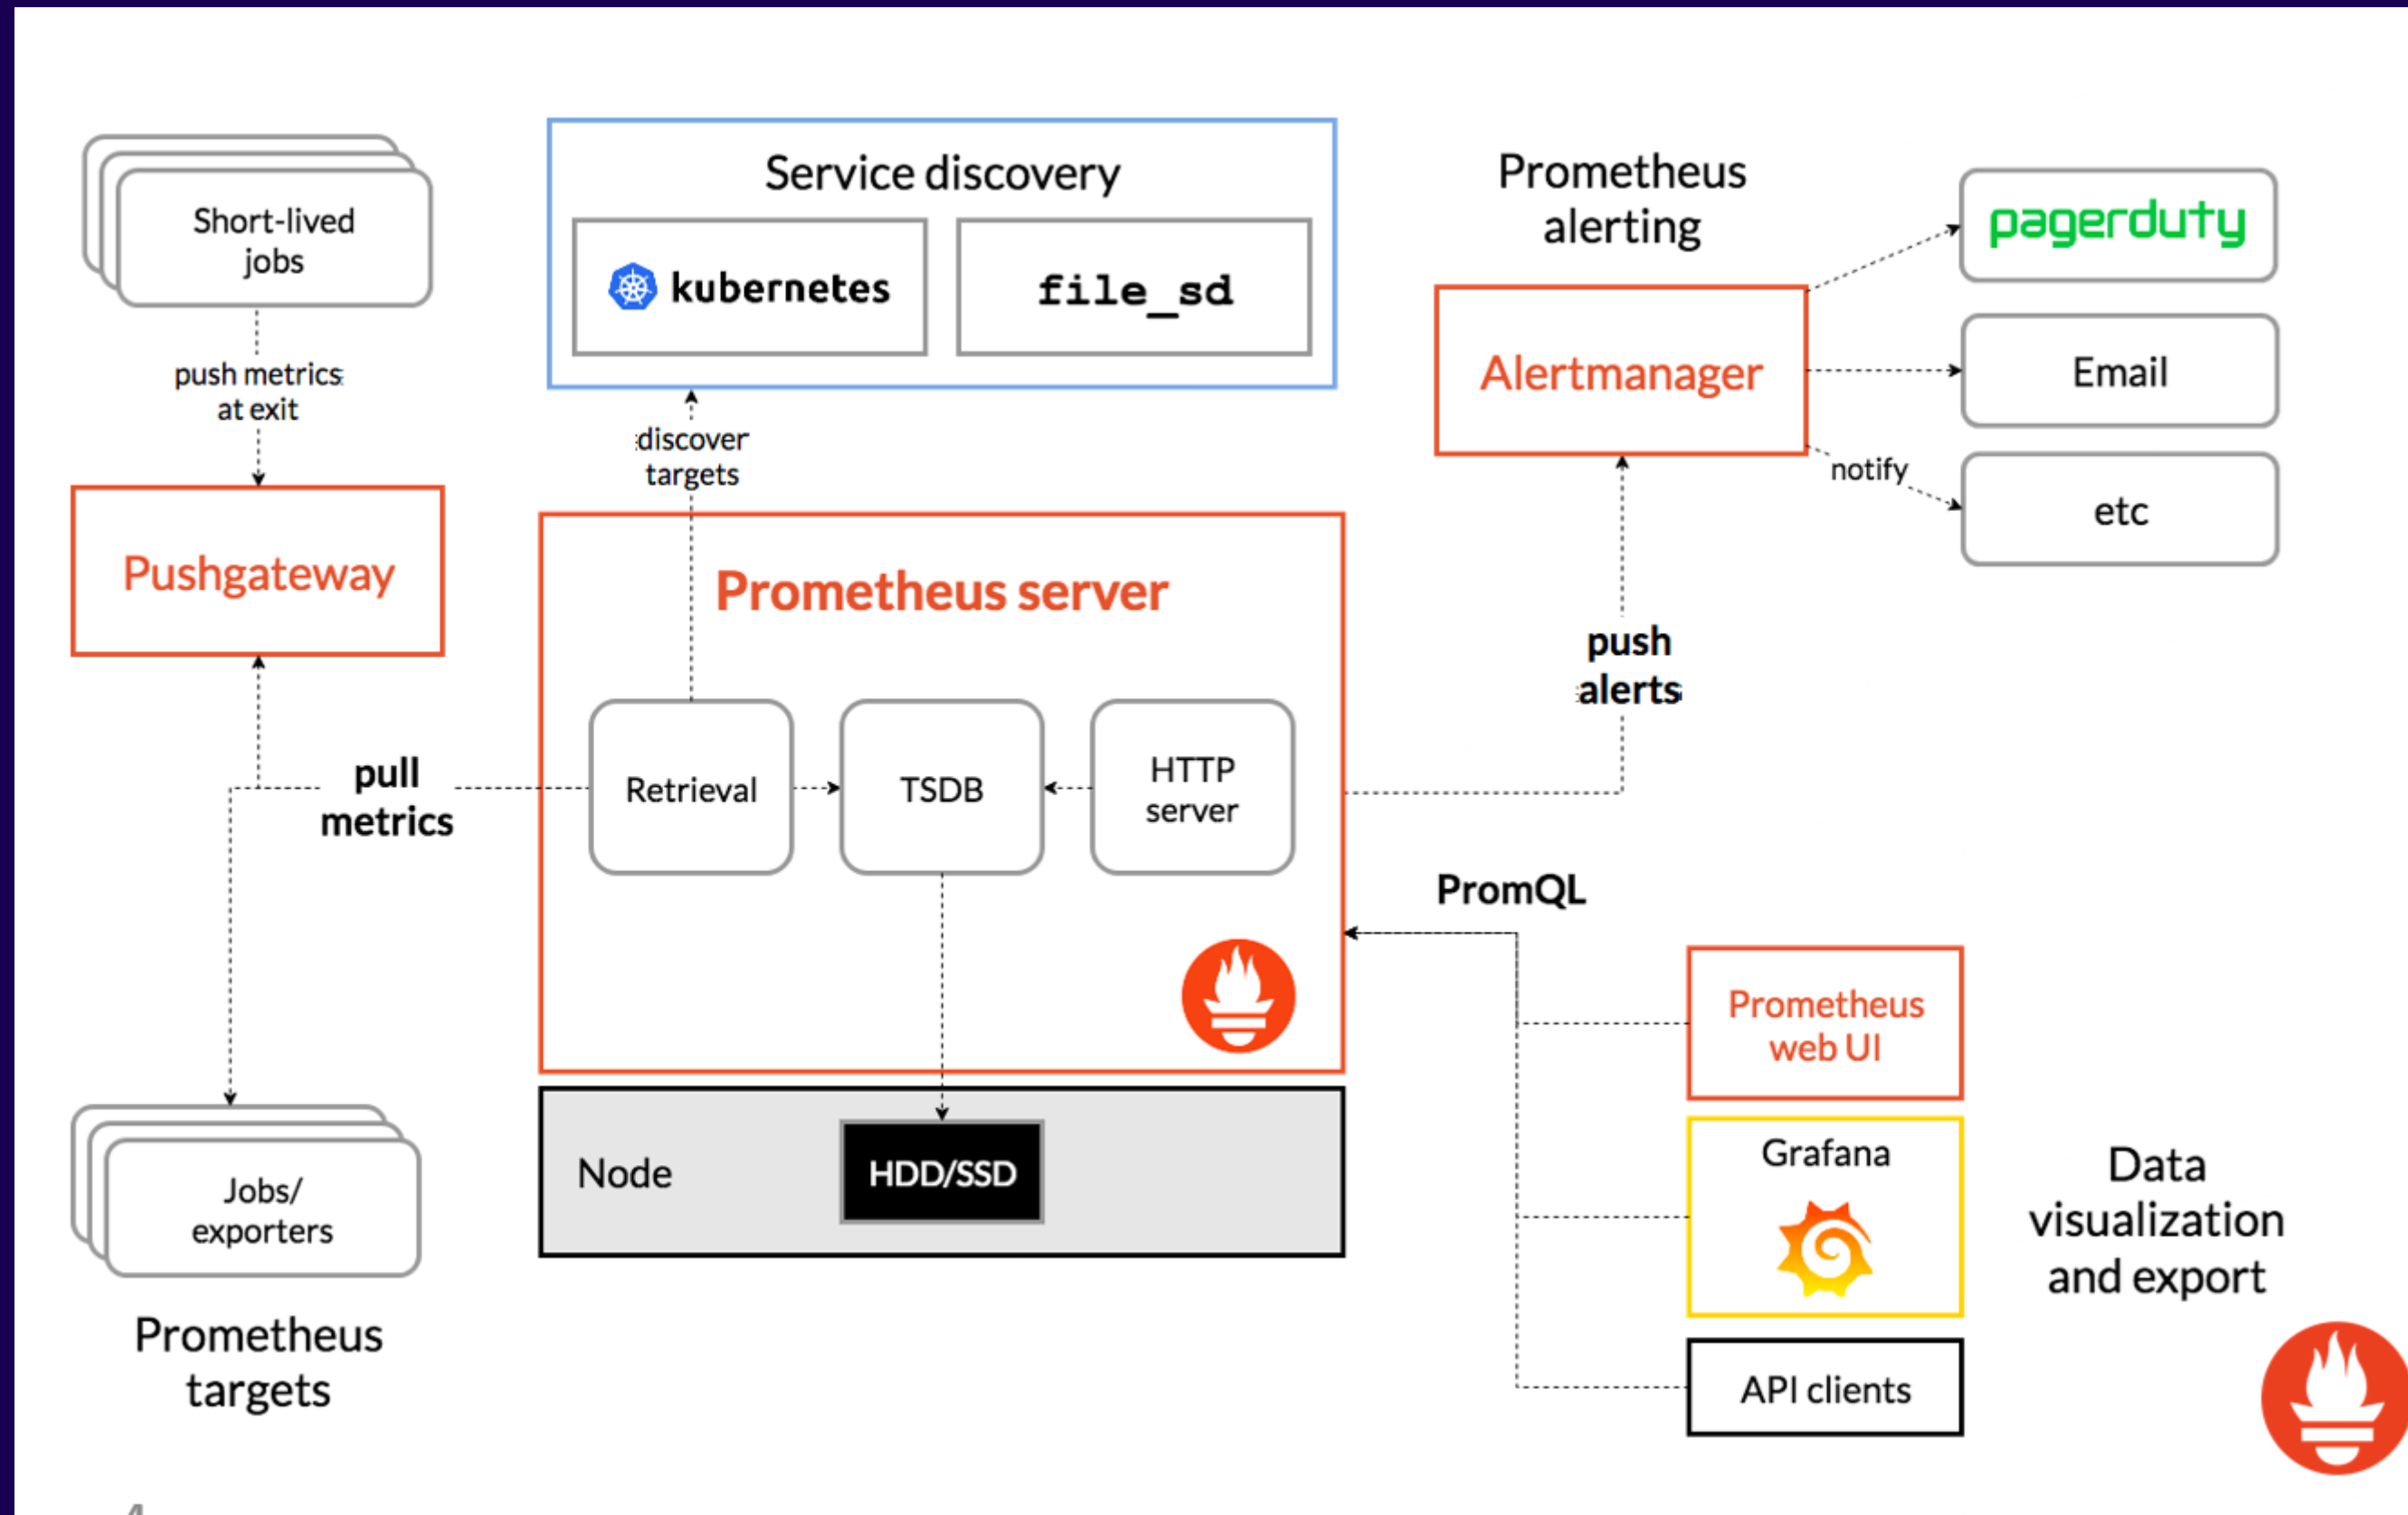
- A data scraper that pulls metrics data over HTTP periodically at a configured interval.
- A time-series database to store all the metrics data.
- A simple user interface where you can visualize, query, and monitor all the metrics.
- Written in Go, fully published in 2015.

Prometheus

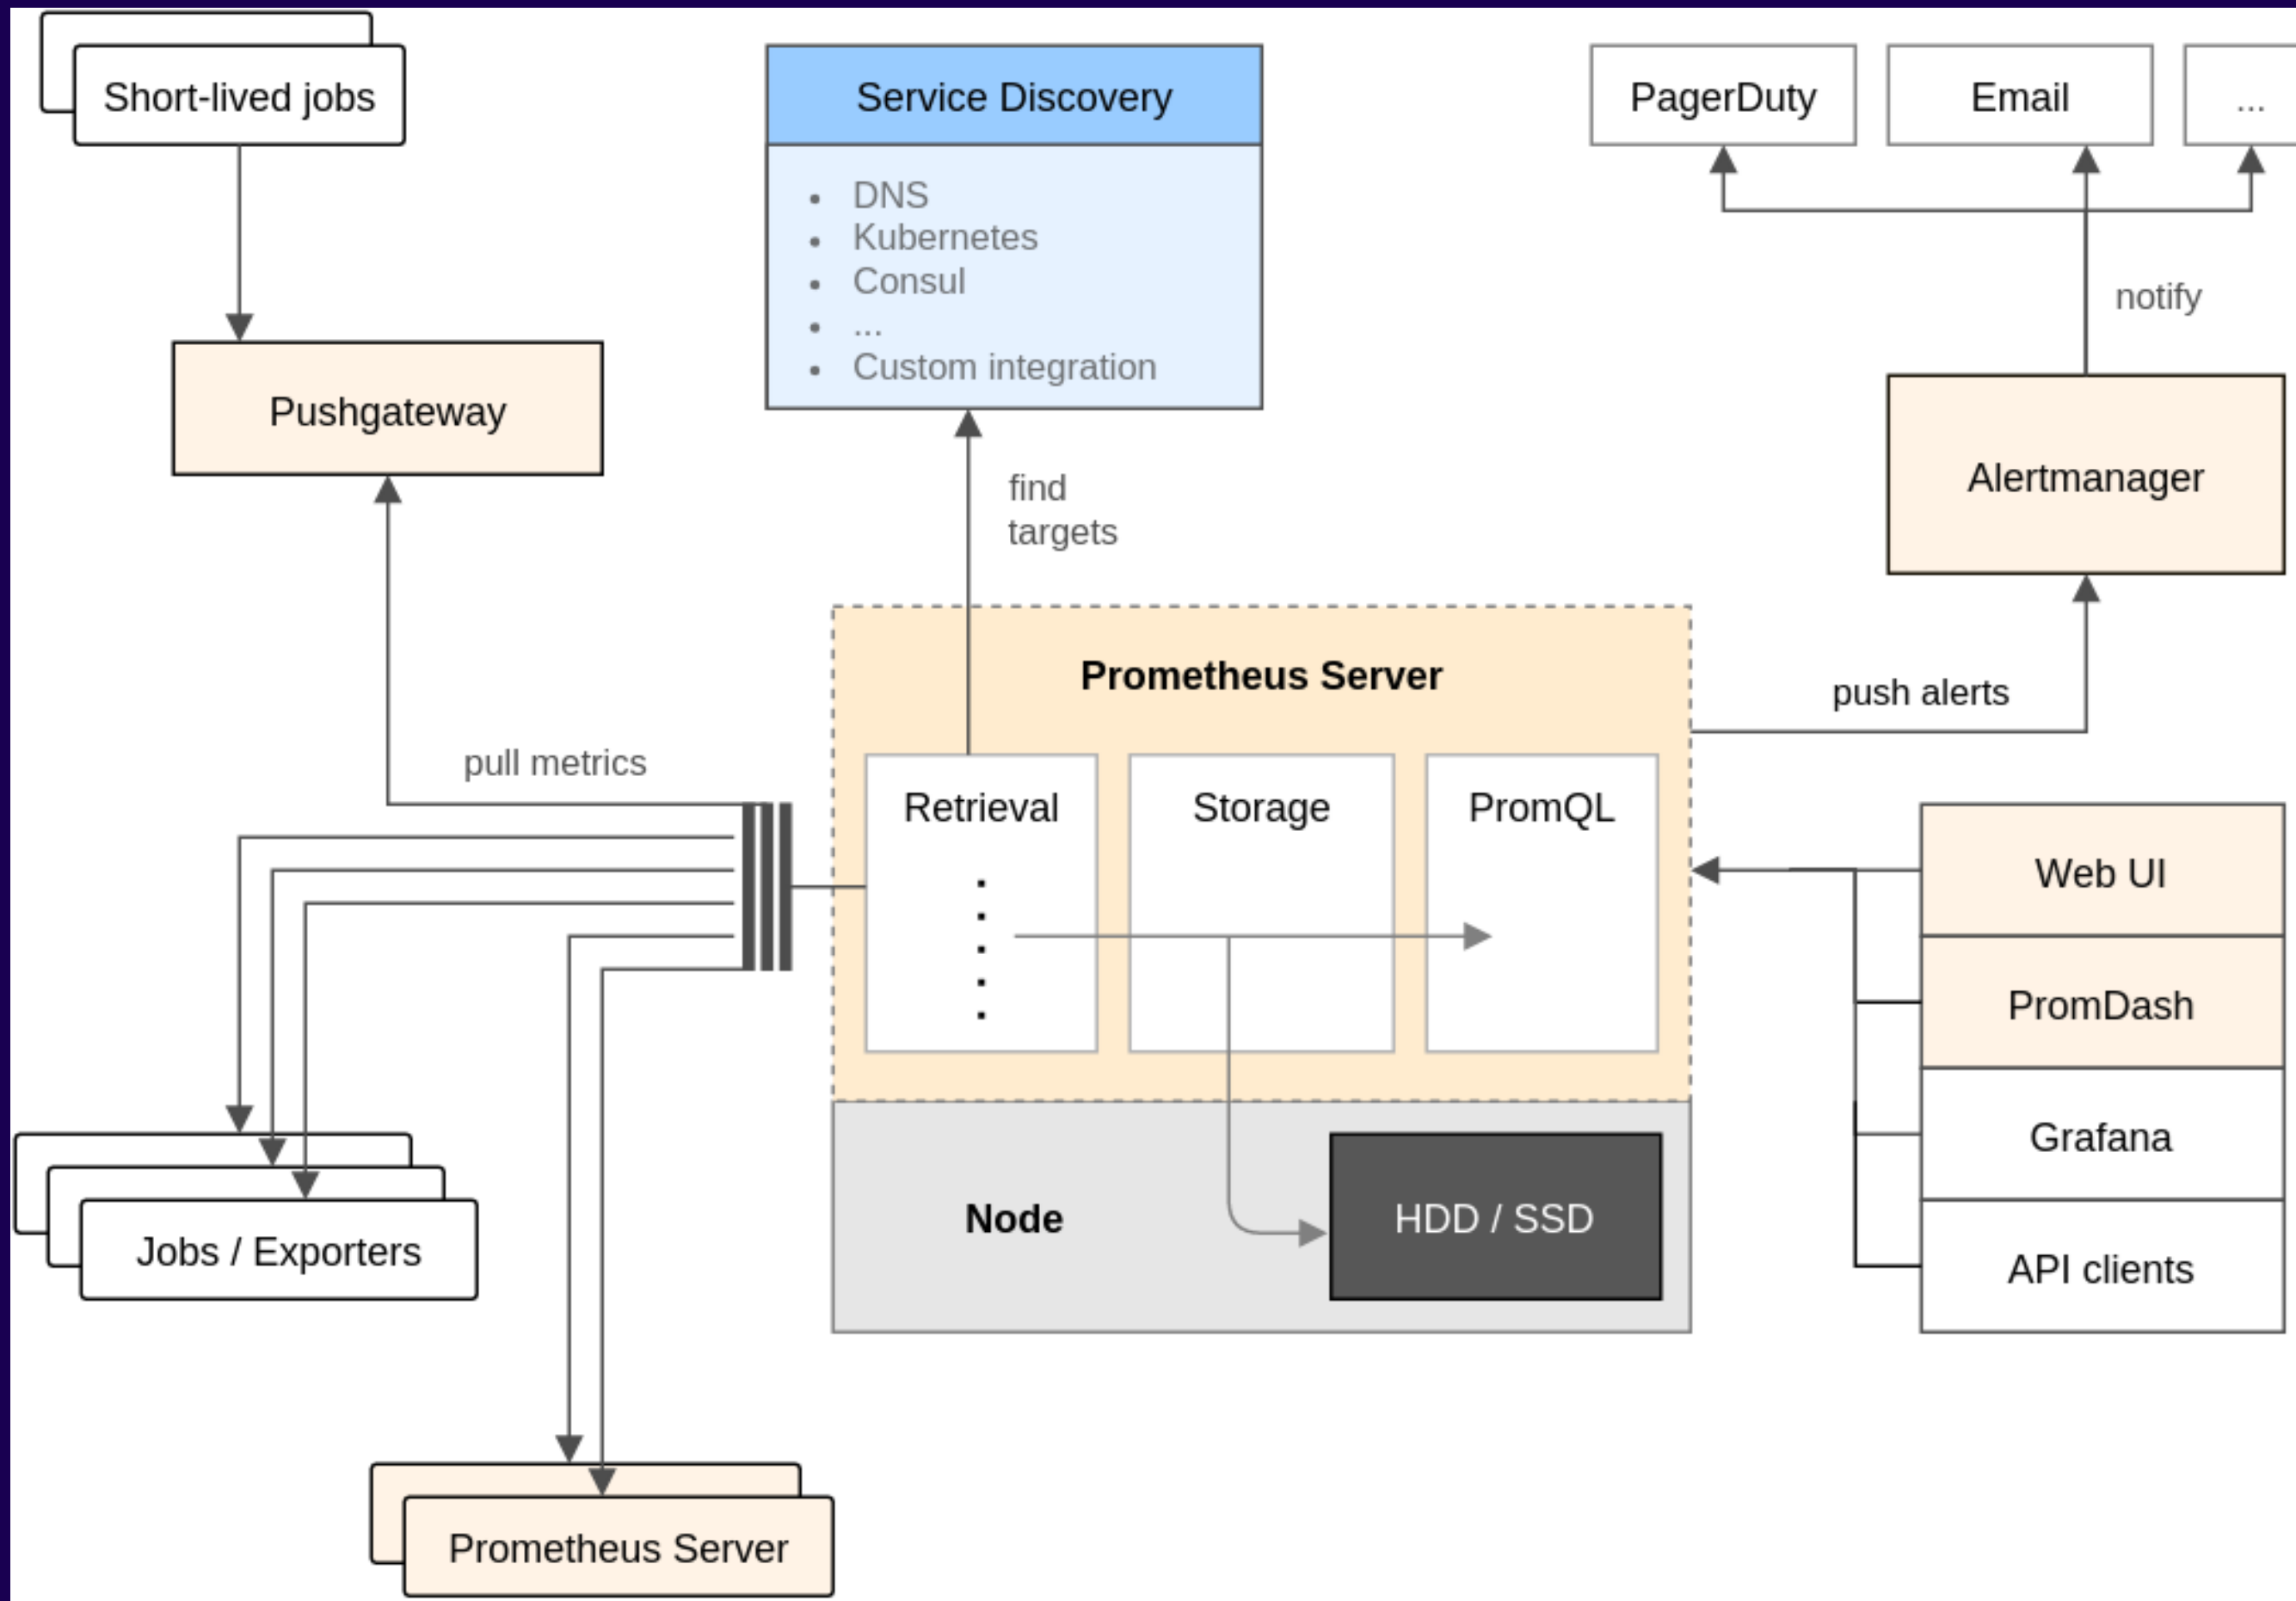


- Monitoring systems and TSDB
 - Instrumentation
 - Metrics collection and Storage
 - Querying
 - Alerting
 - Dashboarding / Graphing / Trending
- Focus on
 - Dynamic Cloud Environments
 - Operational Systems Monitoring

Prometheus



Prometheus



Prometheus



What I can do

Dimensional Data Model

Powerful Query Language

Efficiency

Operational Simplicity

What it cannot do

Raw Log/event Collection

Request Tracing

Anomaly Detection

Automatic horizontal scaling

User Management and authentication has to be handled separately

Expression Browser

Prometheus Alerts Graph Status Help

sort_desc(sum(bazooka_instance_memory_limit_bytes - bazooka_instance_memory_usage_bytes) by (app, proc)) / 1024 / 1024 / 1024

Execute

Graph Console

Element	Value
{app="harsh-dagger",proc="api"}	132.720802
{app="quality-locomotive",proc="web"}	89.547081
{app="husky-long-oyster",proc="web"}	68.982738
{app="vital-albatross",proc="api"}	48.033772
{app="autopsy-gutsy",proc="widget"}	47.410583
{app="western-python",proc="cruncher"}	40.126926
{app="harsh-dagger",proc="api"}	28.527714
{app="outstanding-dagger",proc="api"}	26.119423
{app="gruesome-waterbird",proc="web"}	17.666714
{app="gutsy-square",proc="public"}	15.296242
{app="harsh-dagger",proc="web"}	14.738327
{app="northern-electron",proc="api"}	13.349815


Built in Graphing



Grafana Support

INGESTER STATS


DynamoDB requests [rate-1m]



Used DynamoDB capacity [rate-1m]




DynamoDB latency



S3 requests [rate-1m]



S3 latency



Memcache requests [rate-1m]



Memcache 95th percentile latency



Data Model


Labels

>

Hierarchy



api_http_requests_total{method="post"}
api-server.*.post.*



Prometheus

- Simple Data Model
 $\text{<identifier>} \rightarrow (t_0, v_0), (t_1, v_1), \dots (t_n, v_n)$
- Essentially a time series data
- Timestamps are in milliseconds
- Examples of these include: (See the series selectors below)

```
http_requests_total{job="nginx", instances="1.2.3.4:80", path="/home", status="200"}  
http_requests_total{job="nginx", instances="1.2.3.4:80", path="/home", status="500"}  
http_requests_total{job="nginx", instances="1.2.3.4:80", path="/settings", status="200"}  
http_requests_total{job="nginx", instances="1.2.3.4:80", path="/settings", status="502"}
```

Prometheus Metrics

Prometheus Metrics

Key-Value store (with BigTable semantics) seems suitable.

KEY			VALUE
Metric name	Labels	Timestamp	Sample Value
...			
http_requests_total	{status="200",method="GET"}	@1434317560938	94355
http_requests_total	{status="200",method="GET"}	@1434317561287	94934
http_requests_total	{status="200",method="GET"}	@1434317562344	96483
http_requests_total	{status="404",method="GET"}	@1434317560938	38473
http_requests_total	{status="404",method="GET"}	@1434317561249	38544
http_requests_total	{status="404",method="GET"}	@1434317562588	38663
http_requests_total	{status="200",method="POST"}	@1434317560885	4748
http_requests_total	{status="200",method="POST"}	@1434317561483	4795
http_requests_total	{status="200",method="POST"}	@1434317562589	4833
http_requests_total	{status="404",method="POST"}	@1434317560939	122
...			

PROMQL Query Language

PromQL: `rate(api_http_requests_total[5m])`

SQL: `SELECT job, instance, method, status, path, rate(value, 5m) FROM api_http_requests_total`

PromQL: `avg by(city) (temperature_celsius{country="germany"})`

SQL: `SELECT city, AVG(value) FROM temperature_celsius WHERE country="germany" GROUP BY city`

PromQL: `rate(errors{job="foo"}[5m]) / rate(total{job="foo"}[5m])`

SQL:

`SELECT errors.job, errors.instance, [...more labels...], rate(errors.value, 5m) / rate(total.value, 5m)
FROM errors JOIN total ON [...all the label equalities...] WHERE errors.job="foo" AND total.job="foo"`

PROMQL Query Language

- PromQL has a number of features.
- It can select a vector of values, use functions **and**
- Aggregate by dimension e.g.
 - *sum by (path) (rate(http_requests_total{job="nginx",status =~ "5.."}[1m]))*
- **And** do binary operations e.g.
 - *sum by (path) (rate(http_requests_total{job="nginx",status =~ "5.."}[1m])) / sum by (path) (rate(http_requests_total{job="nginx"}[1m]))*

Metrics

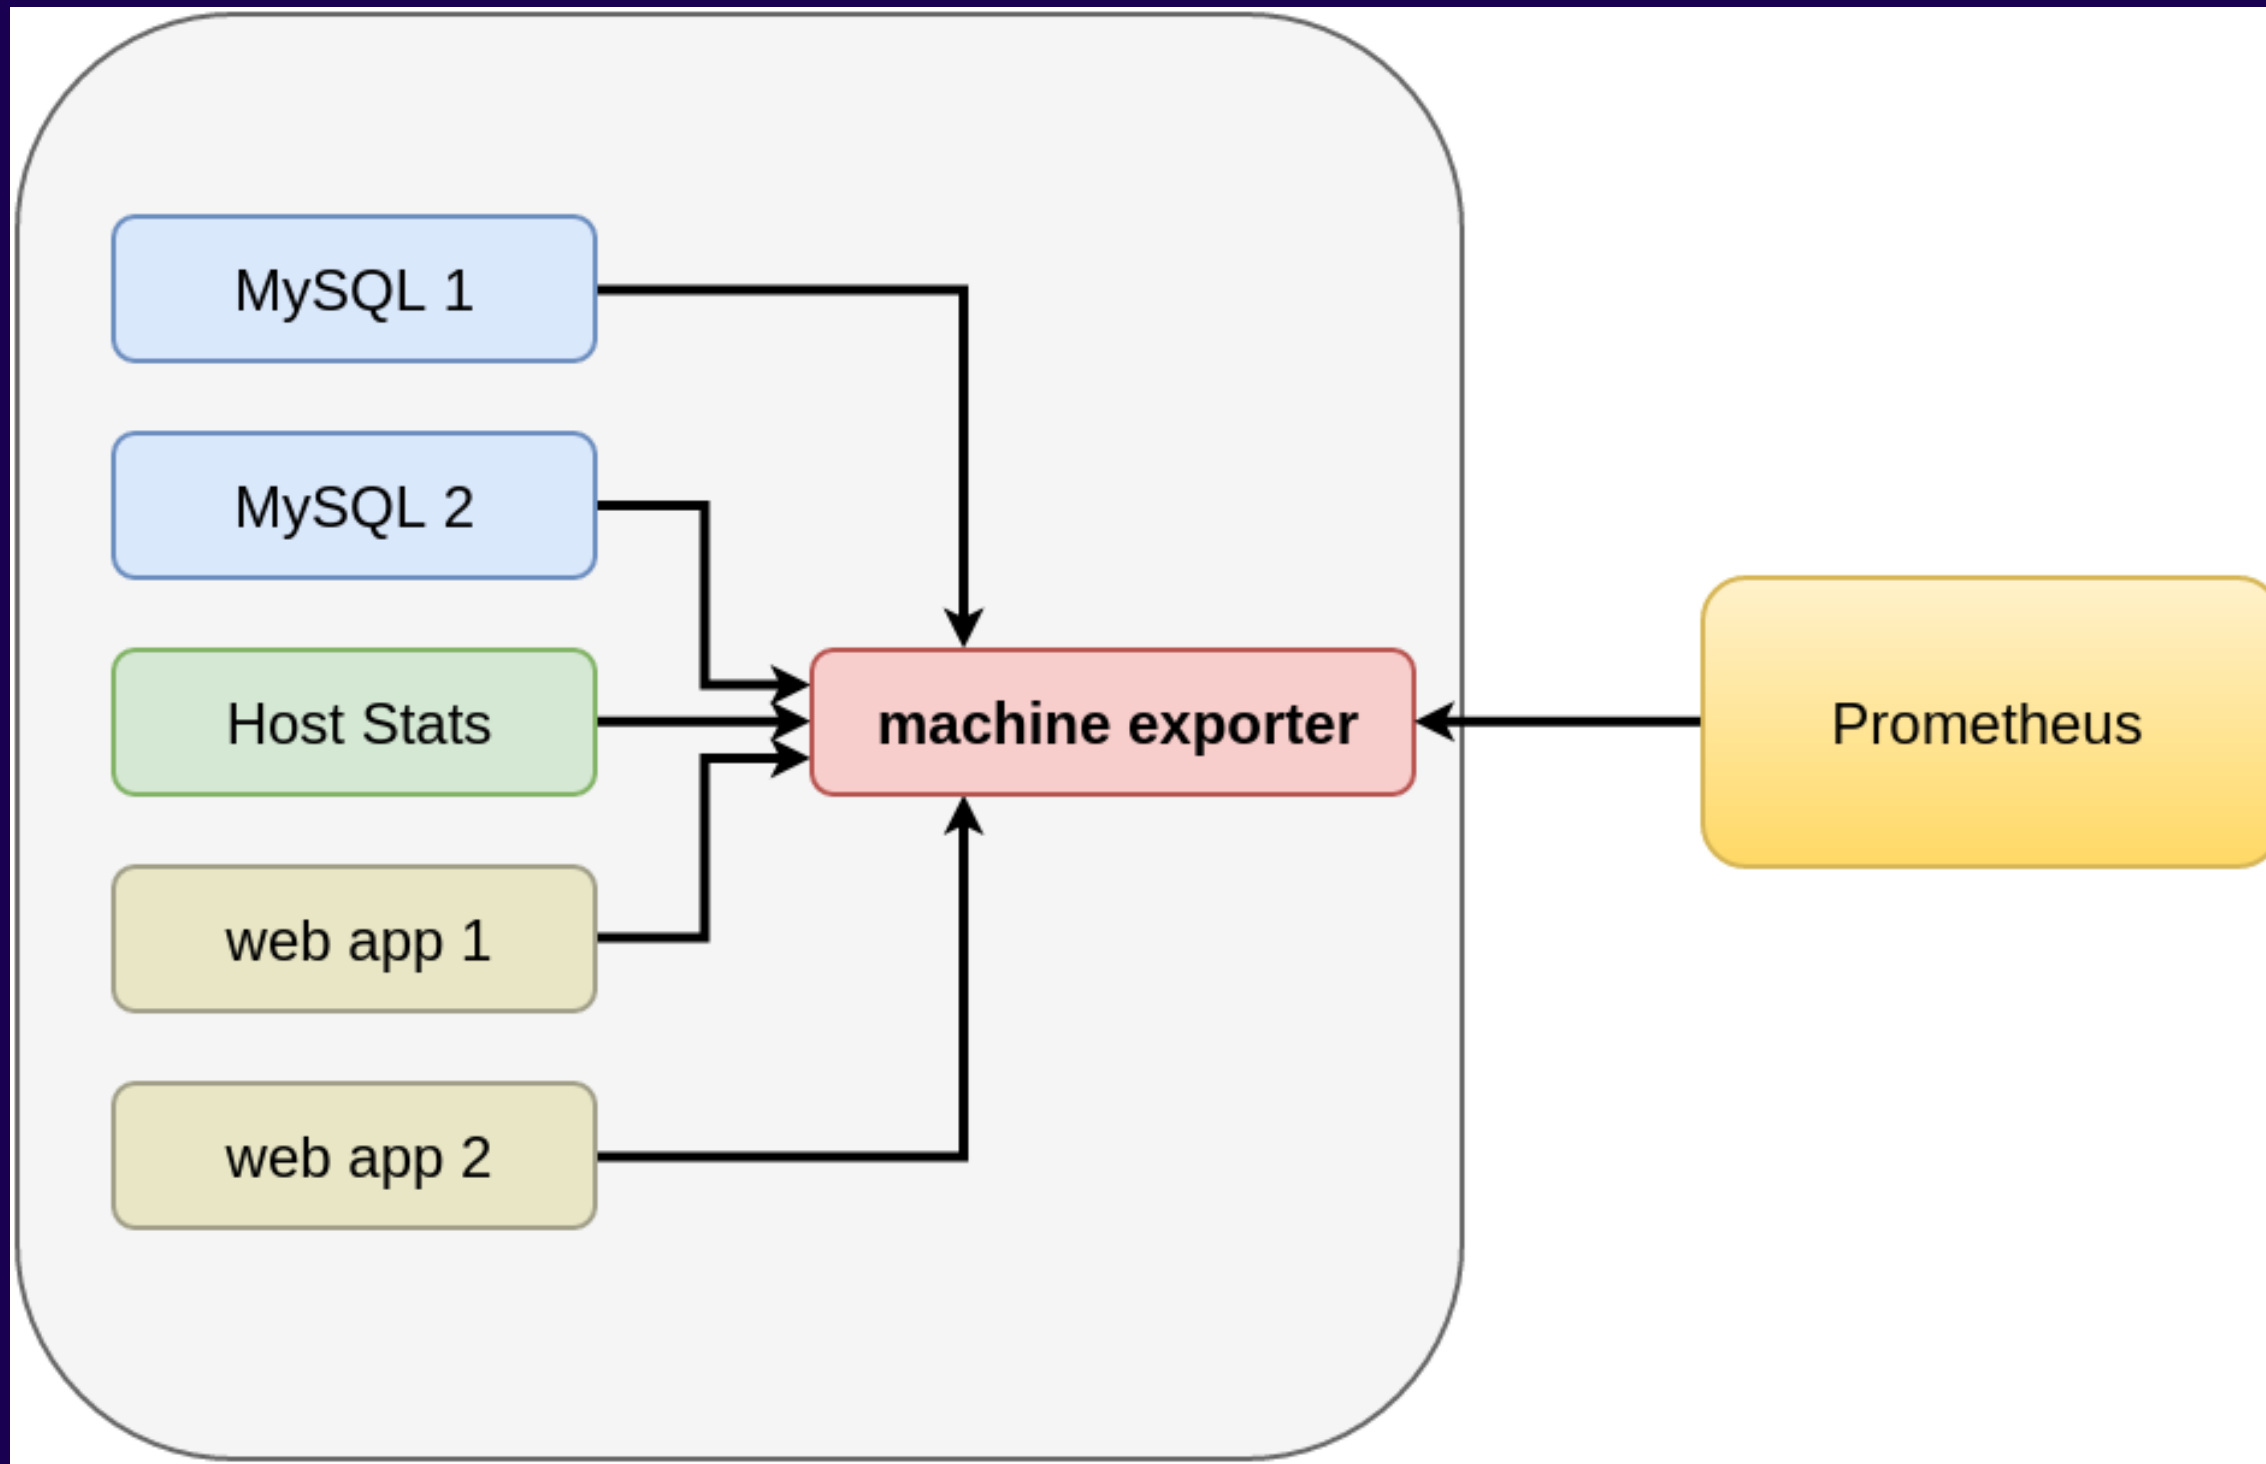
Category of Metrics

- USE
- RED
- AD-HOC

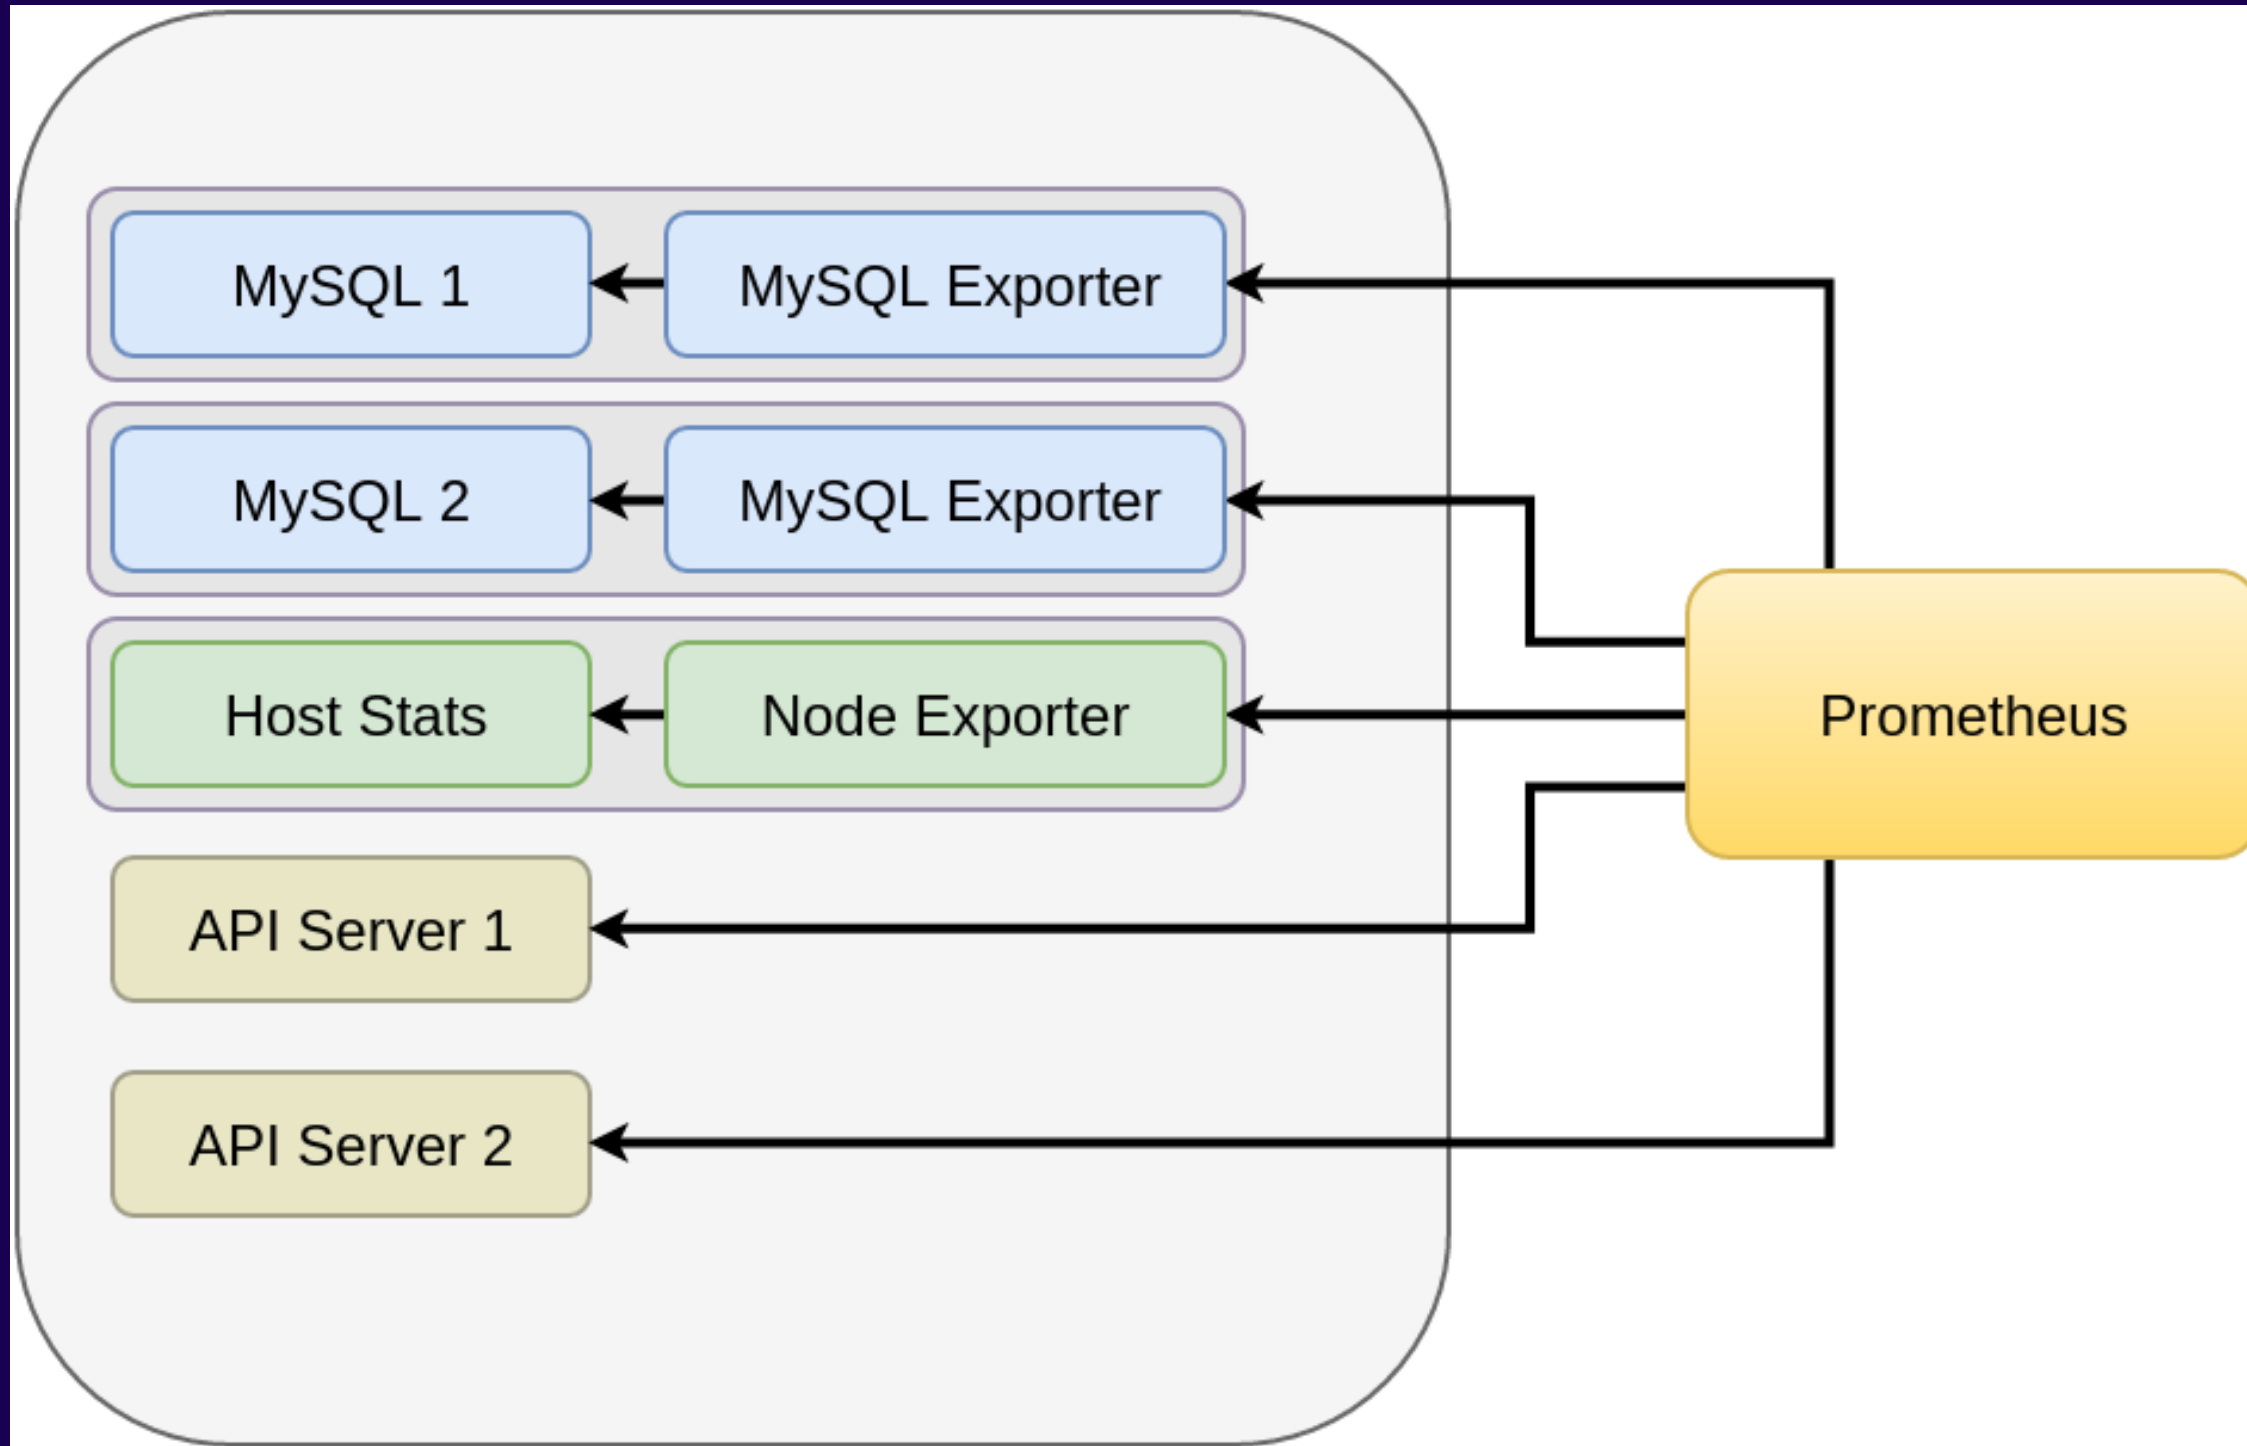
Type of Metrics

- Count
- Gauge
- Histogram

Exporters

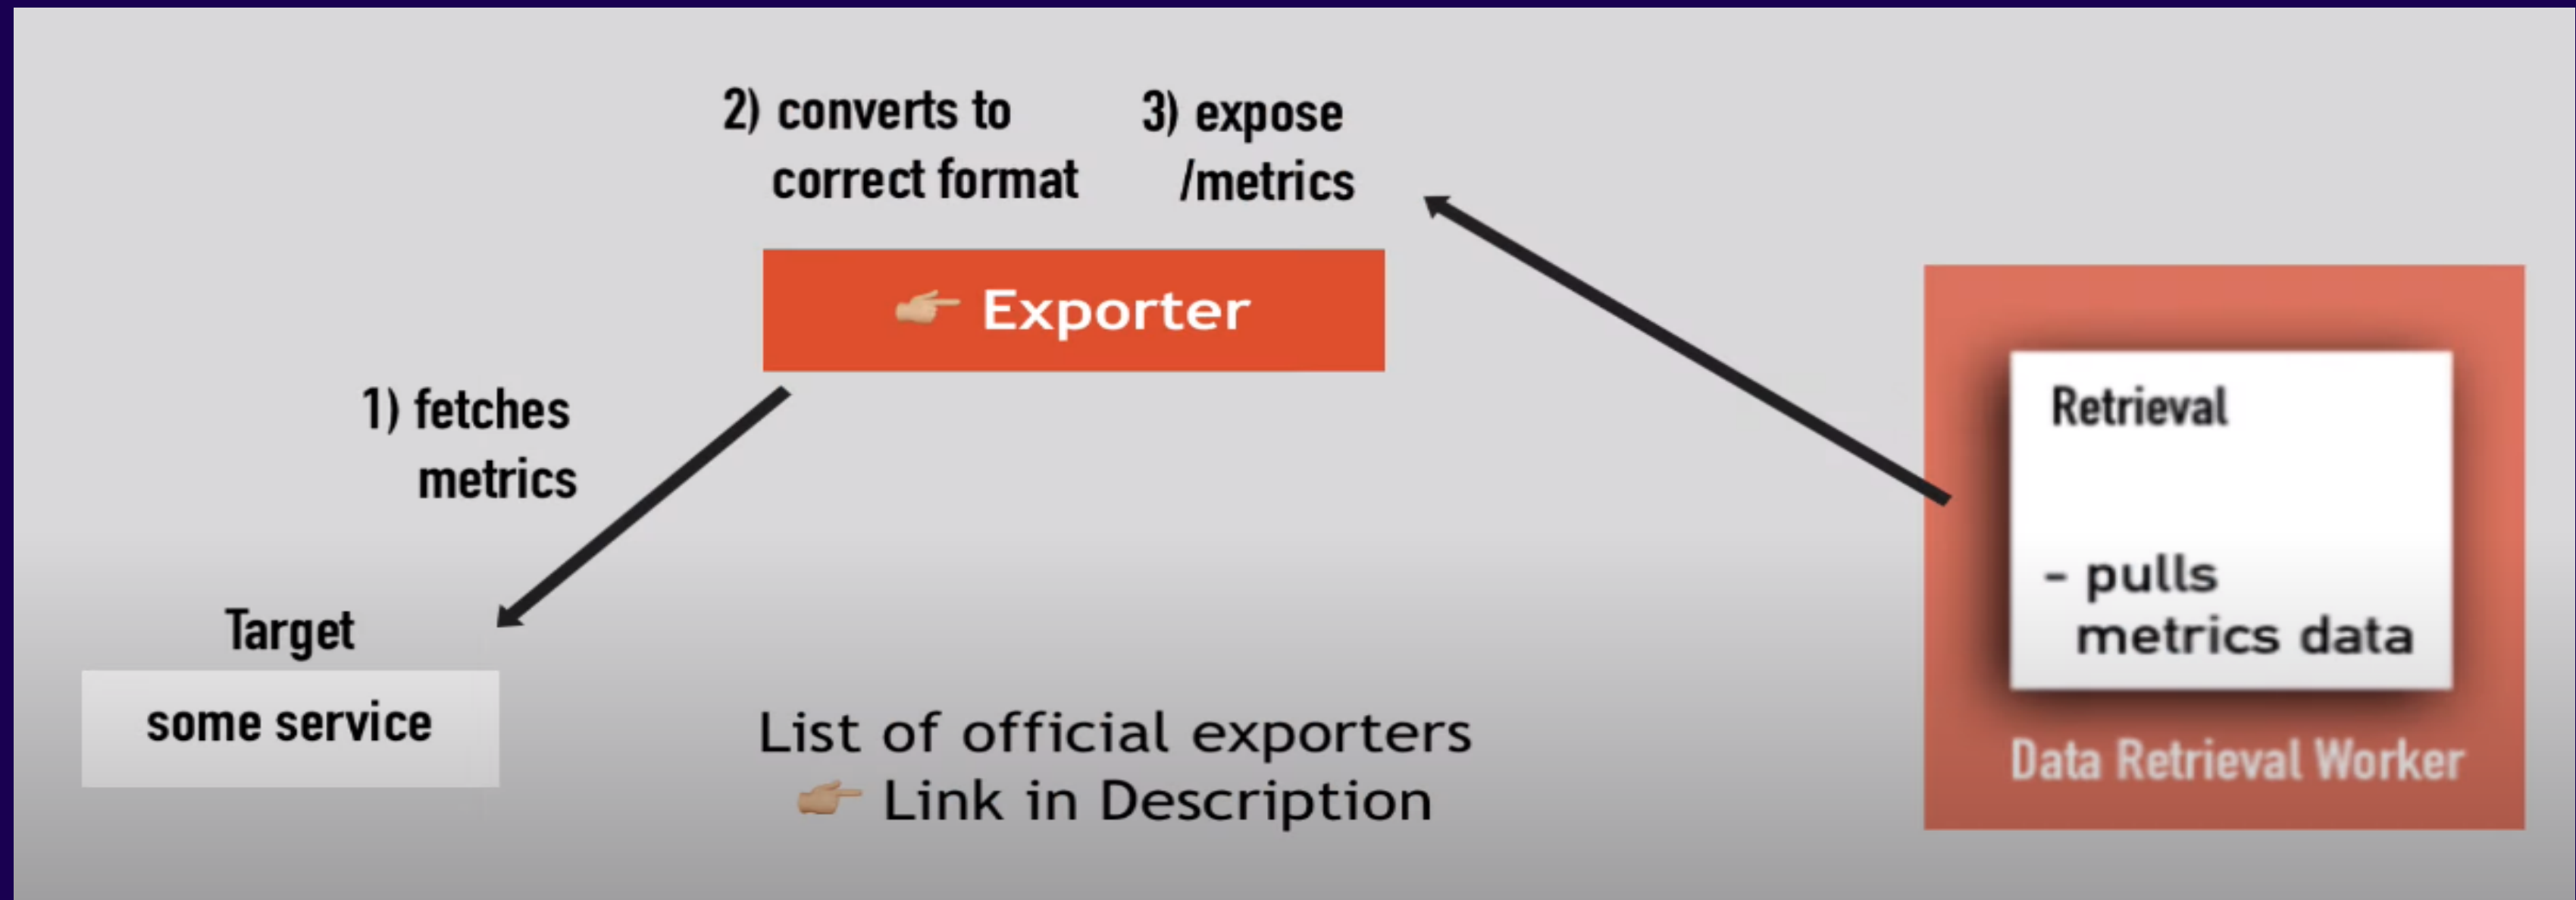


Exporters



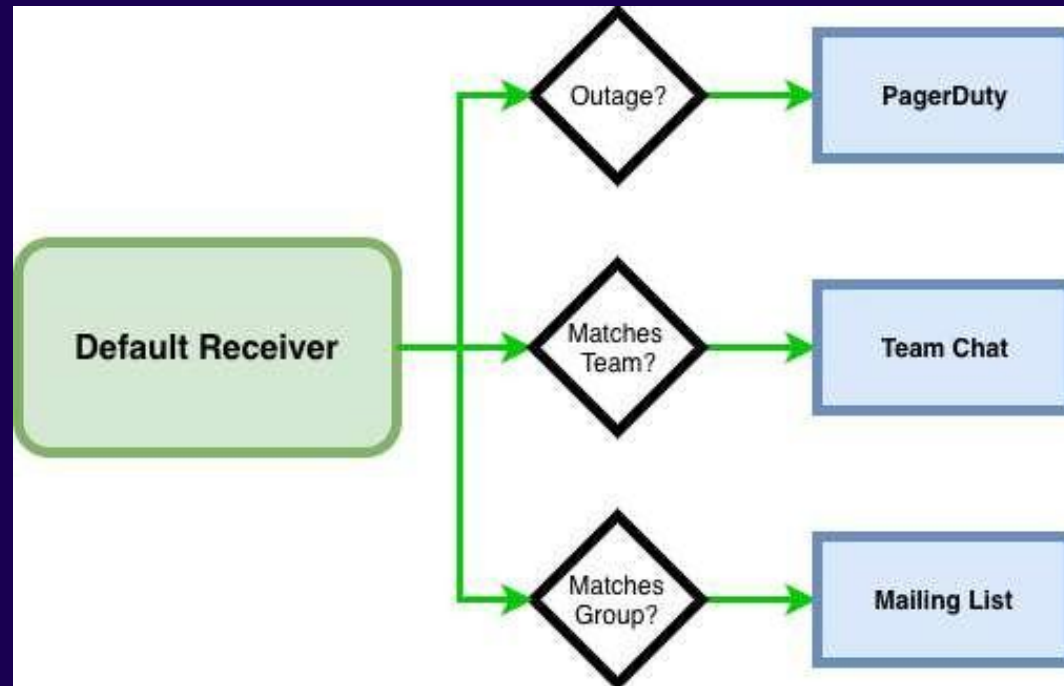
Exporters

A process that exposes Metrics for Prometheus to harvest



The available exporters can be find from here: <https://prometheus.io/docs/instrumenting/exporters/>

Alert Manager Rules



- AlertManager rules are conceptualized as routes, giving you the ability to write sophisticated sets of rules to determine where notifications should end up
- A default receiver should be configured for every notification, and then additional services can be configured through child routes which will match certain conditions

A full configuration reference is available here:

<https://prometheus.io/docs/alerting/configuration>

Alert Manager Rules

- Our config YAML file will be responsible for **setting up routing rules** that will determine how events are triaged
- As mentioned before, all events should **start with a default receiver**, called default-receiver, which will be the starting point for any route
- From there, any number of **sub-receivers can be configured**
- **Sample Configuration one called 'slack'** which will be invoked when the "service" tag of the event that has been triggered matches "activemq"
- Next, **configure our receivers**
- Sample Slack receiver config will **contain WebHook** into Slack

```
global:
  smtp_smarthost: 'localhost:25'
  smtp_from: 'alertmanager@monitoring.com'

route:
  receiver: 'default-receiver'
  group_wait: 30s
  group_interval: 5m
  repeat_interval: 4h
  group_by: [cluster, alertname]
  routes:
    - receiver: 'slack'
      group_wait: 10s
      match_re:
        service: activemq

receivers:
  - name: 'default-receiver'
    email_configs:
      - to: 'justin.reock@roguewave.com'

  - name: 'slack'
    slack_configs:
      - api_url: https://hooks.slack.com/services/
        channel: '#general'
```

Alert Manager Rules

Configure Sample Rules

- **two simple events**, but, events can be created out of a **huge range of possible query configurations**


```
groups:
- name: activemq
  rules:
  - alert: DLQ
    expr: org_apache_activemq_Broker_DLQ > 1
    for: 1m
    labels:
      severity: minor
      service: activemq
    annotations:
      summary: A message has gone into the DLQ
      dashboard: http://192.168.40.120:3000/dashboard/db/activemq-broker
      impact: A message has been misfired
      runbook: http://activemq.apache.org
  - alert: Broker Down
    expr: up{job="activemq"} == 0
    labels:
      severity: major
      service: activemq
    annotations:
      summary: The broker has crashed
      dashboard: http://192.168.40.120:3000/dashboard/db/activemq-broker
      impact: Broker is down
      runbook: http://activemq.apache.org
```

Integrating with Prometheus

15s

- **configure Prometheus to push** alert events into AlertManager
 - Add an alerting section to the Prometheus YAML file
- **Update prom- amq.yml** configuration file from earlier to integrate with our newly configured AlertManager instance
- Upon restarting Prometheus, we **should see our alerts** in the Prometheus dashboard