


# Seshagiri Sriram

## Terraform




# What is Terraform

Terraform is an open source infrastructure as code tool. Terraform was developed by HashiCorp company who is based in San Francisco, CA. We can also say that Terraform is a tool for building, changing, and versioning infrastructure safely and efficiently.

# Terraform principles

Terraform has these three principal simple steps:

- Init
- Plan
- Apply



# Terraform sample code

```
1 references
resource "outscale_keypair" "a_key_pair" {
  key_name = "terraform-key-pair-name"
}

1 references
resource "outscale_firewall_rules_set" "web" {
  group_name = "terraform_acceptance_test_example"
  group_description = "Used in the terraform presentation"
}

0 references
resource "outscale_vm" "basic" {
  image_id = "ami-8a6a0120"
  instance_type = "t2.micro"
  security_group = ["${outscale_firewall_rules_set.web.id}"]
  key_name = "${outscale_keypair.a_key_pair.key_name}"
}
```

# Multi-cloud

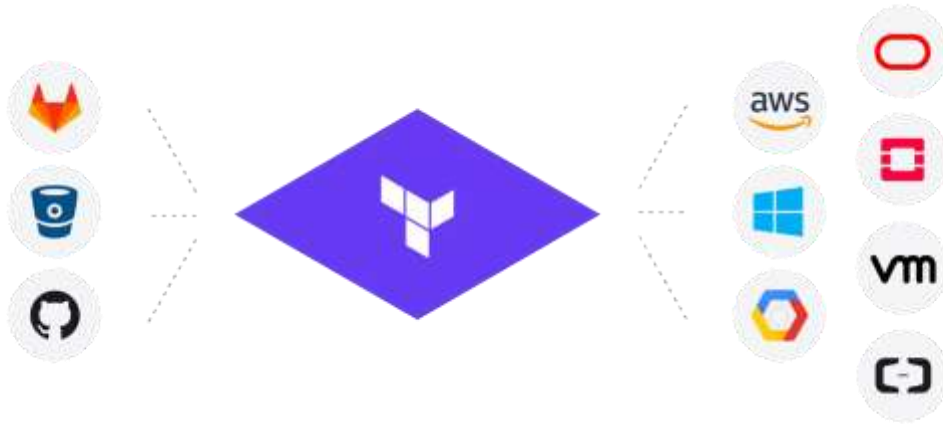
Terraform can be used to manage multi-cloud. It can combine multiple providers in a single workflow, which is a very nice feature.

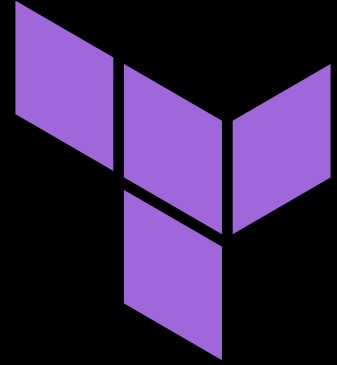
Terraform provides one consistent workflow for developers and operators to provision resources on any infrastructure provider. One workflow to learn increases user productivity, and also reduces organizational risk as that becomes one workflow to secure, one workflow to audit, and one workflow to govern.



# Terraform providers

- Right now, there are more than 100 providers and those individually can manage over a thousand resources. Providers are responsible to provide API interaction. You can find the list of providers on the terraform website.





**Code samples:**

[github.com/gruntwork-io/infrastructure-as-code-training](https://github.com/gruntwork-io/infrastructure-as-code-training)

# Getting started

Terraform is a tool for provisioning infrastructure

It supports many providers (cloud agnostic)

And many resources for each provider

**You define resources as code in Terraform templates**





## PROVIDERS

[Atlas](#)[AWS](#)[Azure \(Service Management\)](#)[Azure \(Resource Manager\)](#)[Chef](#)[CenturyLinkCloud](#)[CloudFlare](#)[CloudStack](#)[Consul](#)[Datadog](#)[DigitalOcean](#)[DNSMadeEasy](#)[DNSimple](#)[Docker](#)[Dyn](#)[Github](#)[Google Cloud](#)[Heroku](#)[InfluxDB](#)[Mailgun](#)[MySQL](#)[OpenStack](#)[Packet](#)[PostgreSQL](#)[Rackspace DNS](#)

## PROVIDERS

Terraform is used to create, manage, and manipulate infrastructure resources. Examples of resources include physical machines, VMs, network switches, containers, etc. Almost any infrastructure noun can be represented as a resource in Terraform.

Terraform is agnostic to the underlying platforms by supporting providers. A provider is responsible for understanding API interactions and exposing resources. Providers generally are an IaaS (e.g. AWS, DigitalOcean, GCE, OpenStack), PaaS (e.g. Heroku, CloudFoundry), or SaaS services (e.g. Atlas, DNSimple, CloudFlare).

Use the navigation to the left to read about the available providers.

[DOCUMENTATION HOME](#)[AWS PROVIDER](#)[EC2 RESOURCES](#)

[aws\\_ami](#)  
[aws\\_ami\\_copy](#)  
[aws\\_ami\\_from\\_instance](#)  
[aws\\_app\\_cookie\\_stickiness\\_policy](#)  
[aws\\_autoscaling\\_group](#)  
[aws\\_autoscaling\\_lifecycle\\_hook](#)  
[aws\\_autoscaling\\_notification](#)  
[aws\\_autoscaling\\_policy](#)  
[aws\\_autoscaling\\_schedule](#)  
[aws\\_ebs\\_volume](#)  
[aws\\_eip](#)  
[aws\\_elb](#)  
[aws\\_instance](#)  
[aws\\_key\\_pair](#)  
[aws\\_launch\\_configuration](#)  
[aws\\_lb\\_cookie\\_stickiness\\_policy](#)  
[aws\\_placement\\_group](#)  
[aws\\_proxy\\_protocol\\_policy](#)  
[aws\\_spot\\_instance\\_request](#)  
[aws\\_volume\\_attachment](#)

## AWS\_INSTANCE

Provides an EC2 Instance resource. This allows instances to be created, updated, and deleted. Instances also support [provisioning](#).

### Example Usage

```
# Create a new instance of the 'ami-408c7f28' (Ubuntu 14.04) on an
# t1.micro node with an AWS Tag naming it "HelloWorld"
provider "aws" {
  region = "us-east-1"
}


resource "aws_instance" "web" {
  ami = "ami-408c7f28"
  instance_type = "t1.micro"
  tags {
    Name = "HelloWorld"
  }
}
```

### Argument Reference

The following arguments are supported:

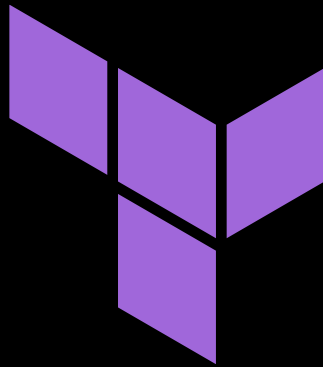
- ami** - (Required) The AMI to use for the instance.
- availability\_zone** - (Optional) The AZ to start the instance in.
- placement\_group** - (Optional) The Placement Group to start the instance in.
- tenancy** - (Optional) The tenancy of the instance (if the instance is

```
provider "aws" {  
  region = "us-east-1"  
}  
  
resource "aws_instance" "example" {  
  ami = "ami-408c7f28"  
  instance_type = "t2.micro"  
  tags { Name = "terraform-example" }  
}
```



This **template** creates a single EC2 instance in AWS

```
> terraform plan
+ aws_instance.example
    ami:                "" => "ami-408c7f28"
    instance_type:      "" => "t2.micro"
    key_name:           "" => "<computed>"
    private_ip:         "" => "<computed>"
    public_ip:          "" => "<computed>"
```



Plan: 1 to add, 0 to change, 0 to destroy.

**Use the `plan` command to see what you're about to deploy**

```
> terraform apply
```

```
aws_instance.example: Creating...
```

```
    ami:                "" => "ami-408c7f28"
```

```
    instance_type:      "" => "t2.micro"
```

```
    key_name:           "" => "<computed>"
```

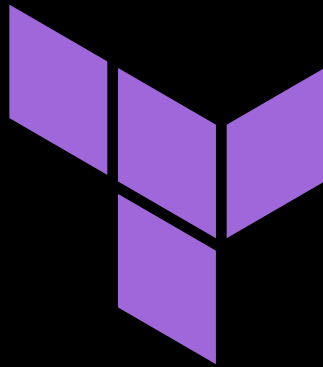
```
    private_ip:         "" => "<computed>"
```

```
    public_ip:          "" => "<computed>"
```

```
aws_instance.example: Creation complete
```

```
Apply complete! Resources: 1 added, 0 changed, 0  
destroyed.
```

**Use the `apply` command to apply  
the changes**



# Terraform

Templates can be parameterized using Variables

```
variable "name" {  
    description = "The name of the EC2 instance"  
}
```

Description, default and types are optional


# Terraform

```
variable "name" {  
    description = "The name of the EC2 instance"  
}  
resource "aws_instance" "example" {  
    ami = "ami-408c7f28"  
    instance_type = "t2.micro"  
    tags { Name = "${var.name}" }  
}
```

Terraform will prompt for a variable value when plan is done.  
Variables can also be passed with `-var <varname>=<varvalue>`

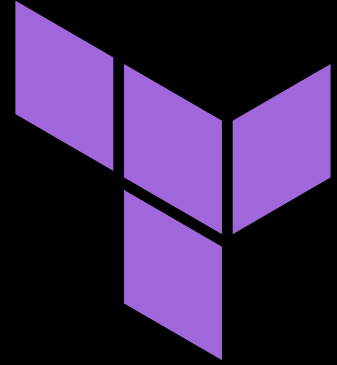
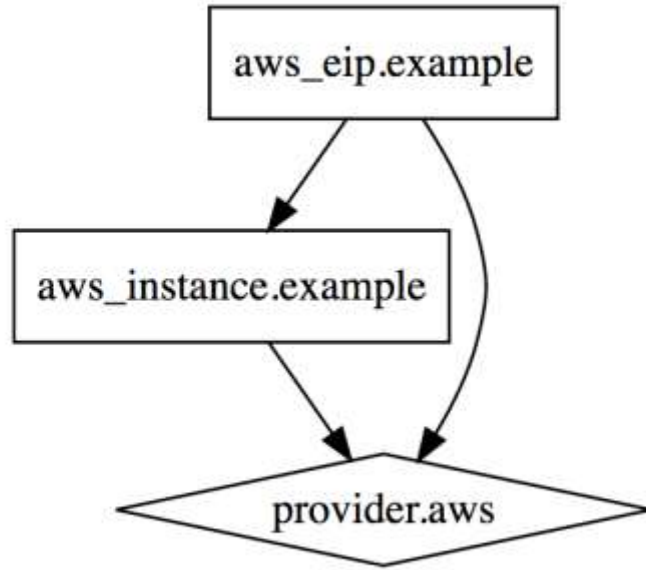
```
resource "aws_eip" "example" {  
    instance = "${aws_instance.example.id}"  
}
```

```
resource "aws_instance" "example" {  
    ami = "ami-408c7f28"  
    instance_type = "t2.micro"  
    tags { Name = "${var.name}" }  
}
```



**Notice the use of `${}` to **depend** on  
the id of the `aws_instance`**





**Terraform automatically builds a dependency graph**

```
> terraform destroy
```

```
aws_instance.example: Refreshing state... (ID: i-f3d58c70)
```

```
aws_elb.example: Refreshing state... (ID: example)
```

```
aws_elb.example: Destroying...
```

```
aws_elb.example: Destruction complete
```

```
aws_instance.example: Destroying...
```

```
aws_instance.example: Destruction complete
```

```
Apply complete! Resources: 0 added, 0 changed, 2 destroyed.
```

# Use the **destroy** command to clean up

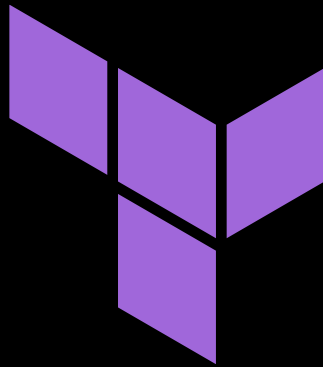
# Terraform

Terraform maintains state.

By default, stored in `.tfstate` files locally.

Preferred to store state remotely in S3, Atlas, Consul etc.

```
> terraform remote config \  
-backend=s3 \  
-backend-config=bucket=my-s3-bucket \  
-backend-config=key=terraform.tfstate \  
-backend-config=encrypt=true \  
-backend-config=region=us-east-1
```



You can enable **remote state storage** in S3, Atlas, Consul, etc.

# Terraform

Atlas is expensive

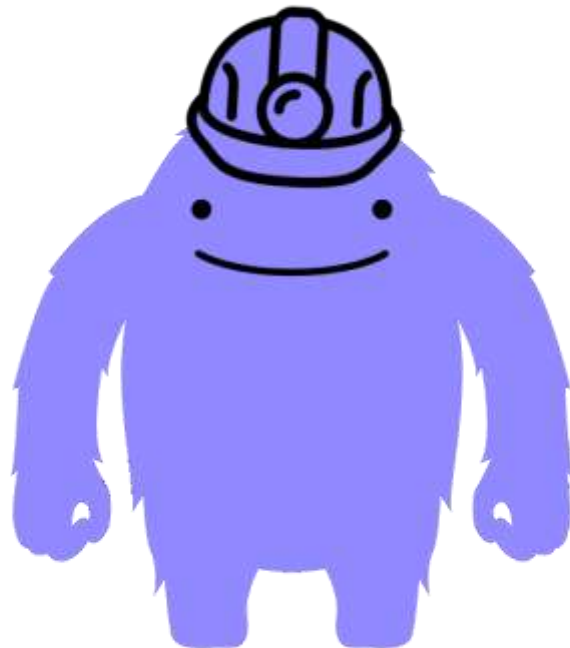
An alternative is Terragrunt

Terragrunt is an open source wrapper for Terraform

Uses Dynamodb and looks for a .terragrunt file

An example

```
terragrunt plan == terraform plan
```



# Terraform Modules

A module is similar to a blueprint

In plain words, it is a folder with terraform templates

A normal convention is to

<code>vars.tf</code>	Keep Variables and module inputs here
----------------------	---------------------------------------

<code>main.tf</code>	Main module
----------------------	-------------

<code>outputs.tf</code>	Keep outputs from application here
-------------------------	------------------------------------

# Terraform Modules

# Terraform Modules

Source specifies the folder where the module is located.

Can be re-used any number of times

It can even refer to a versioned GIT URL

```
module "example_rails_app" {  
    source = "./rails-module"  
    name = "Example Rails App"    # module variables  
    ami = "ami-123456"           # Module variables  
}
```



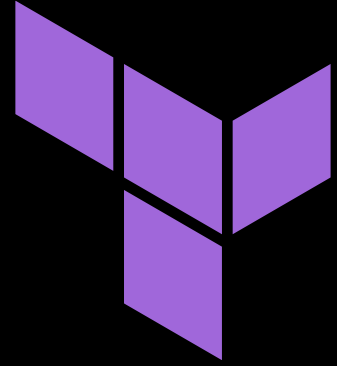
# Terraform Modules

When using GIT modules, run the ***terraform get -update*** command.


# Terraform Best Practices

01. Plan before apply

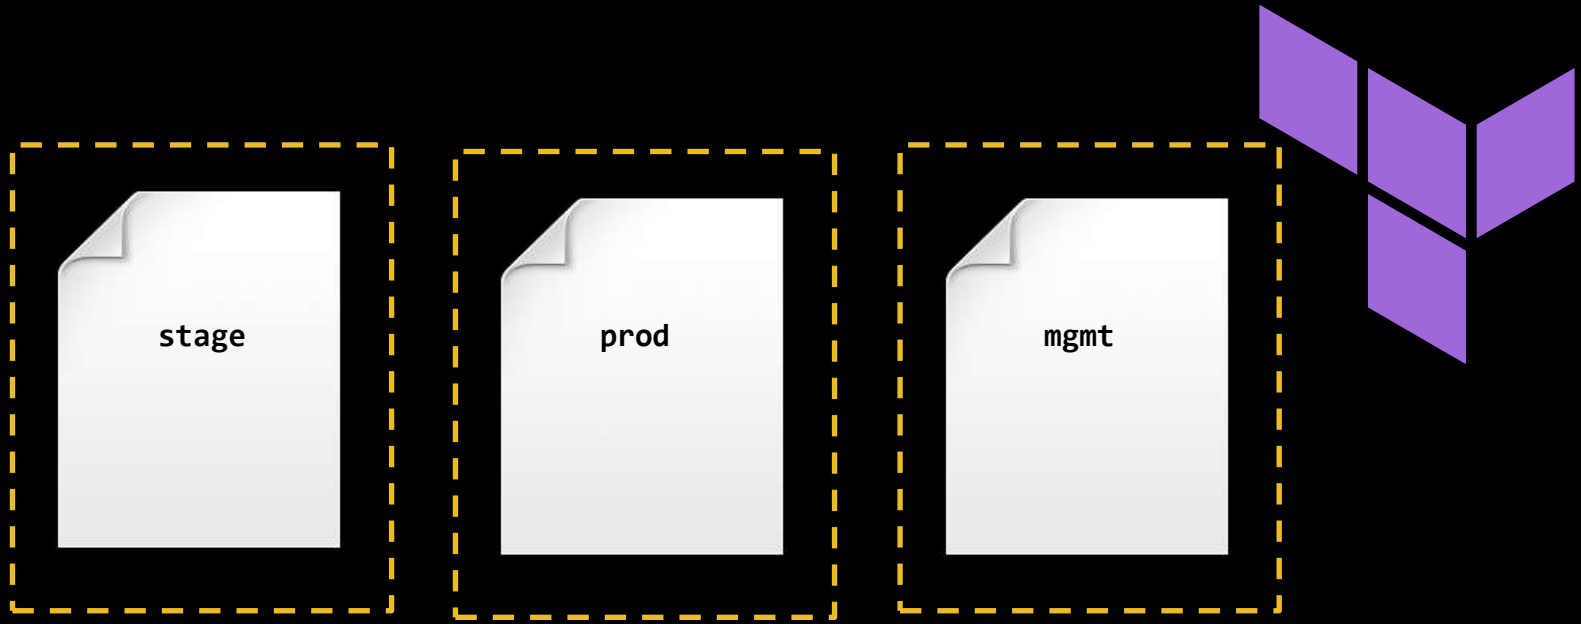
02. Stage before moving to Production



### **3. Isolated environments**



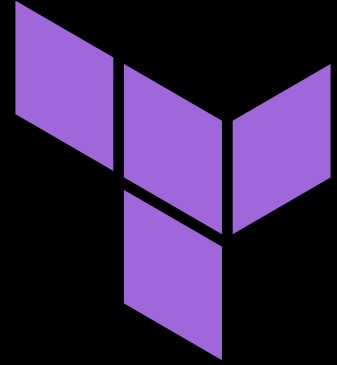
**It's tempting to define everything  
in 1 template**



What you really want is **isolation**  
for each environment



**That way, a problem in stage  
doesn't affect prod**



# Recommended folder structure (simplified):

global (Global resources such as IAM, SNS, S3)

- L `main.tf`
- L `.terragrunt`

stage (Non-production workloads, testing)

- L `main.tf`
- L `.terragrunt`

prod (Production workloads, user-facing apps)

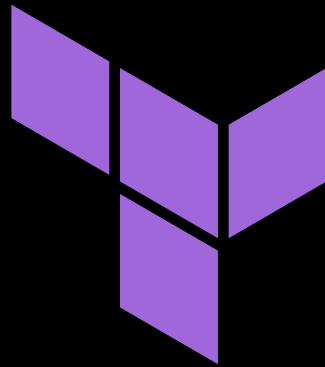
- L `main.tf`
- L `.terragrunt`

mgmt (DevOps tooling such as Jenkins, Bastion

Host)

**Each folder gets its own `.tfstate`**

- L `main.tf`
- L `.terragrunt`





global (Global resources such as IAM, SNS, S3)

- └─ main.tf
- └─ .terragrunt

stage (Non-production workloads, testing)

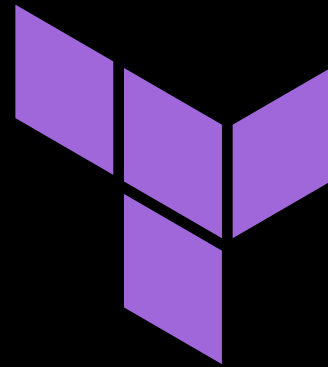
- └─ main.tf
- └─ .terragrunt

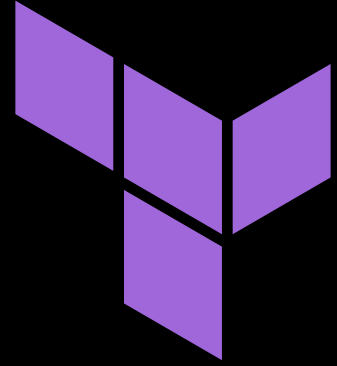
prod (Production workloads, user-facing apps)

- └─ main.tf
- └─ .terragrunt

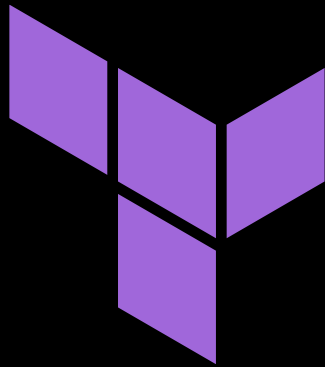
mgmt (DevOps tooling such as Jenkins, Bastion  
Host)

**Use terraform\_remote\_state to share  
state between them**

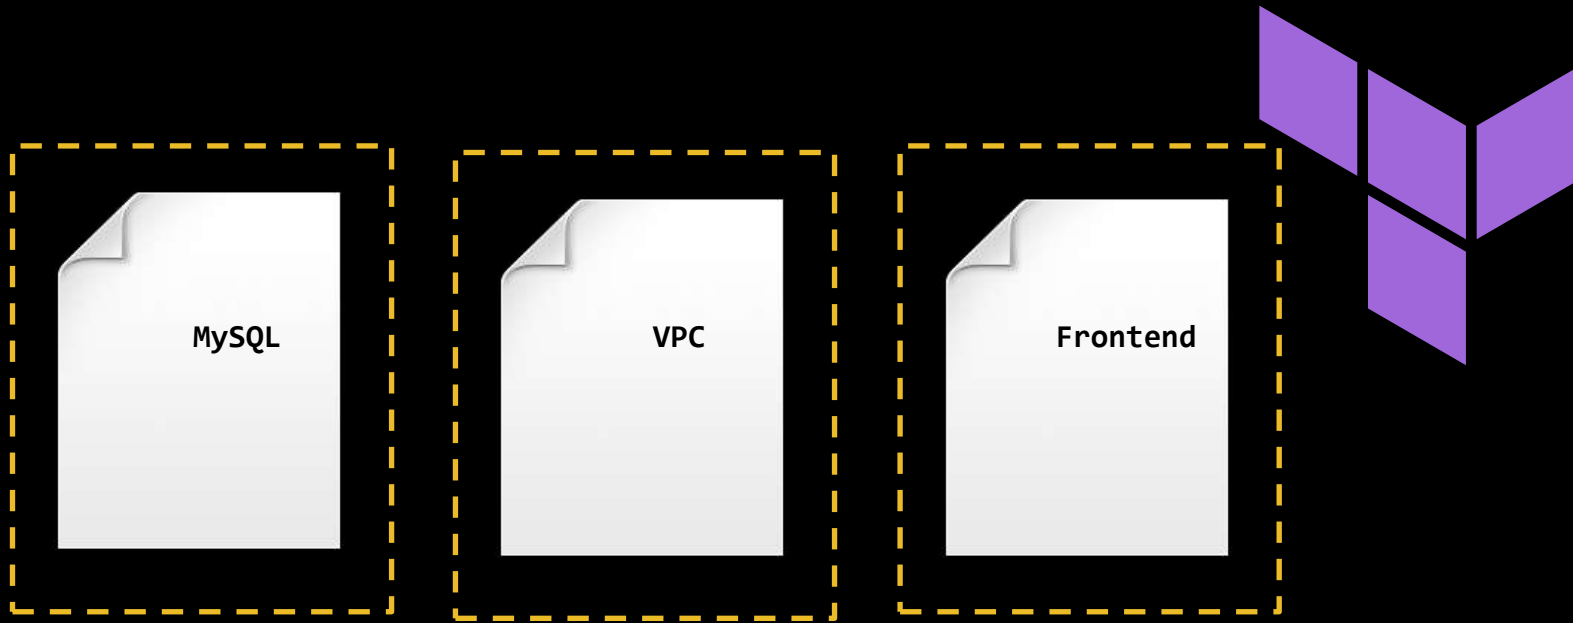




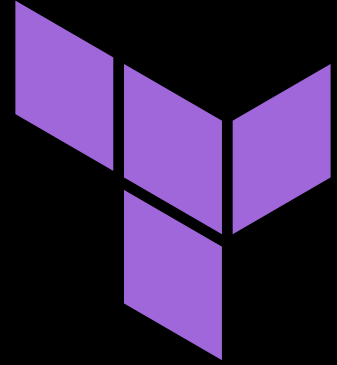
## 4. Isolated components



**It's tempting to define everything  
in 1 template for all components.**



What you really want is **isolation**  
for each component



# Recommended folder structure (full):

global (Global resources such as IAM, SNS, S3)

- L iam

- L sns

stage (Non-production workloads, testing)

- L vpc

- L mysql

- L frontend

prod (Production workloads, user-facing apps)

- L vpc

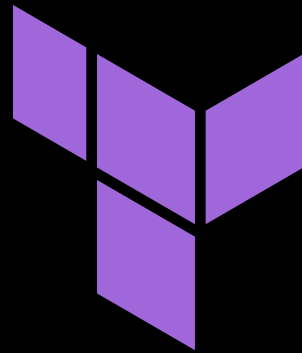
- L mysql

- L frontend

**Each component in each environment  
gets its own .tfstate**

mgmt (DevOps tooling such as Jenkins, Bastion)

- L vpc



global (Global resources such as IAM, SNS, S3)

- └ iam

- └ sns

stage (Non-production workloads, testing)

- └ vpc

- └ mysql

- └ frontend

prod (Production workloads, user-facing apps)

- └ vpc

- └ mysql


- └ frontend

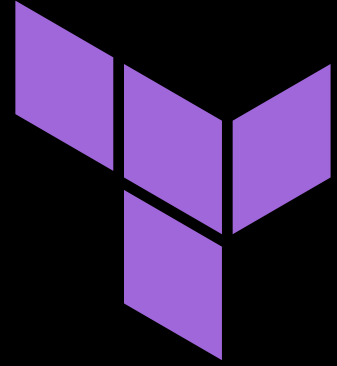
**Use terraform\_remote\_state to share state between them**

mgmt (DevOps tooling such as Jenkins, Bastion)

hosted

- └ vpc

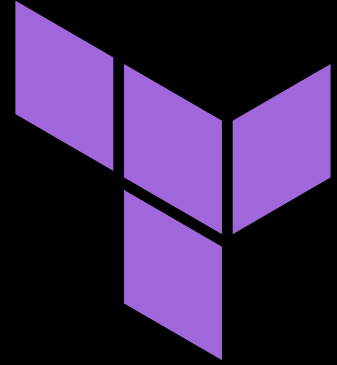




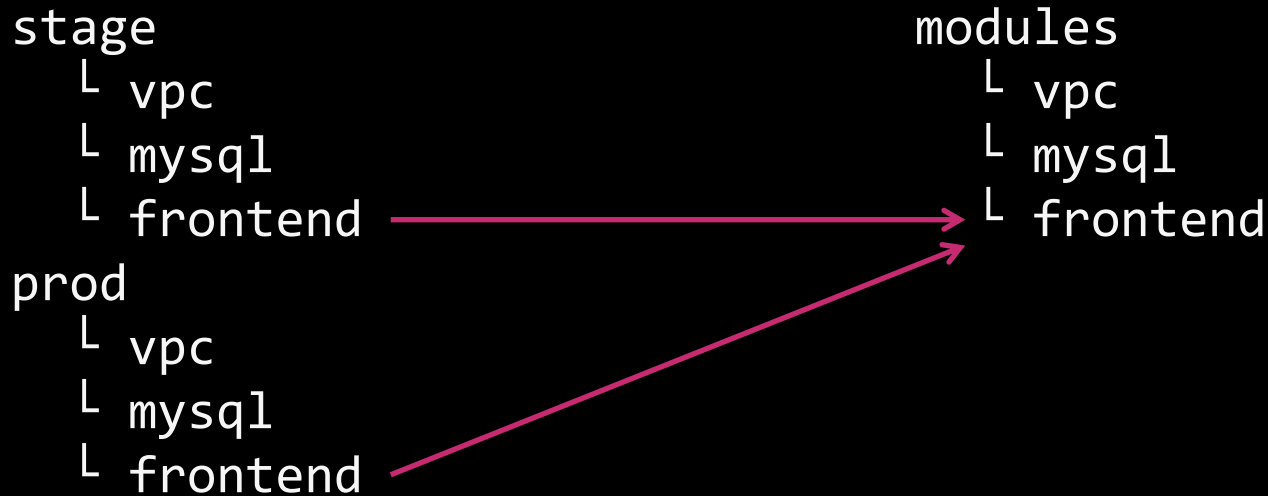
## 5. Use modules



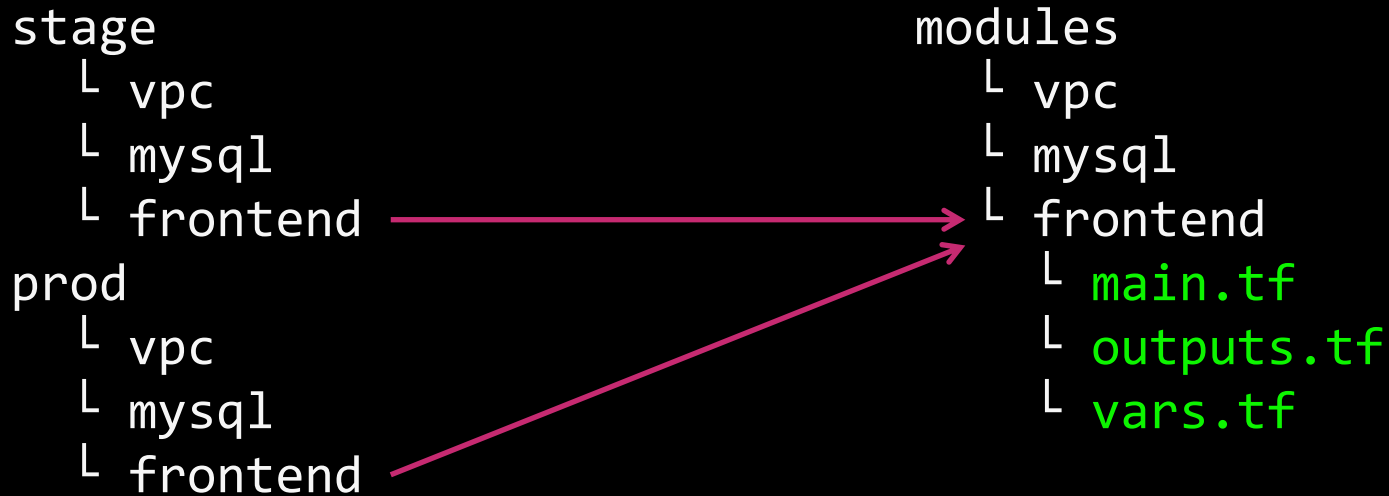
```
stage
├─ vpc
├─ mysql
├─ frontend
prod
├─ vpc
├─ mysql
├─ frontend
```



How do you **avoid copy/pasting** code  
between stage and prod?



**Define reusable modules!**



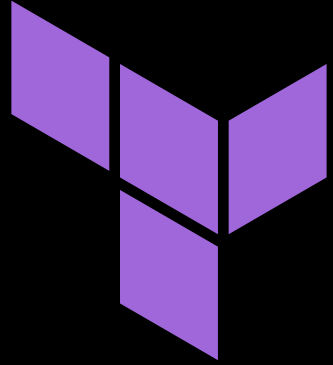
Each module defines one **reusable component**

```
variable "name" {  
    description = "The name of the EC2  
instance"  
}
```

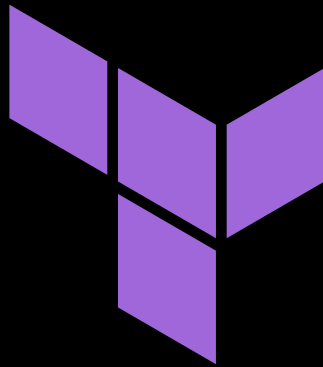
```
variable "ami" {  
    description = "The AMI to run on the EC2  
instance"  
}
```

```
variable "memory" {  
    description = "The amount of memory to  
allocate"  
}
```

**Define inputs in vars.tf to  
configure the module**

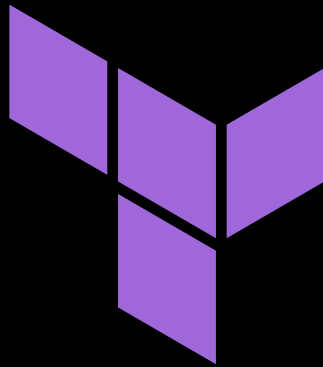


```
module "frontend" {  
    source = "../modules/frontend"  
  
    name = "frontend-stage"  
    ami = "ami-123asd1"  
    memory = 512  
}
```

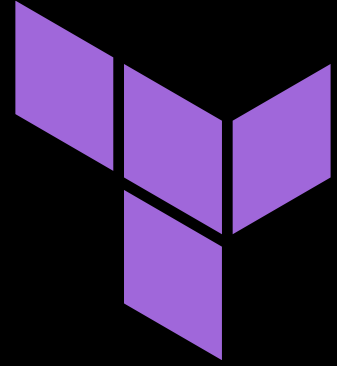


**Use the module in stage**  
**(stage/frontend/main.tf)**

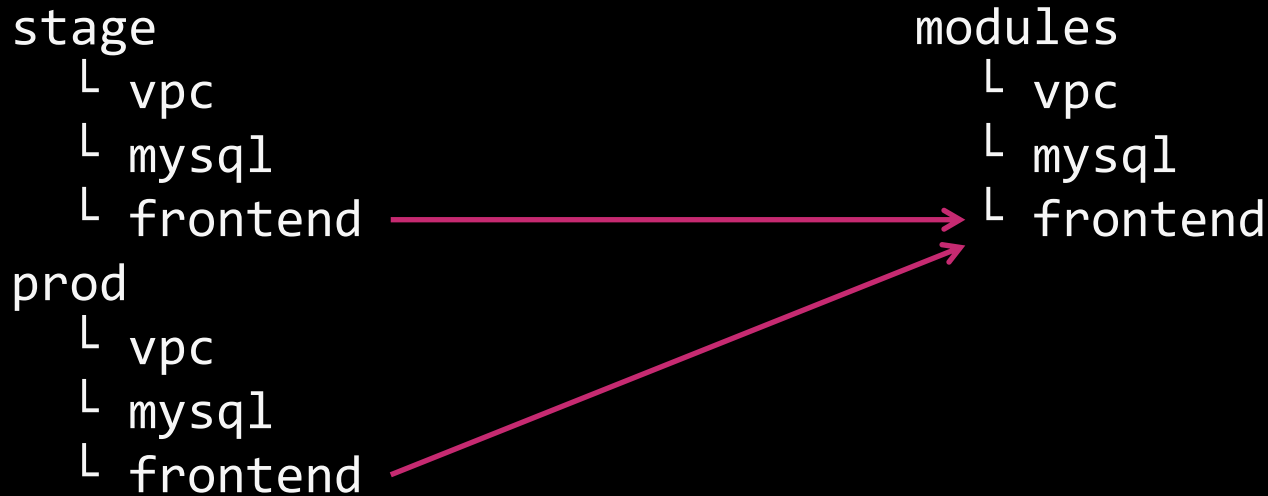
```
module "frontend" {  
    source = "../modules/frontend"  
  
    name = "frontend-prod"  
    ami = "ami-123abcd"  
    memory = 2048  
}
```



**And in prod**  
(prod/frontend/main.tf)

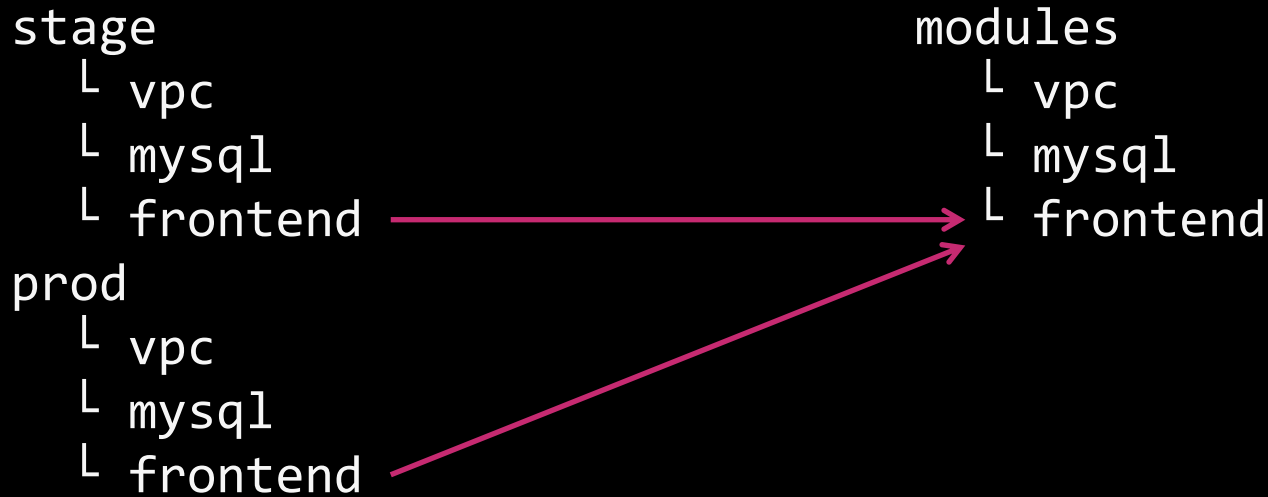


## **6. Use versioned modules**



If stage and prod point to the same folder, you lose **isolation**





**Any change in modules/frontend affects both stage and prod**

infrastructure-live

- └ stage
  - └ vpc
  - └ mysql
  - └ frontend
- └ prod
  - └ vpc
  - └ mysql
  - └ frontend

infrastructure-modules

- └ vpc
- └ mysql
- └ frontend

**Solution: define modules in a  
separate repository**


infrastructure-live

- └ stage
  - └ vpc
  - └ mysql
  - └ frontend

- └ prod
  - └ vpc
  - └ mysql
  - └ frontend

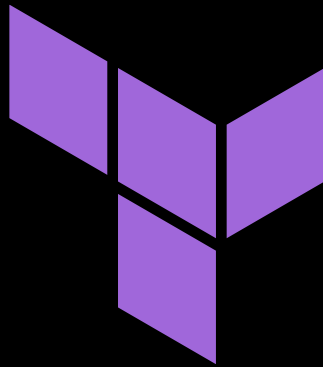
infrastructure-modules

- └ vpc
- └ mysql
- └ frontend

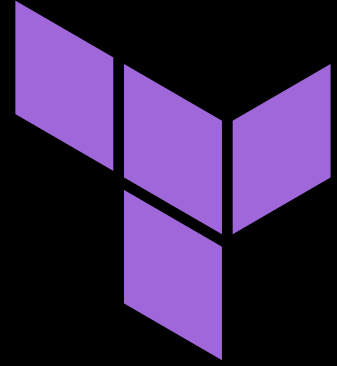


Now stage and prod can use  
different **versioned URLs**

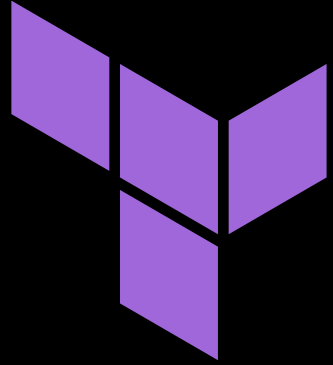
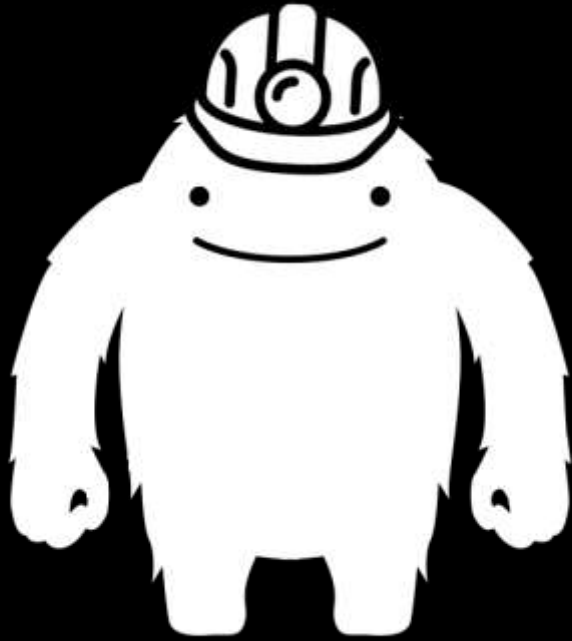
```
module "frontend" {  
    source =  
    "git::git@github.com:foo/infrastructure-  
modules.git//frontend?ref=0.2"  
  
    name = "frontend-prod"  
    ami = "ami-123abcd"  
    memory = 2048  
}
```



**Example Terraform code**  
(prod/frontend/main.tf)



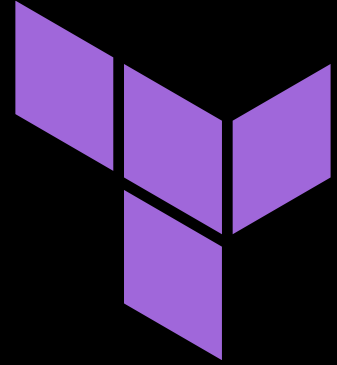
## 7. State file storage



Use **terragrunt**

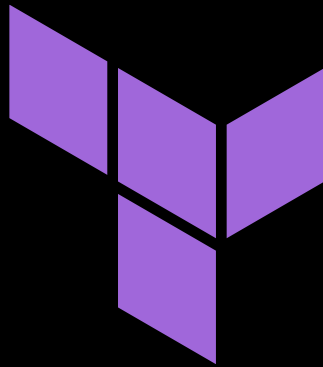
[github.com/gruntwork-io/terragrunt](https://github.com/gruntwork-io/terragrunt)

```
dynamoDbLock = {  
    stateFileId = "mgmt/bastion-host"  
}
```



**Use a custom lock (`stateFileId`) for each set of templates**

```
remoteState = {  
  backend = "s3"  
  backendConfigs = {  
    bucket = "acme-co-terraform-state"  
    key = "mgmt/bastion-  
host/terraform.tfstate"  
    encrypt = "true"  
  }  
}
```



Use an **S3 bucket** with **encryption** for  
remote state storage

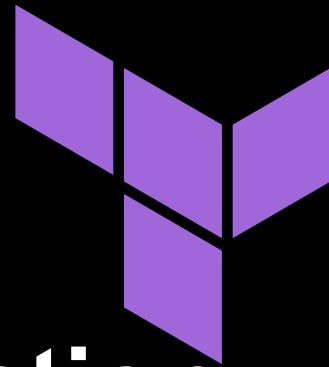


# Terraform Loops

Terraform is Declarative. Very little logic is possible.

Count can be used limitedly

```
resource "aws_instance" "example" {  
    count = 3  
    ami = "${var.ami}"  
    instance_type = "t2.micro"  
    tags { Name = "${var.name}-${count.index}" }  
}
```



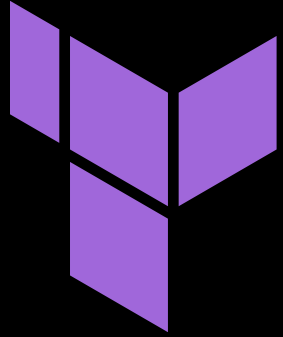
# Do even more with interpolation functions:

[terraform.io/docs/configuration/interpolation.html](https://terraform.io/docs/configuration/interpolation.html)

```
resource "aws_instance" "example" {  
    count = 3  
    ami = "${element(var.amis, count.index)}"  
    instance_type = "t2.micro"  
    tags { Name = "${var.name}-${count.index}" }  
}
```

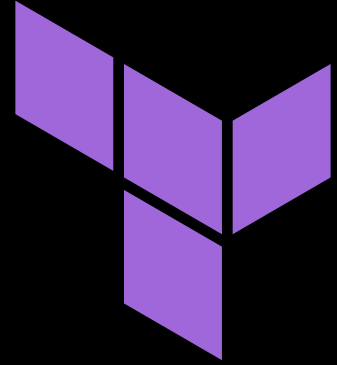
```
variable "amis" {  
    type = "list"  
    default = ["ami-abc123", "ami-abc456", "ami-  
abc789"]  
}
```

**Create three EC2 Instances, each with a  
different AMI**



```
output "all_instance_ids" {  
    value = ["${aws_instance.example.*.id}"]  
}
```

```
output "first_instance_id" {  
    value = "${aws_instance.example.0.id}"  
}
```



**Note: resources with count are actually **lists** of resources!**

# Terraform if-else

Terraform is Declarative. Very little logic is possible.

Count can be used limitedly for same purpose

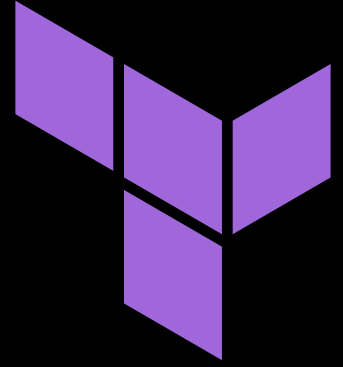
```
resource "aws_instance" "example" {  
    count = "${var.should_create_instance}"  
    ami = "ami-abcd1234"  
    instance_type = "t2.micro"  
    tags { Name = "${var.name}" }  
}  
  
variable "should_create_instance" {    default = true }
```

# Advantages of Terraform

- 1. Define infrastructure-as-code**
- 2. Concise, readable syntax**
- 3. Reuse: inputs, outputs, modules**
- 4. Plan command!**
- 5. Cloud agnostic**
- 6. Very active development**

# Disadvantages of Terraform

- 1. Maturity. You will hit bugs.**
- 2. Collaboration on Terraform state is tricky (but not with terragrunt)**
- 3. No rollback**
- 4. Poor secrets management**



**Thank You**