



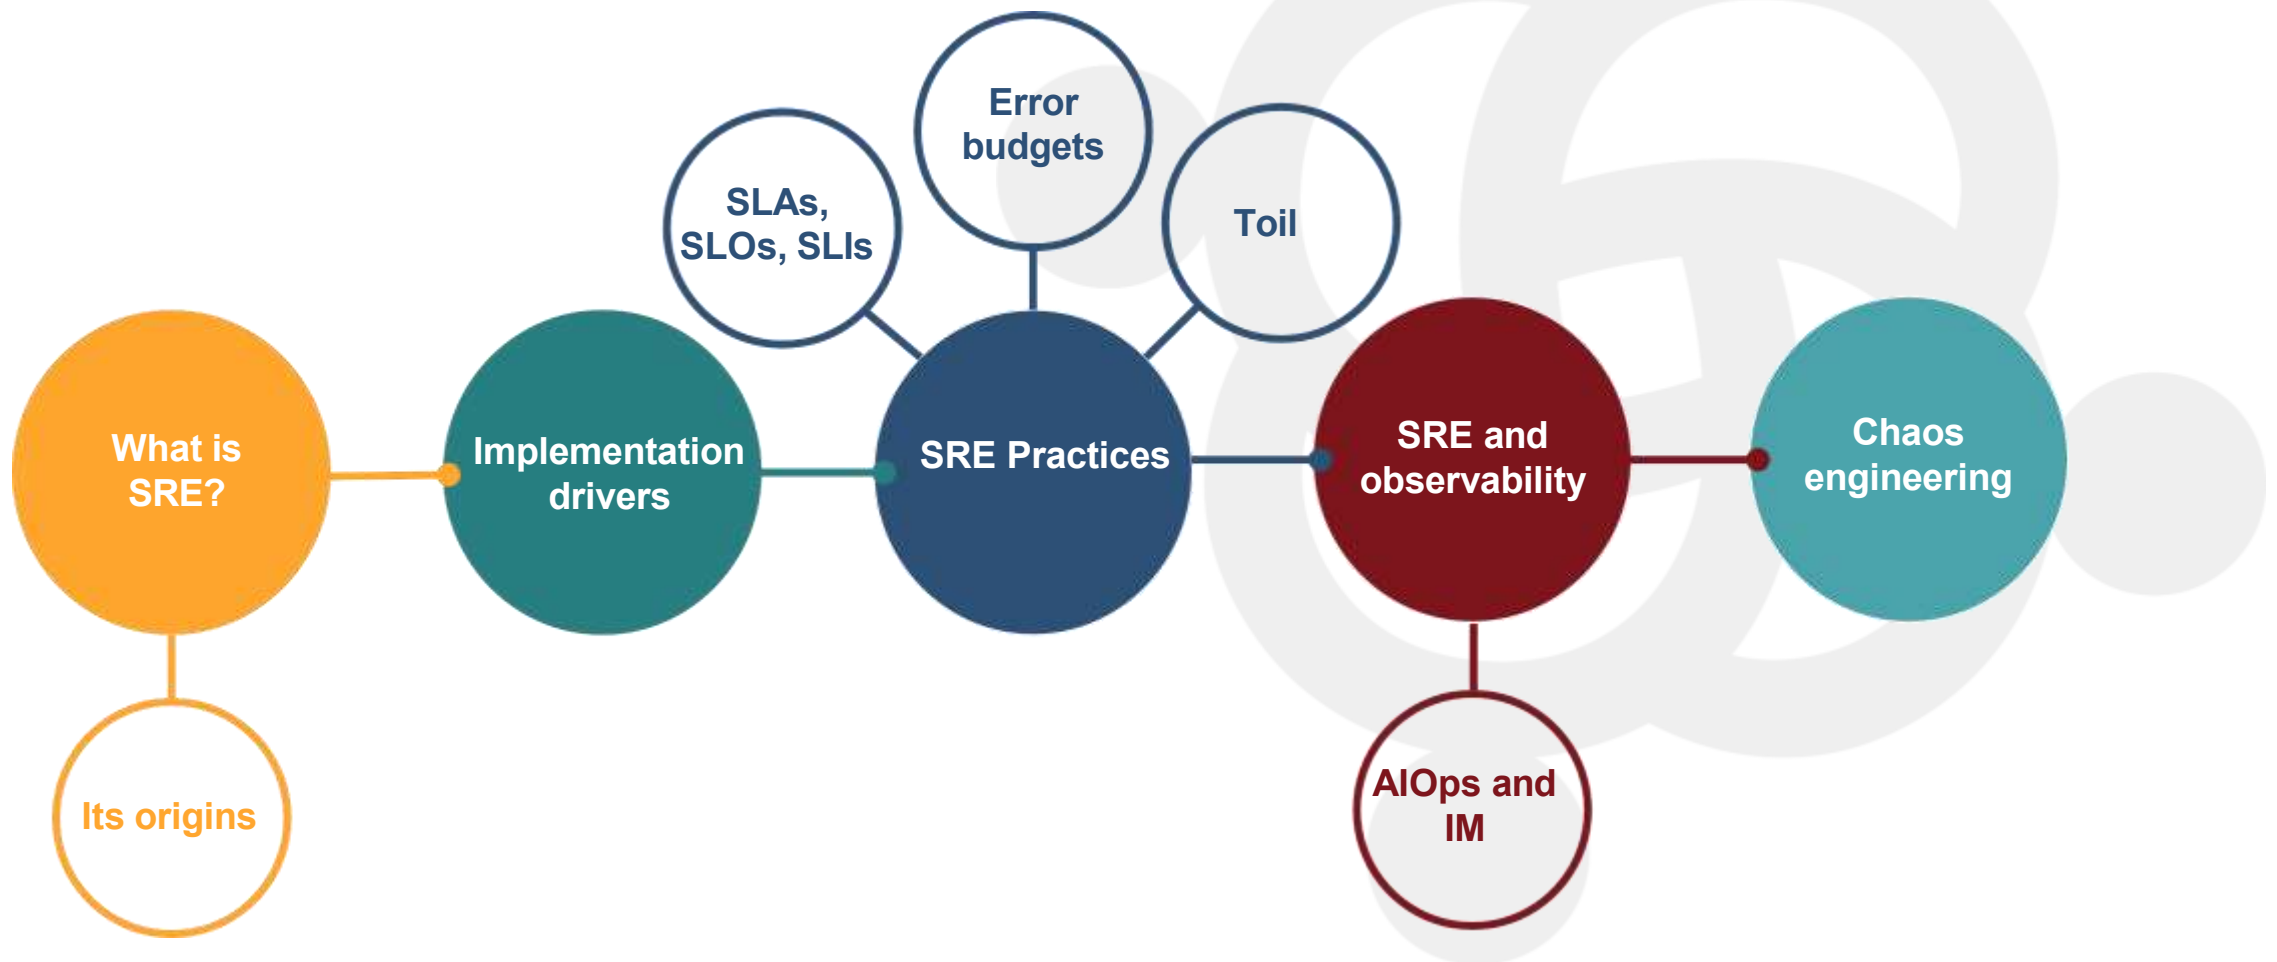
DevOps Institute
ADVANCING THE HUMANS OF DEVOPS




SRE and Incident Management



Flow: Talk map



What is SRE?



- *The goal is to create ultra-scalable and highly reliable distributed software systems*
- *SRE's spend **50% of their time doing "ops" related work** such as issue resolution, on-call, and manual interventions*
- *SRE's spend **50% of their time on development tasks** such as new features, scaling or automation*
- *Observability, monitoring, alerting and automation are a large part of SRE*




Time Spent On Dev Work (Versus Ops Work) and On Call

- % of work on dev
- % of work on call





Where has the concept come from?



- *Site Reliability Engineering (SRE) is a discipline that incorporates aspects of software engineering and applies them to infrastructure and operations problems*
- *Created at Google around 2003 and publicized via SRE books*

“What happens when a software engineer is tasked with what used to be called operations.”

Ben Treynor, Google



class SRE implements DevOps

▶ DevOps

is a set of practices, guidelines and culture designed to break down silos in IT development, operations, architecture, networking and security.

▶ Site Reliability Engineering

is a set of practices we've found to work, some beliefs that animate those practices, and a job role.



Successful SRE Implementation Drivers

- 60%** | How quickly we resolve incidents
- 43%** | The amount of time between failures
- 41%** | How quickly we do root cause incident analysis
- 40%** | How quickly we push product updates
- 33%** | How quickly our business can expand to new markets
- 22%** | How quickly we can understand the cause of social media sentiment





Toil and the wisdom of production



- *Any manual, mandated operational task is bad*
- *If a task can be automated, then it should be automated*
- *Tasks can provide the "wisdom of production" that will inform better system design and behavior*

SREs must have time to make tomorrow better than today




Causes Of Toil

Minor Moderate Major N/A





Moving forward to SRE at Slack




- *Slack moved from 100 AWS instances to 15,000 instances over 4 years*
- *Excessive toil caused by low-quality, noisy alerting*
- *Ops teams were so consumed by interrupt-driven toil that they were unable to make progress on improving reliability*
- *Slack **explicitly committed to the importance of reliability over feature velocity***
- *Operational ownership of services pushed back into the dev teams resulting in the teams making the code fixes necessary to stop the incident alerts*



SRE Activity Breakdown

Responding to incidents or outages



Post-mortem analysis and/or write-ups



Participating in on-call rotation



Developing applications or capabilities



Experimenting or receiving training to expand knowledge or skills



Authoring business processes, rules, or best practices



Performing audits of usage/cost allocation



Spinning up new hosts/instances



Planning release roadmaps



Performing chaos engineering exercises



Providing trainings on third-party platform capabilities



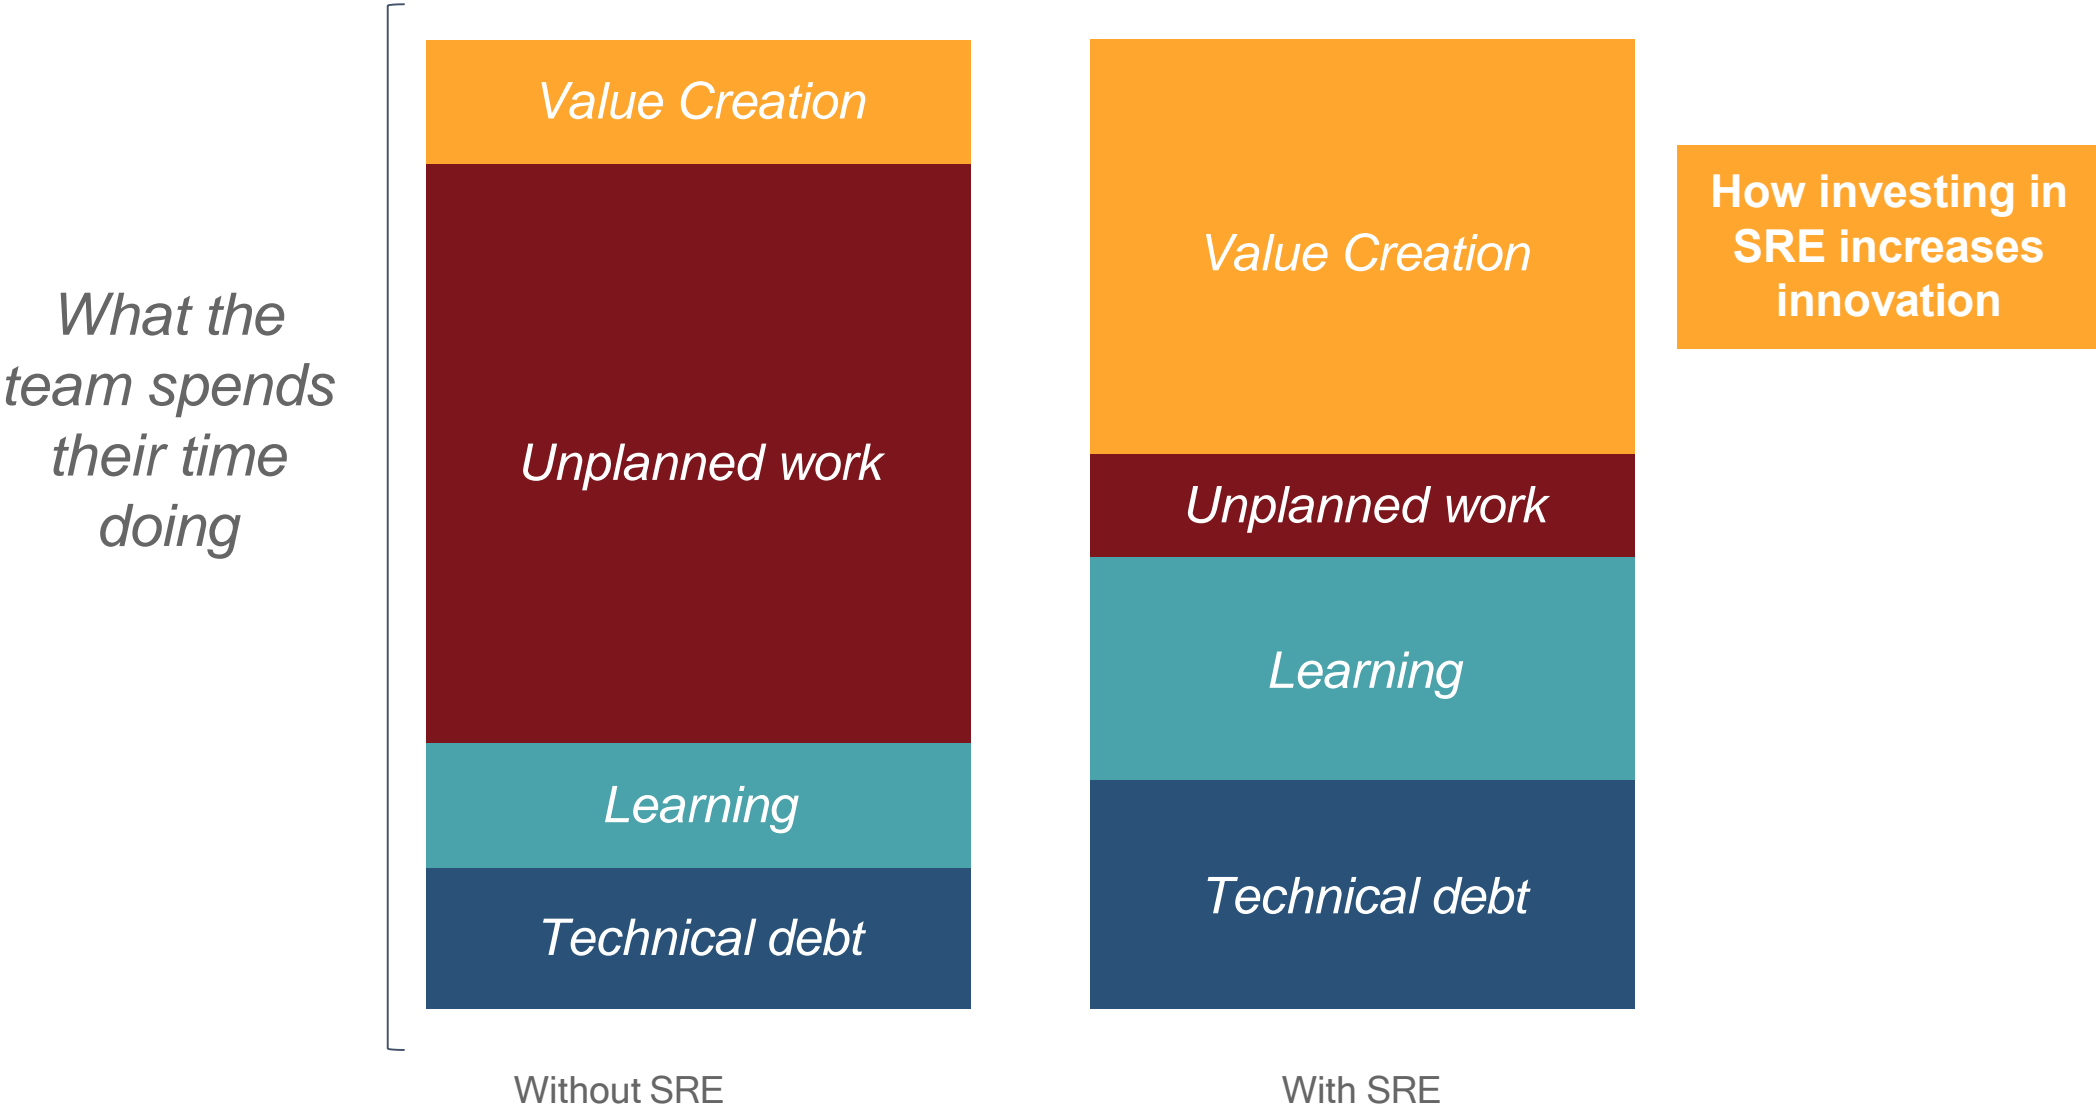
Load testing or other capacity management activities



● N/A ● Minor ● Moderate ● Major



Reducing unplanned work and technical debt



SRE principles and practices

Culture

Reliability @ Scale, Shift-Left "Wisdom of Production", and Continuous Improvement

Toil Reduction

Reduce Non-Value Add Work using Tooling and Automation

SLAs/SLOs/SLIs

Metrics such as Availability, Latency, and Response Time with Error Budgets

Measurements

Observability, Monitoring, Telemetry, and Instrumentation

Anti-Fragility

Improve Resilience using Fire Drills, Chaos Monkey, Security and Automation

Continuous Integration (CI)

Backlog & Design

Code & Test

Commit & Merge

Build & Test

Pipeline

Artifacts

Continuous Delivery / Deployment (CD)

SAT & UAT

Approve Release

Deploy to Prod

Post-Prod Tests

Operate

Work Sharing

Work Technical Debt in Small Increments

Manage Load % for Ops, Dev and On-Call Work

Deployments

Gradual Releases using Green/Blue, A/B, Canary Deployments, Automation Scripts, Testing and Monitoring

Performance Management

Monitoring, APM, Capacity Testing & Auto-Scaling

Incident Management

Emergency Response, 50% Ops/Dev Load, 25% On-Call Load, and Blameless Retrospectives



SLAs, SLOs and SLIs

Service Level...

In SRE services are managed to the SLO

SLA	SLO	SLI
<i>Agreement</i>	<i>Objective</i>	<i>Indicator</i>
<i>A business contract that comes into effect when your users are so unhappy you have to compensate them in some fashion</i>	<i>Specify a target level for the reliability of your service e.g., what the success rate should be 98% (it's never 100%)</i>	<i>An indicator of the level of service that you are providing e.g., http request success rate 99%</i>

SLOs need consequences if they are violated




The VALET dimensions of SLO

	<i>Dimension</i>	<i>SLO</i>	<i>Budget</i>	<i>Policy</i>
V	<i>Volume/traffic</i>	<i>Does the service handle the right volumes of data or traffic?</i>	<i>Budget: 99.99% of HTTP requests per month succeed with 200 OK</i>	<i>Address scalability issues</i>
A	<i>Availability</i>	<i>Is the service available to users when they need it?</i>	<i>Budget: 99.9% availability/uptime</i>	<i>Address downtime issues/outages, zero downtime deployments</i>
L	<i>Latency</i>	<i>Does the service deliver in a user-acceptable period of time?</i>	<i>Payload of 90% of HTTP responses returned in under 300ms</i>	<i>Address performance issues</i>
E	<i>Errors</i>	<i>Is the service delivering the capabilities being requested?</i>	<i>0.01% of HTTP requests return 4xx or 5xx status codes</i>	<i>Analyze and respond to main status codes, new functionality or infrastructure may be required</i>
T	<i>Tickets</i>	<i>Are our support services efficient?</i>	<i>75% of service tickets are automatically resolved</i>	<i>Automate more manual processes</i>



Error budgets



The SLO is a proxy for customer happiness



Error budgets by SLI and SLO

SLI

[Metric identifier] [Operator] [Metric]

Home page request served in < 100 ms

95th percentile of Home page latency over 5 mins < 200ms

Requests should be completed within 250 ms

Services should be available for 99.99% of time (based on heartbeat events from bounded system)

Book page request response code != 5xx

SLO

[Objective] [SLI] [Period]

95% of home page requests served in < 100ms over past 24 hours

99% of 95th percentile of home page latency over 5 mins < 200ms for the past month

95% of requests should be completed within 250 ms over 24 hours

95% of Services should be available for 99.99% of time over 30 days

99% of book page request response code != 5xx over the past 7 days

ERROR BUDGETS

[Error Budget] [SLI]

Allow 5% failure of home page requests served in < 100ms over past 24 hours

Allow 1% failure of 95% percentile home latency over 5 minutes < 200ms for the past month

Allow 5% failure of requests should be completed within 250 ms over 24 hours

Allow 5% failure of services availability over 30 days

Allow for 1% failure of book page request response code != 5xx over the last 7 days

Should you automate everything?

HOW LONG CAN YOU WORK ON MAKING A ROUTINE TASK MORE EFFICIENT BEFORE YOU'RE SPENDING MORE TIME THAN YOU SAVE?
(ACROSS FIVE YEARS)

		HOW OFTEN YOU DO THE TASK					
		50/DAY	5/DAY	DAILY	WEEKLY	MONTHLY	YEARLY
HOW MUCH TIME YOU SHAVE OFF	1 SECOND	1 DAY	2 HOURS	30 MINUTES	4 MINUTES	1 MINUTE	5 SECONDS
	5 SECONDS	5 DAYS	12 HOURS	2 HOURS	21 MINUTES	5 MINUTES	25 SECONDS
	30 SECONDS	4 WEEKS	3 DAYS	12 HOURS	2 HOURS	30 MINUTES	2 MINUTES
	1 MINUTE	8 WEEKS	6 DAYS	1 DAY	4 HOURS	1 HOUR	5 MINUTES
	5 MINUTES	9 MONTHS	4 WEEKS	6 DAYS	21 HOURS	5 HOURS	25 MINUTES
	30 MINUTES		6 MONTHS	5 WEEKS	5 DAYS	1 DAY	2 HOURS
	1 HOUR		10 MONTHS	2 MONTHS	10 DAYS	2 DAYS	5 HOURS
	6 HOURS				2 MONTHS	2 WEEKS	1 DAY
	1 DAY					8 WEEKS	5 DAYS



SRE and Observability/Monitoring


Observability is a characteristic of systems; that they can be observed. It's closely related to a DevOps tenet: 'telemetry everywhere', meaning that anything we implement is emitting data about its activities. It requires intentional behavior during digital product and platform design and a conducive architecture. It's not monitoring. Monitoring is what we do when we observe our observable systems and the tools category that largely makes this possible.

Monitoring Tool Usage

Infrastructure monitoring (e.g., disk or CPU)



Network performance monitoring or diagnostics (e.g., latency or saturation)



Application performance monitoring (e.g., tracing or events)



Digital experience monitoring (e.g., Synthetics or RUM)



Artificial intelligence for ITOps (e.g., anomaly detection or self-healing)



Public/social sentiment monitoring




Competitive benchmarking intelligence



Never Rarely Sometimes Always

Received AIOps Value



Low (1-3)
Moderate (4)
High (7-9)

*Based on a 1-9 value scale.



SRE persona


How Observability Supports SRE's Goals

- *Reducing the toil associated with incident management – particularly around cause analysis – improving uptime and MTTR*
- *Providing a platform for inspecting and adapting according to SLOs and ultimately improving teams' ability to meet them*
- *Offering a potential solution to improve when SLOs are not met, and error budgets are over-spent*
- *Relieving team cognitive load when dealing with vast amounts of data – reducing burnout*
- *Releasing humans and teams from toil, improving productivity, innovation and the flow and delivery of value*
- *Supporting multifunctional, autonomous teams and the “we build it, we own it” DevOps mantra*
- *Completing the value stream cycle by providing insights around value outcomes that can be fed back into the innovation phase*





Monitoring Data Usage Drivers





Use Case Automation Levels

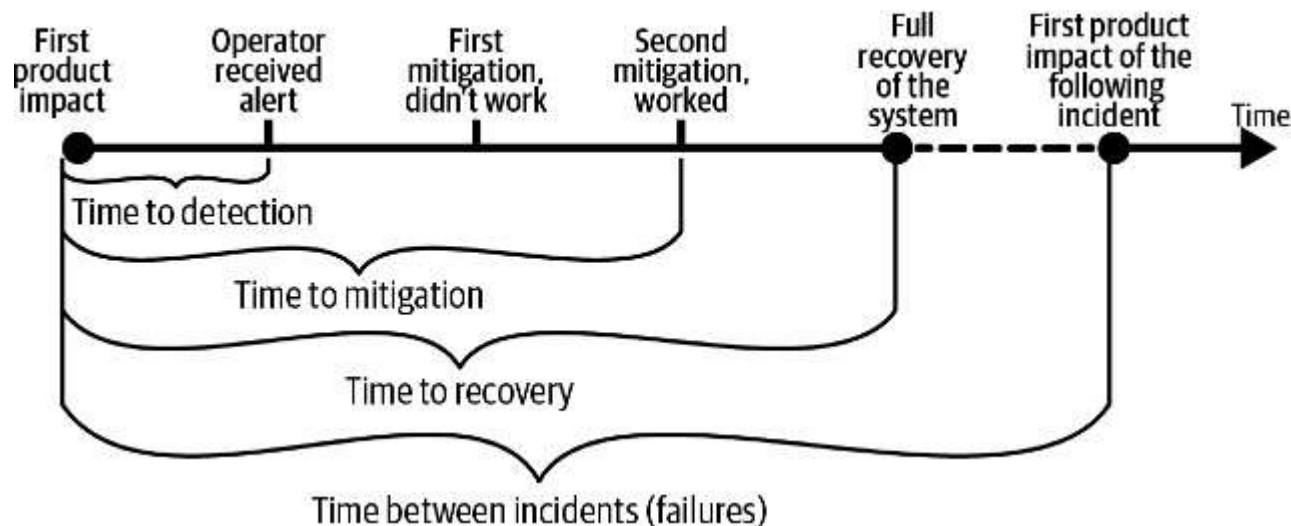


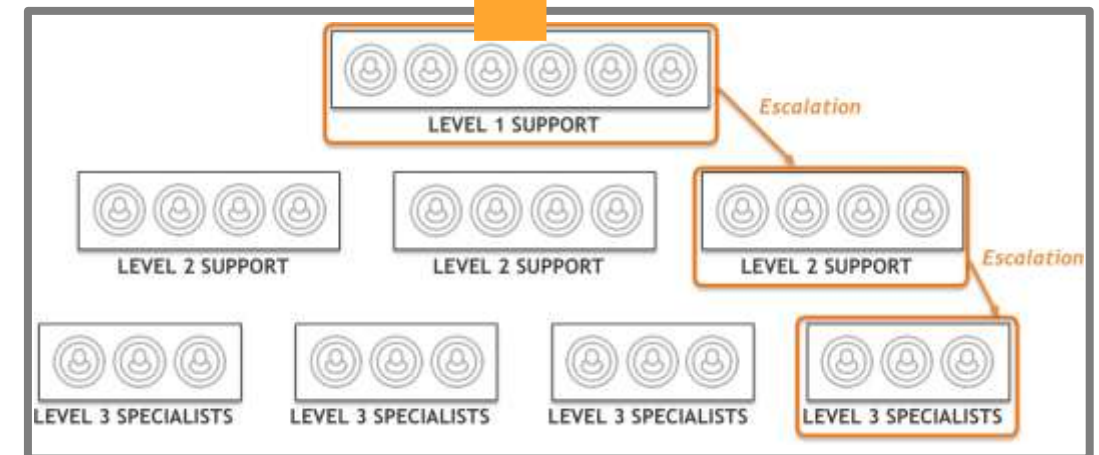
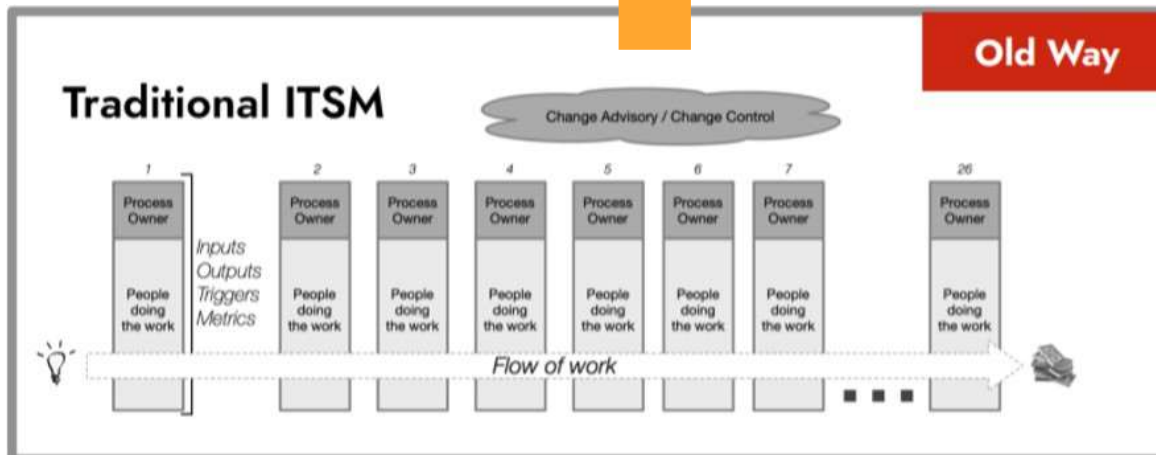
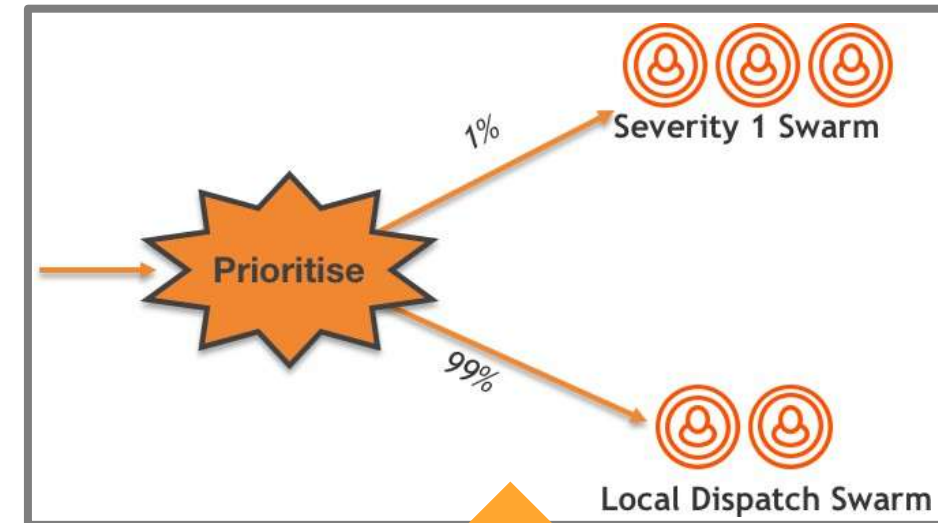
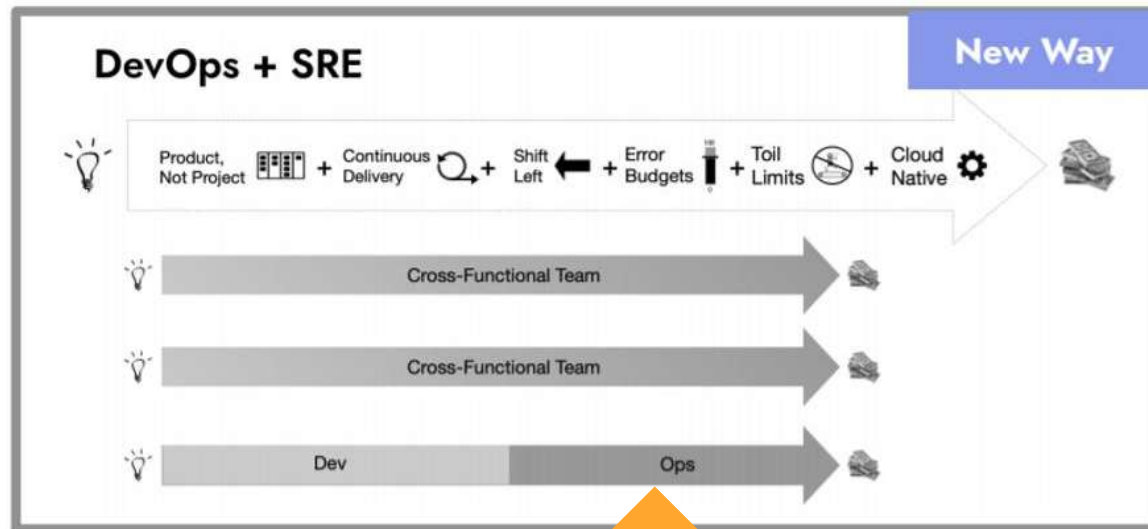


Good practices for Incident Management

One of the key responsibilities of SRE is to manage incidents of the production system(s) that they are responsible for. Within an incident, SREs contribute to debugging the system, choosing the right immediate mitigation, and organizing the incident response if it requires broader coordination

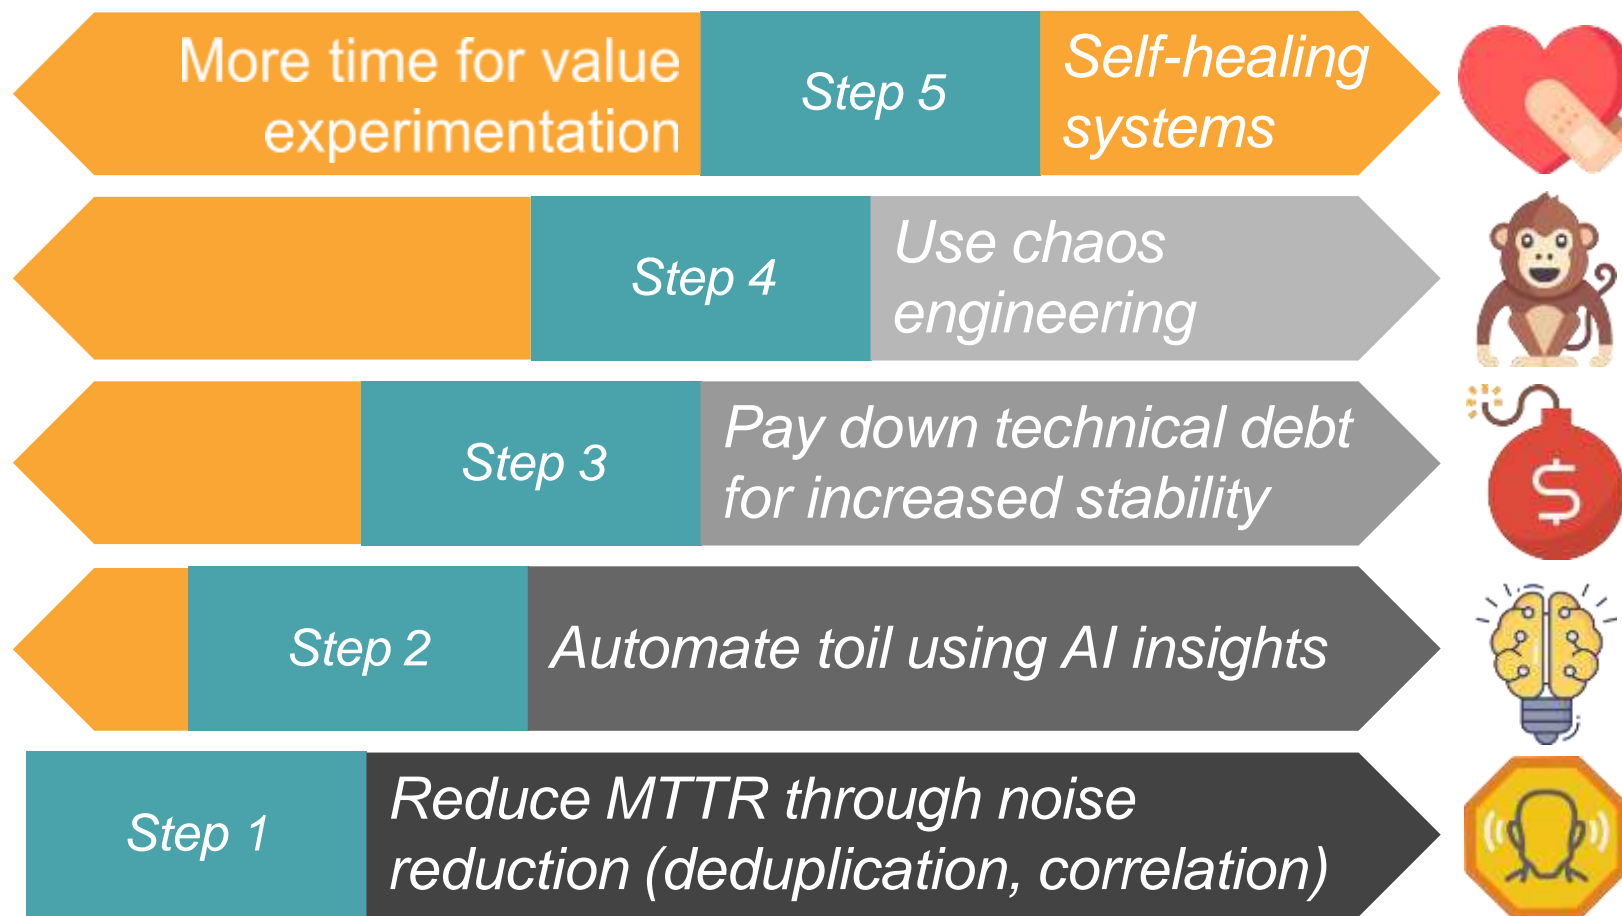
- Defect prevention
- Strategies for deploy/roll back/ roll forward (Feature Flags, Blue-Green, Canary)
- Auto remediation
- Reverting system





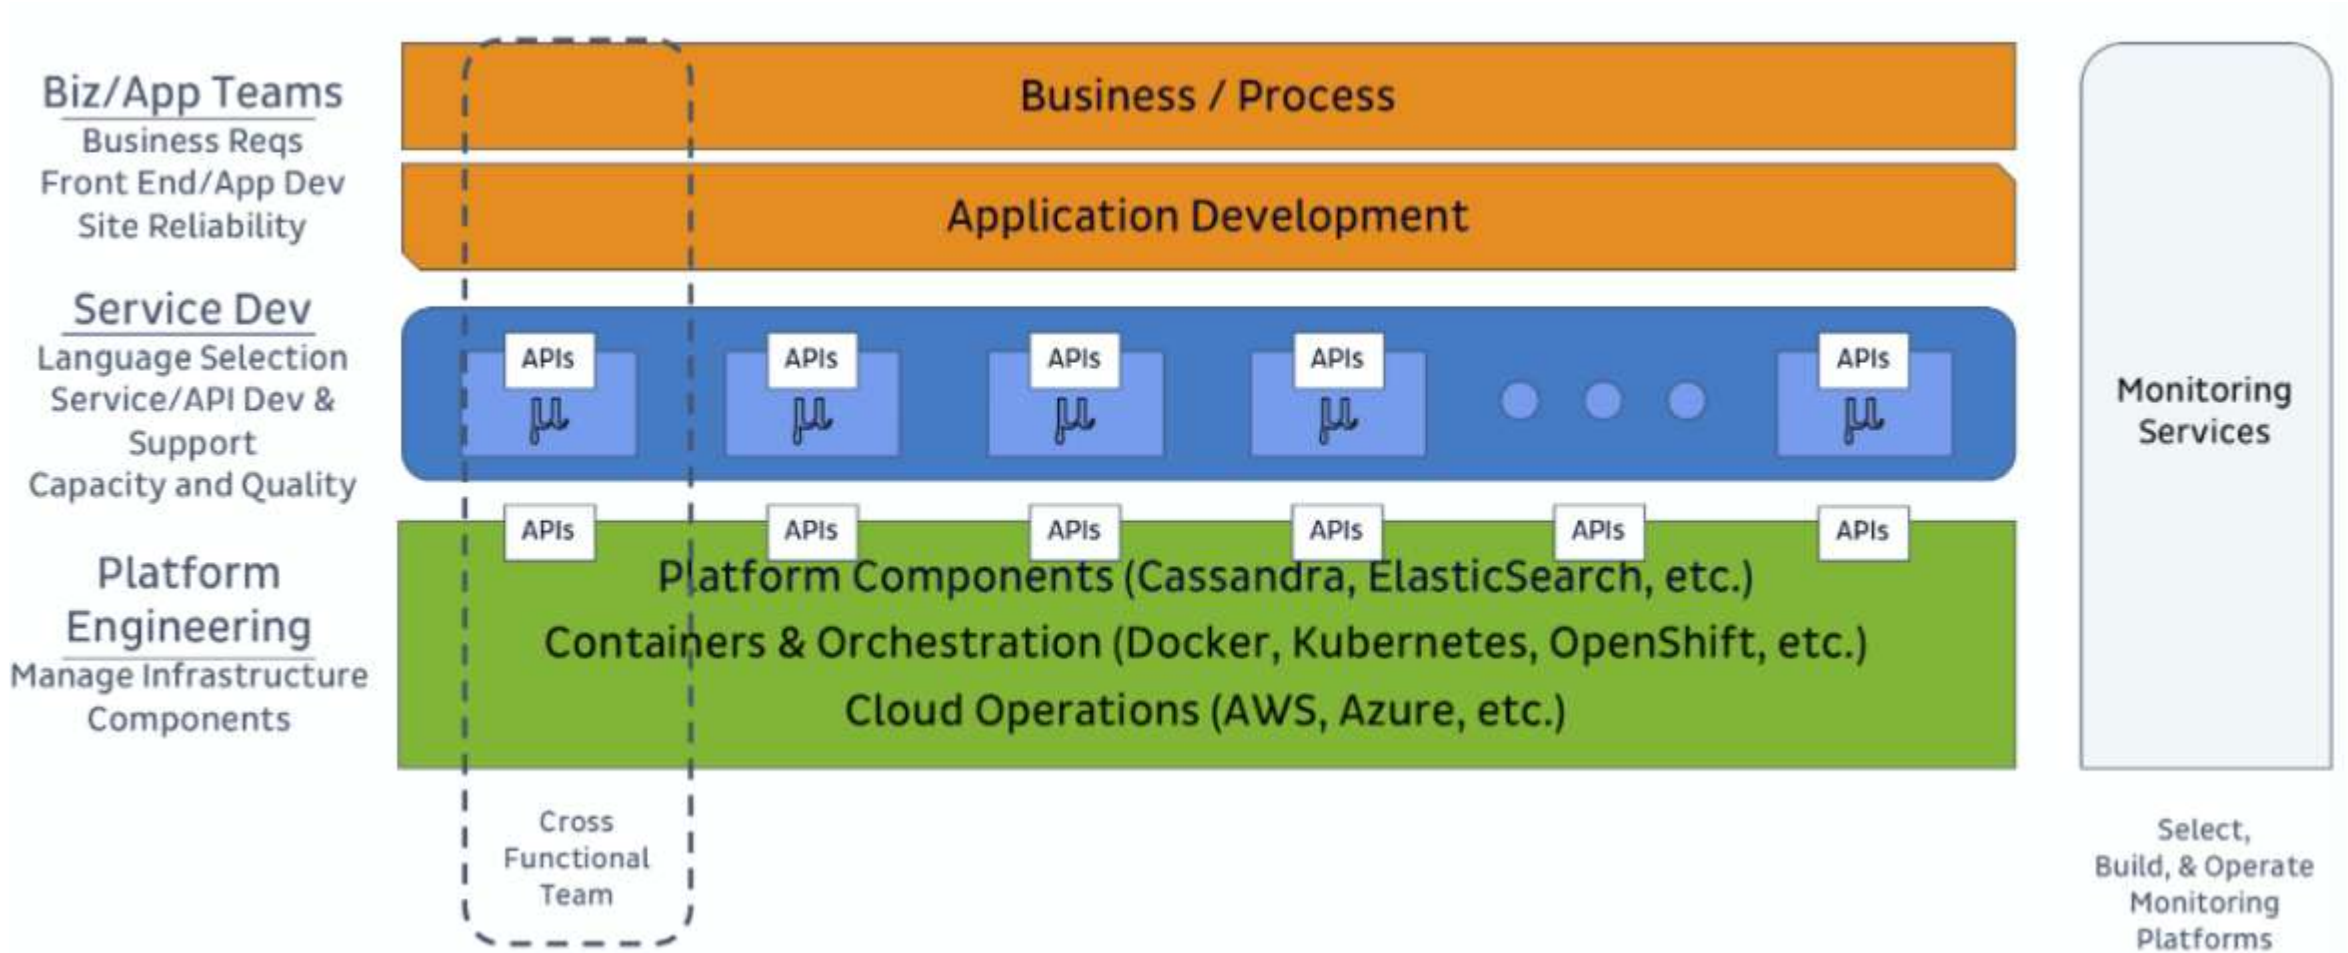


IM, Observability, AIOps and the SRE





Platform SRE





Platform SRE ushers in self-service

- *The Platform provides “self-service” provision of infrastructure, functionalities, configurations and environments that can be consumed by development teams , third parties e.g. distributed teams and partners*
- *Embedded governance, controls and standards are built-in*
- *End-to-end deployment automation, infrastructure playbooks of a service or application*
- *Abstraction of infrastructure specific implementations for multi/hybrid cloud through runbooks and playbooks*
- *In-Source code, products built by platform teams can be extended or enhanced by SRE/Dev/Ops or any other*



Chaos Engineering

The discipline of experimenting on a distributed system in order to build confidence in the system's ability to withstand turbulent conditions.

Properties of a Chaos Experiment

- *Define steady state*
- *Formulate hypothesis*
- *Outline methodology*
- *Identify blast radius*
- *Observability is key*
- *Readily abortable*



Getting started with Chaos Engineering

From a technical point of view, they are easy to set up and do not have to be sophisticated in terms of implementation

- *Get relevant people in a room who are responsible for a system or set of systems*
- *Shut off a component that the system should be robust enough to tolerate without it*
- *Record the learning outcomes and bring the component back online*

Apply It: Build a Game Day event

What / Who / When / Where / How

Note: *You don't need to run your GameDay in production! Insights can come from conducting experiments first in a staging or test environment*



Implementation challenges

And how to overcome them

You don't have enough cross-team usage or buy-in

Your difficult and dense process is slowing down incident response

Postmortems are underutilized and don't encompass in-depth learnings

You wait for incidents to happen

You stop at incident management without SLOs

Loop in sales, support and customer success

Customize lightweight, lean checklists

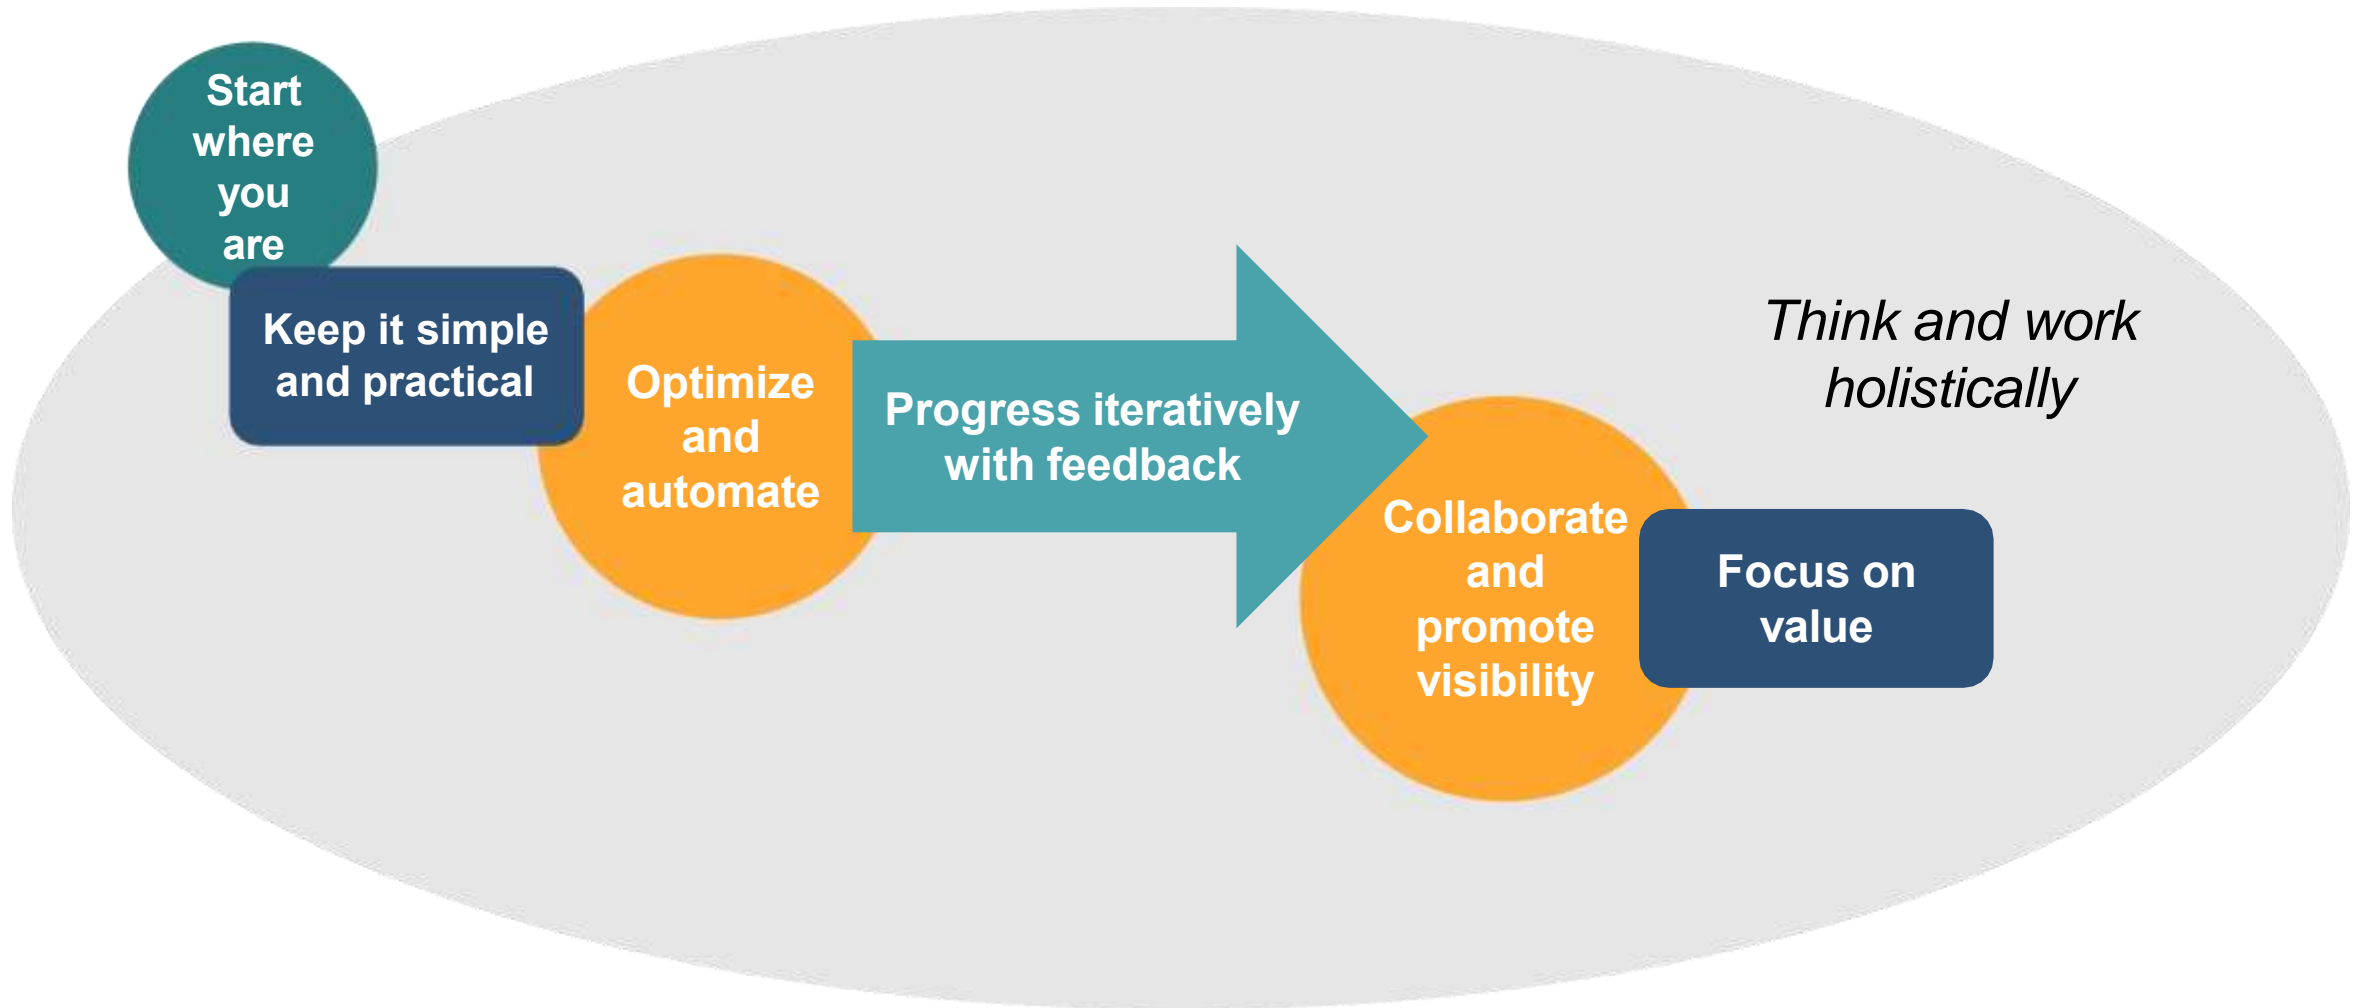
Automate: turn unstructured postmortems into a taxonomy with metadata & data

Practice chaos engineering

Automate: visualize your metrics



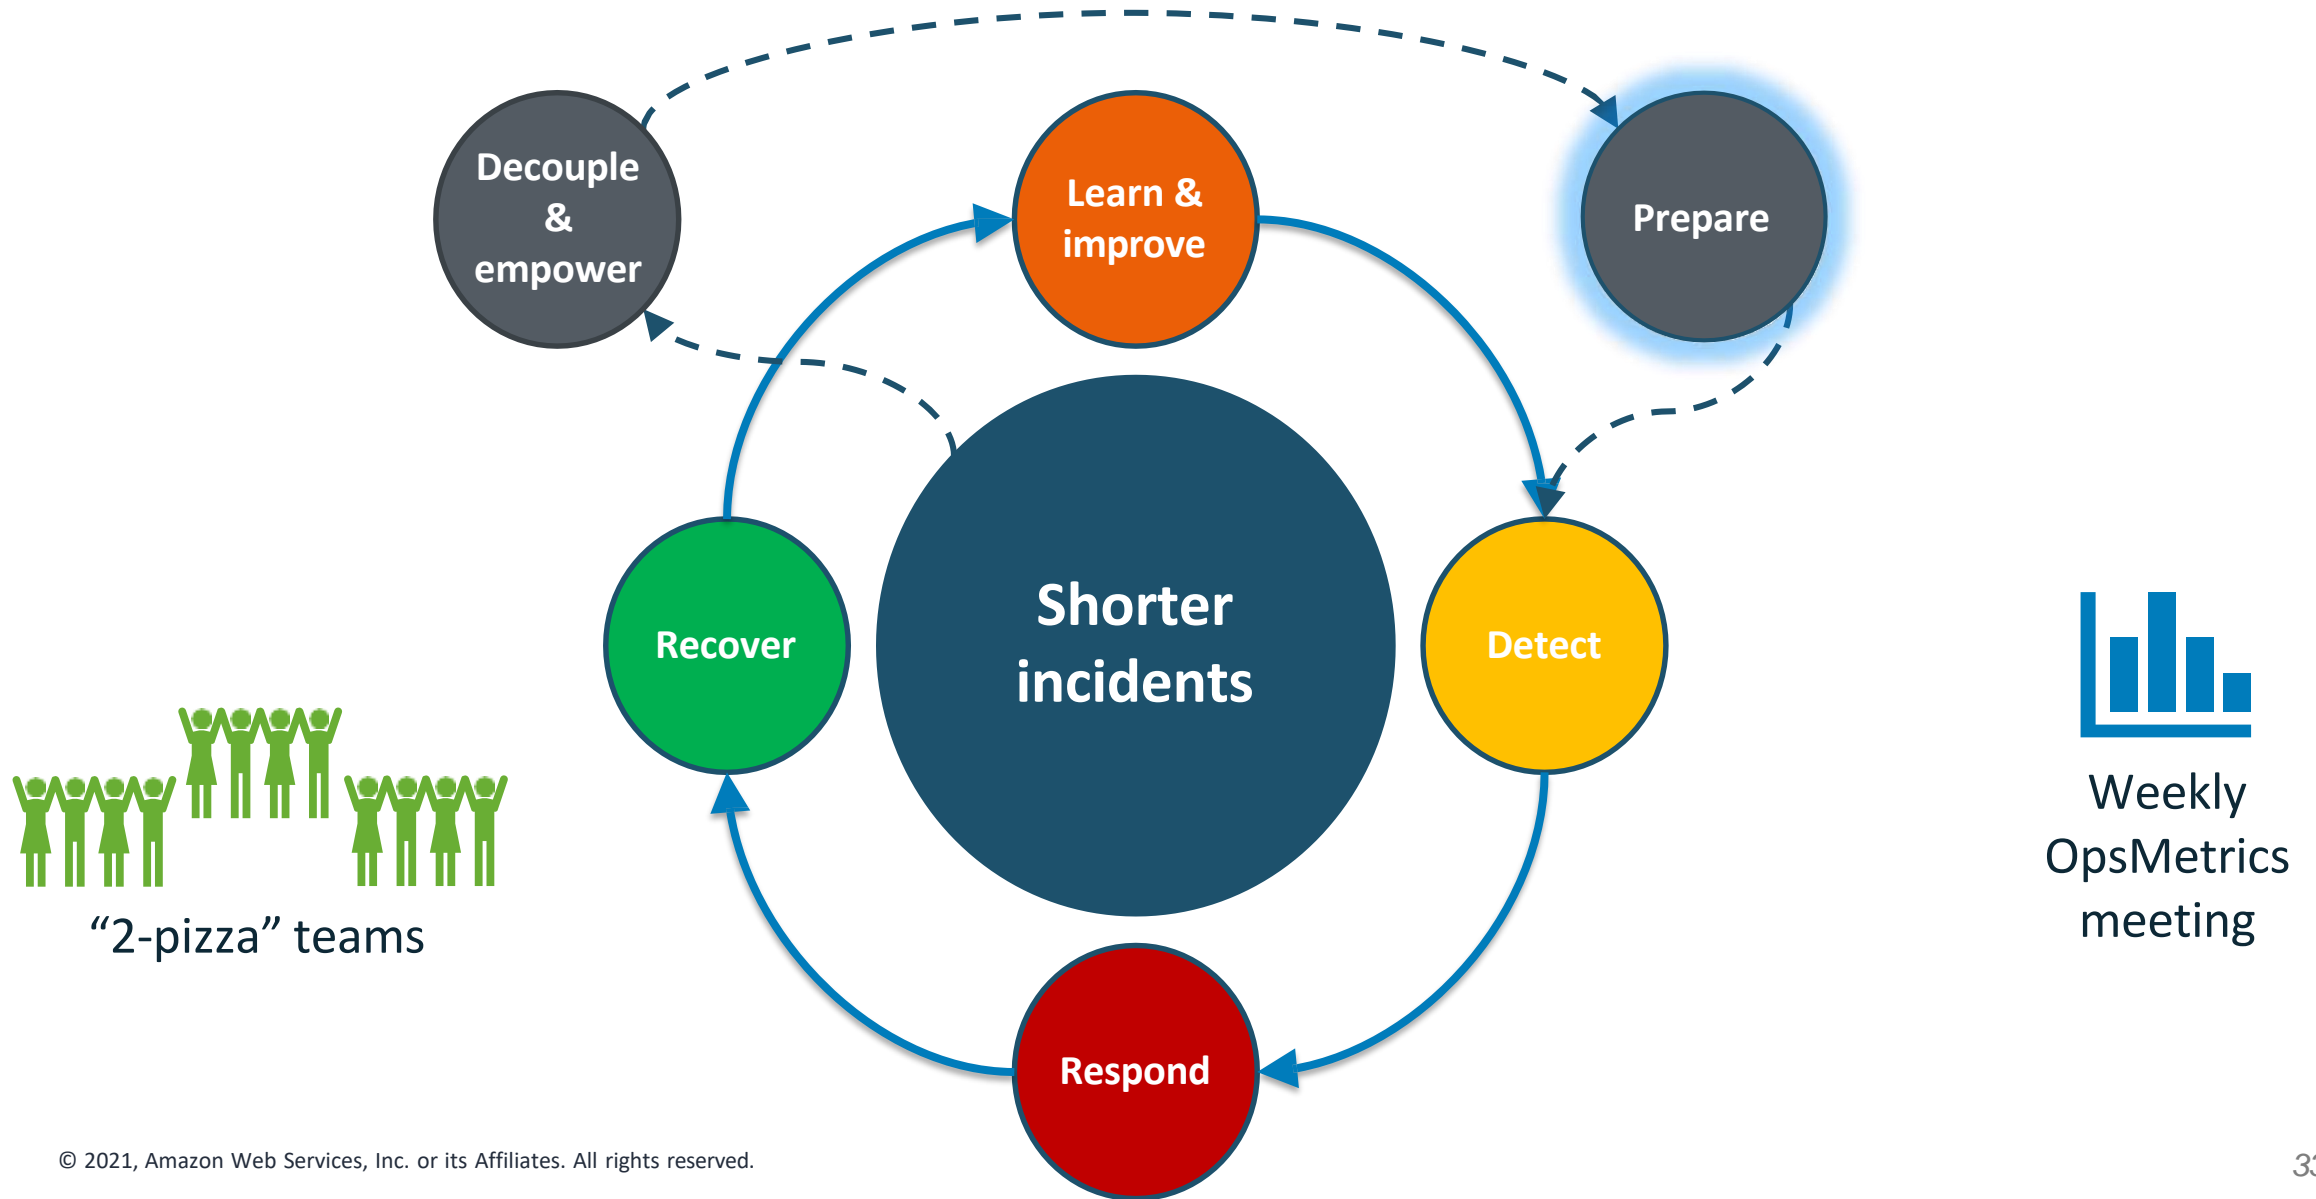
Where to start



Agenda

- The uptime flywheel - prepare
- Prepare for downtime
- AWS Services for incident management
- The uptime flywheel – learn
- Chaos engineering
- AWS Fault Injection Simulator
- Summary and Marketplace next steps

The uptime flywheel @ AWS



Prepare for downtime

1. Detect

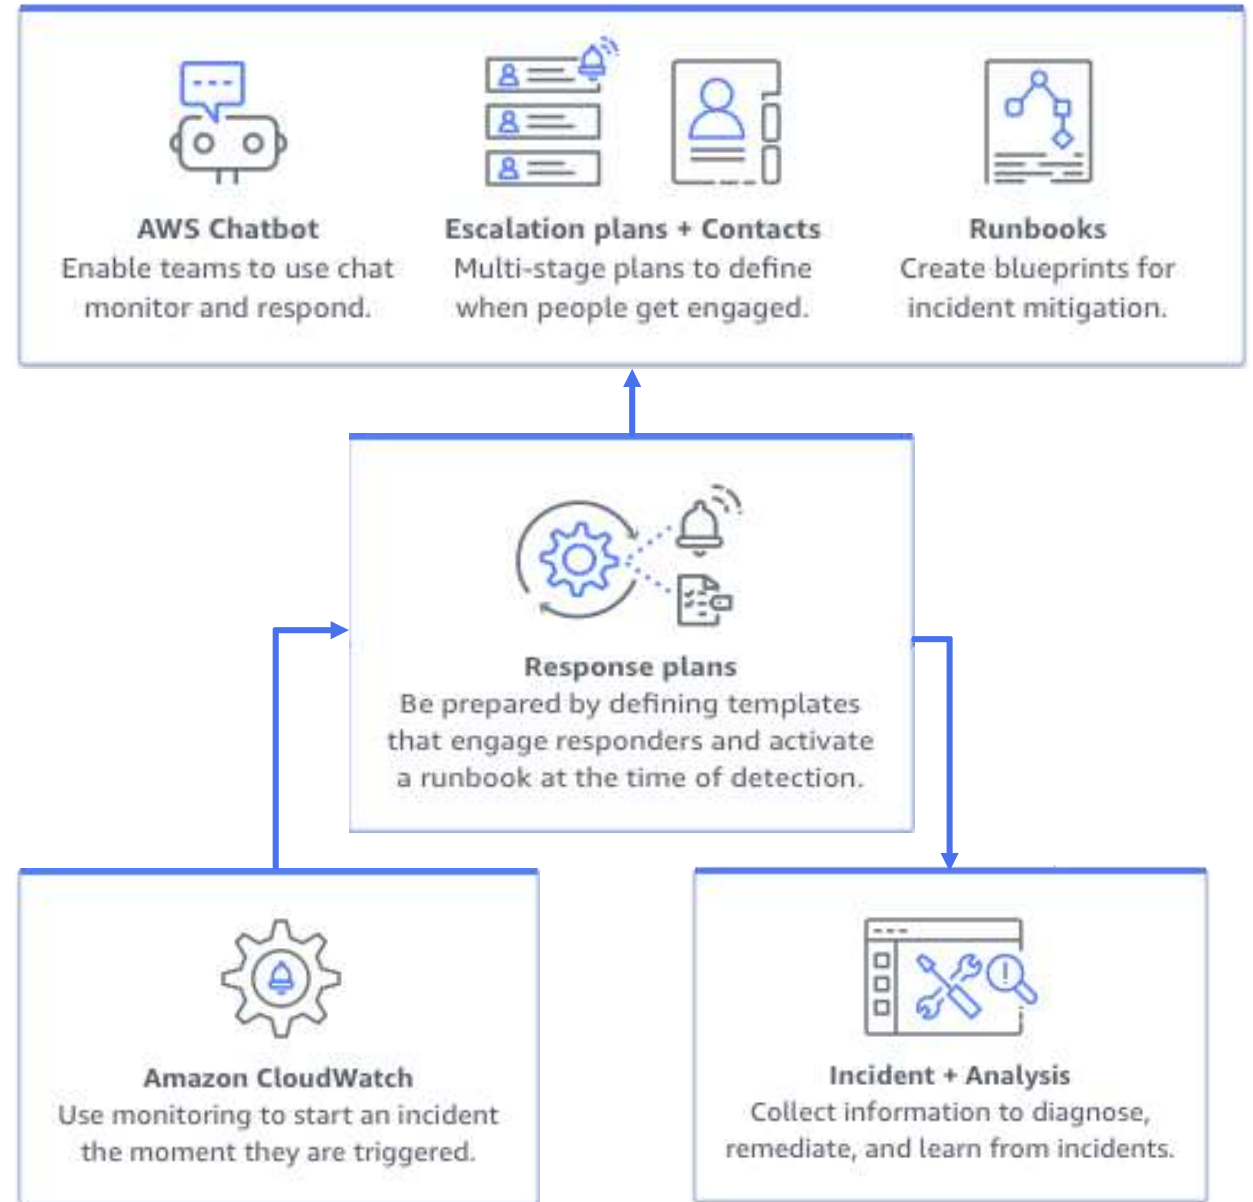
- What could go wrong?
- How would I know?

2. Respond and recover

- Who needs to be engaged?
- What do they need to do to diagnose?
 - Procedures and scripts
 - Where do we collaborate?

3. Learn and improve

- How did we respond?
- What actions will we take?



AWS Systems Manager

Centralize operational data from multiple AWS services and automate tasks across your AWS resources

Benefits

- Simplify resources and application management
- Easy to operate and securely manage multi-cloud infrastructures at scale
- Resolve critical application availability and performance
- Prepare for and manage incidents efficiently with automated response

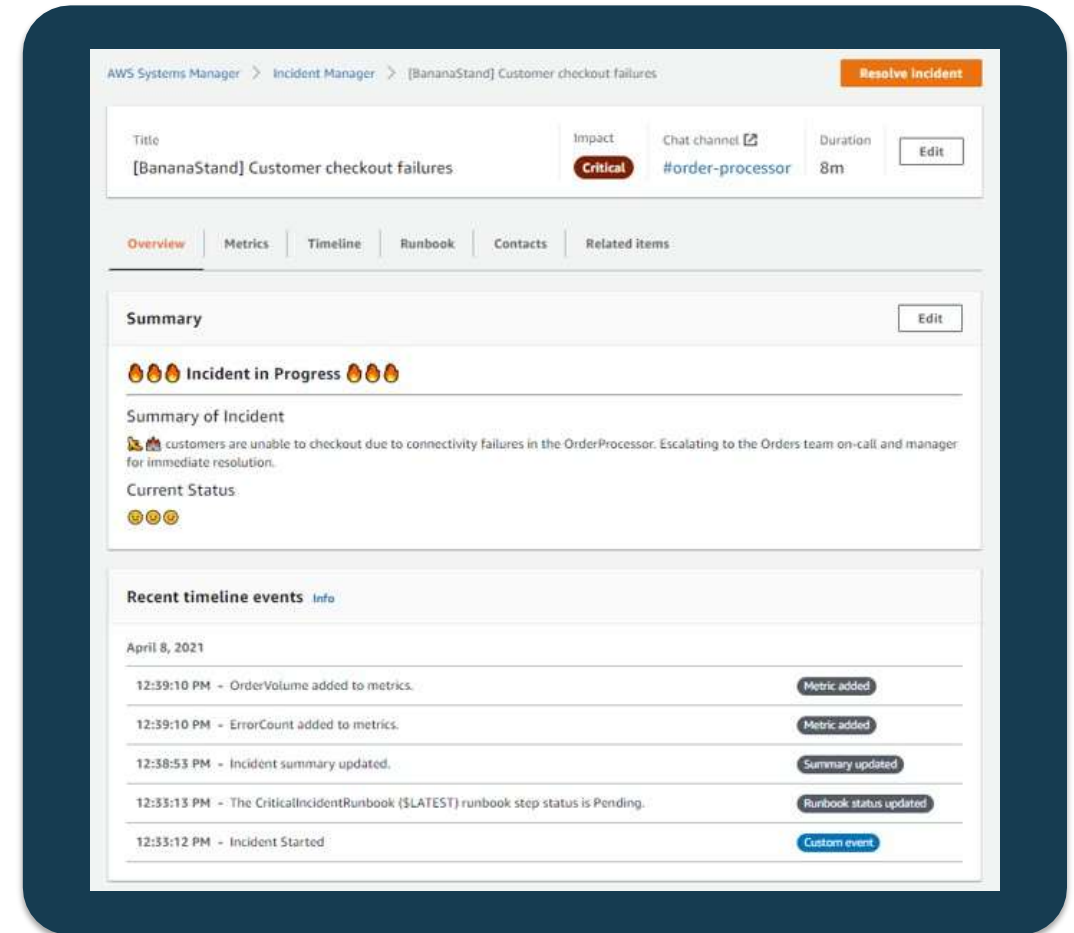


Incident Manager

Resolve application issues faster with automated response plans



- ➔ Specify a response plan to critical application alarms, including who to engage, what runbook to follow, and where to collaborate
- ➔ Notify the right people immediately with SMS, voice, and escalations (additional partner integrations coming soon)
- ➔ Single console to track incidents from detection to mitigation and post-incident analysis, including timeline, runbooks, metrics, etc.
- ➔ Collaborate in Slack via AWS Chatbot to resolve incidents
- ➔ Identify post-incident action items, such as improving alarms or automating runbooks steps, using Amazon's post-incident analysis template and track them in OpsCenter



OverviewMetricsTimelineRunbookEngagementsRelated itemsProperties

RunbookInfo

Cancel

RunbookBananaStand-MajorIncident (v1)

Execution detailsExecution details: BananaStand-MajorIncident

Status⌚ Waiting

00:00:00(UTC-7:00)

13:23:28

Triage

⌚ Waiting

Determine customer impact

- View the Metrics tab of the incident or navigate to your CloudWatch Dashboards to find key performance indicators (KPIs) that show the extent of customer impact.
- Use CloudWatch Synthetics and Contributor Insights to identify real-time failures in customer workflows.

Communicate customer impact

Update the following fields to accurately describe the incident:

- Title - The title should be quickly recognizable by the team and specific to the particular incident.
- Summary - The summary should contain the most important and up-to-date information to quickly onboard new responders to the incident.
- Impact - Select one of the following impact ratings to describe the incident:
 - 1 – Critical impact, full application failure that impacts many to all customers.
 - 2 – High impact, partial application failure with impact to many customers.
 - 3 – Medium impact, the application is providing reduced service to many customers.
 - 4 – Low impact, the application is providing reduced service to few customers.
 - 5 – No impact, customers are not currently impacted but urgent action is needed to avoid impact.

Resume

Diagnosis

⌚ Pending

Rollback

- Look for recent changes to the production environment that might have caused the incident. Engage the responsible team using the Contacts tab of the incident.
- Rollback these changes if possible.

Locate failures

- Review metrics and alarms related to your Application. Add any related metrics and alarms to the Metrics tab of the incident.
- Use CloudWatch ServiceLens to troubleshoot issues across multiple services.
- Investigate the possibility of ongoing incidents across your organization. Check for known incidents and issues in AWS using Personal Health Dashboard. Add related links to the Related Items tab of the incident.
- Avoid going too deep in diagnosing the failure and focus on how to mitigate the customer impact. Update the Timeline tab of the incident when a possible diagnosis is identified.

© 2021, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

37

←

●

Diagnosis

Pending

Rollback

- Look for recent changes to the production environment that might have caused the incident. Engage the responsible team using the **Contacts** tab of the incident.
- Rollback these changes if possible.

Locate failures

- Review metrics and alarms related to your [Application](#). Add any related metrics and alarms to the **Metrics** tab of the incident.
- Use [CloudWatch ServiceLens](#) to troubleshoot issues across multiple services.
- Investigate the possibility of ongoing incidents across your organization. Check for known incidents and issues in AWS using [Personal Health Dashboard](#). Add related links to the **Related Items** tab of the incident.
- Avoid going too deep in diagnosing the failure and focus on how to mitigate the customer impact. Update the **Timeline** tab of the incident when a possible diagnosis is identified.

←

●

Mitigation

Pending

Collaborate

- Communicate any changes or important information from the previous step to the members of the associated chat channel for this incident. Ask for input on possible ways to mitigate customer impact.
- Engage additional contacts or teams using their escalation plan from the **Contacts** tab.
- If necessary, prepare an emergency change request in [Change Manager](#).

Implement mitigation

- Consider re-routing customer traffic or throttling incoming requests to reduce customer impact.
- Look for common runbooks in [Automation](#) or run commands using [Run Command](#).
- Update the **Timeline** tab of the incident when a possible mitigation is identified. If needed, review the mitigation with others in the associated chat channel before proceeding.

←

●

Recovery

Pending

Monitor customer impact

- View the **Metrics** tab of the incident to monitor for recovery of your key performance indicators (KPIs).
- Update the **Impact** field in the incident when customer impact has been reduced or resolved.

Identify action items

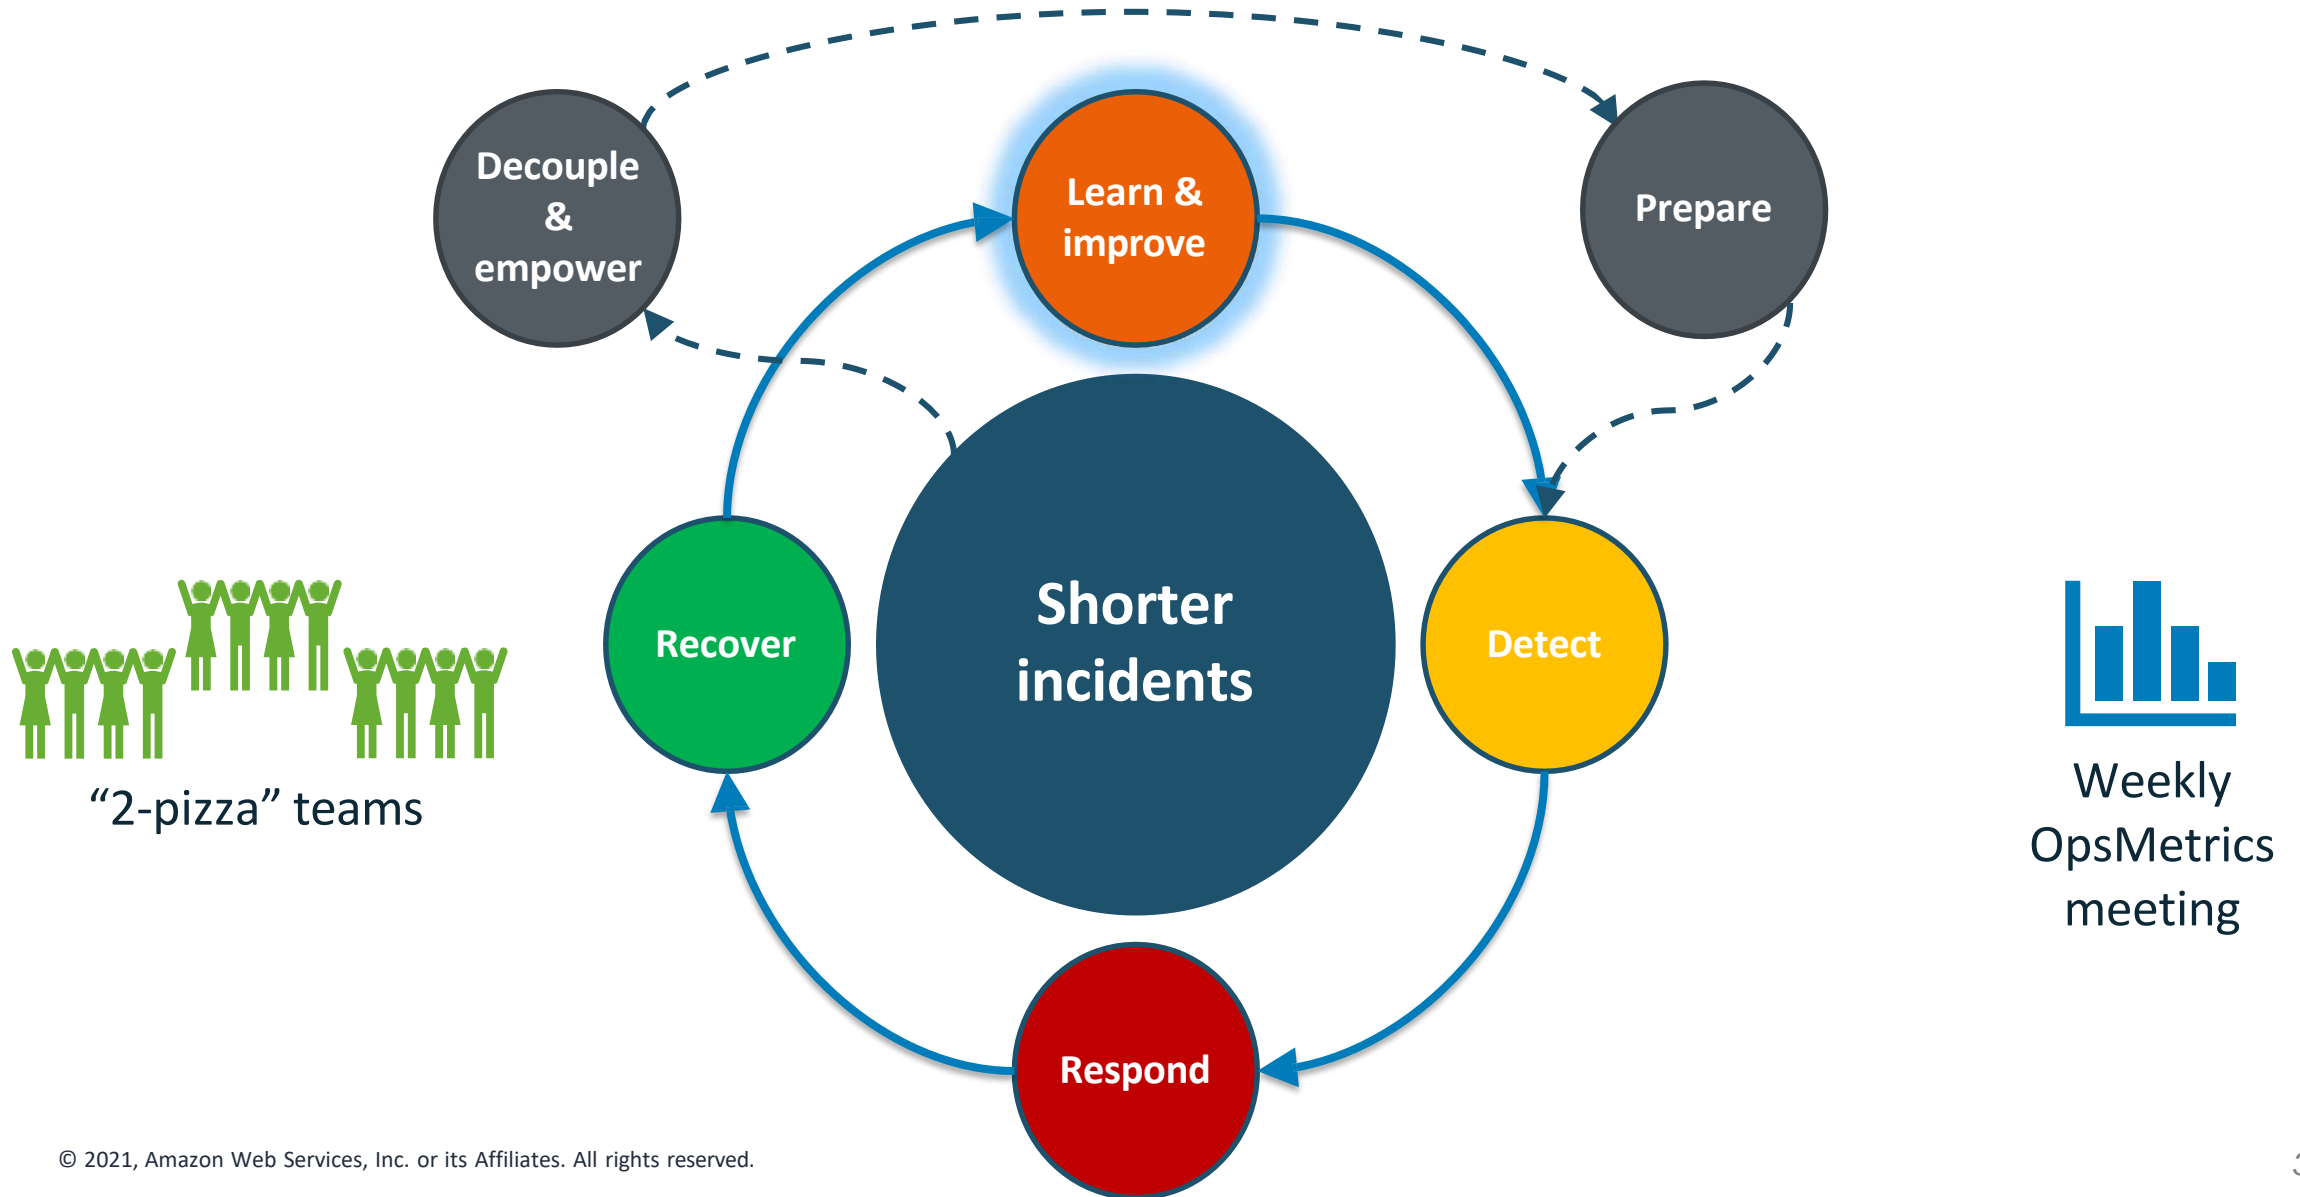
- Add entries in the **Timeline** tab of the incident to record key decisions and actions taken, including temporary mitigations that might have been implemented.
- Create a **Post-Incident Analysis** when the incident is closed in order to identify and track action items in [OpsCenter](#).

23:58:59

© 2021, Amazon Web Services, Inc. or its Affiliates. All rights reserved.

38

The uptime flywheel @ AWS




Chaos engineering

Experiment to ensure that the impact of failures is mitigated

Chaos experiment

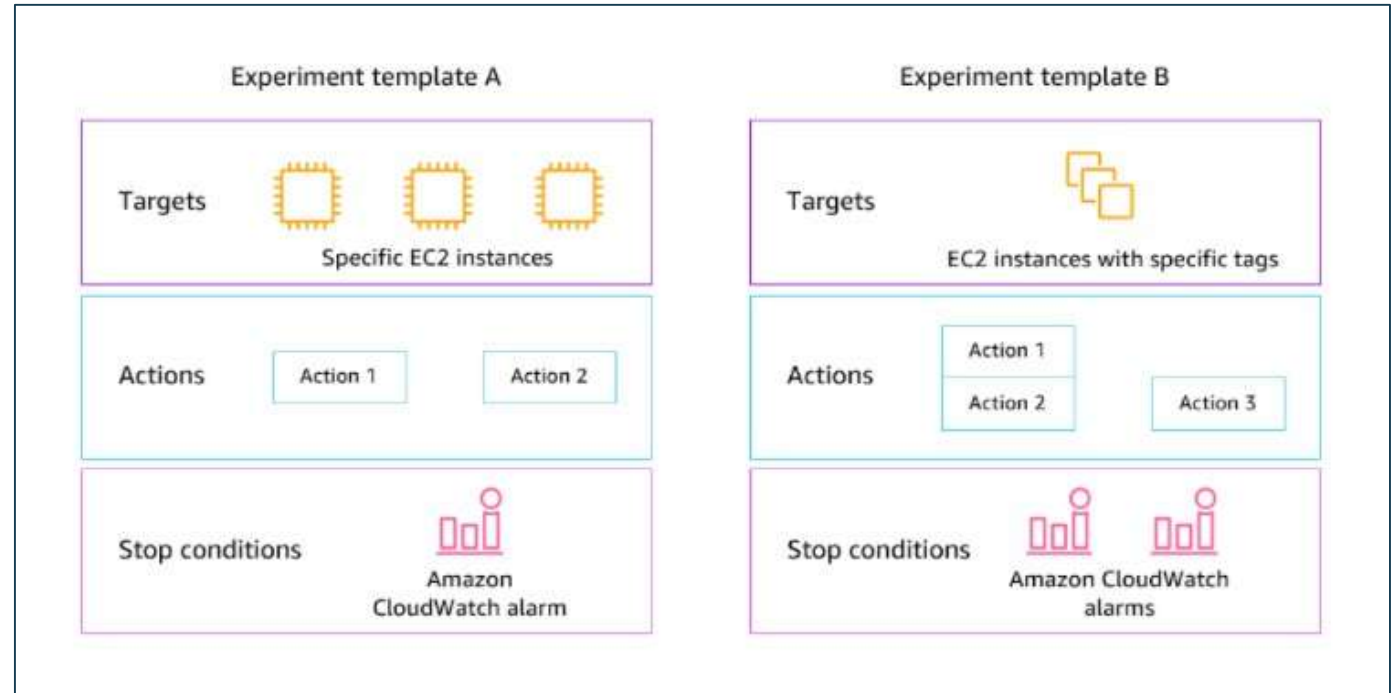
- Inject events that simulate
 - Hardware failures, like servers dying
 - Software failures, like malformed responses
 - Nonfailure events, like spikes in traffic or scaling events
 - Any event capable of disrupting steady state



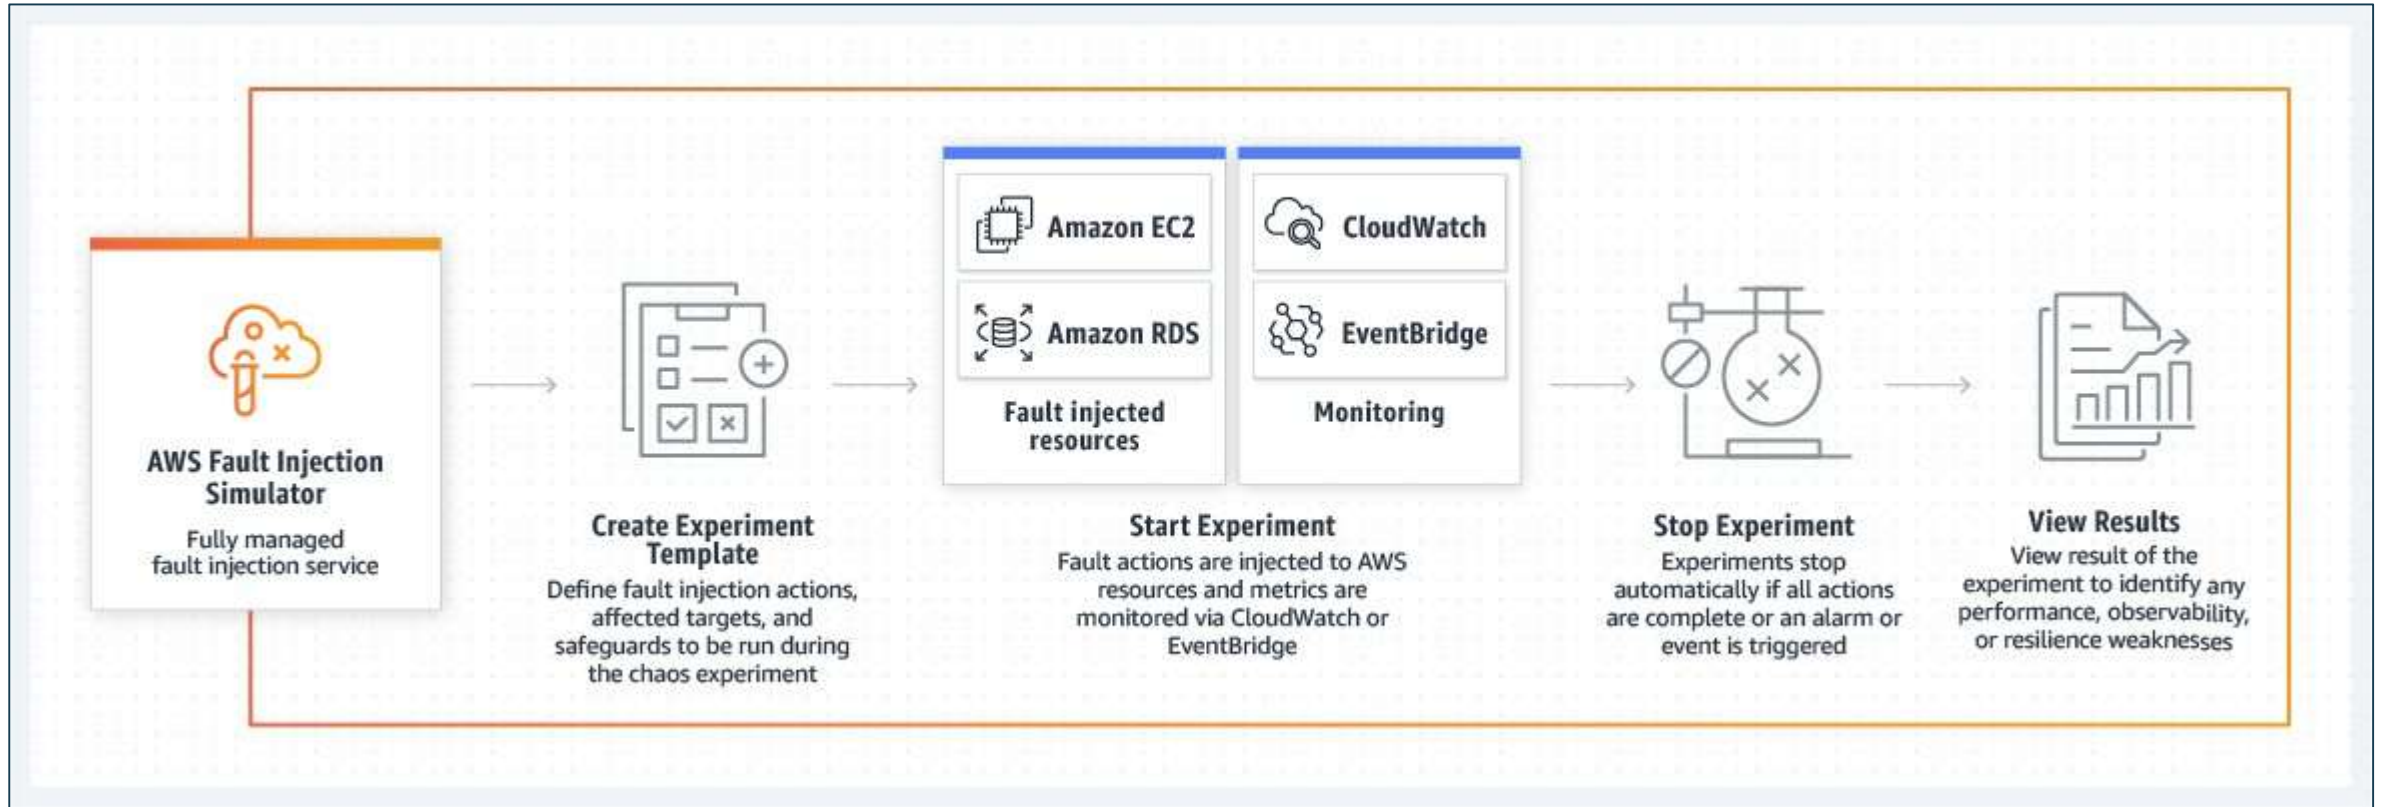
AWS Fault Injection Simulator

Improve resiliency and performance with controlled experiments

- A fast and easy way to get started with fault injection experiments
- Validate how your application performs on AWS
- Safeguard fault injection experiments
- Improve application performance, resiliency, and observability
- Get comprehensive insights by generating real-world failure conditions



AWS Fault Injection Simulator




AWS Fault Injection Simulator use case: periodic game days

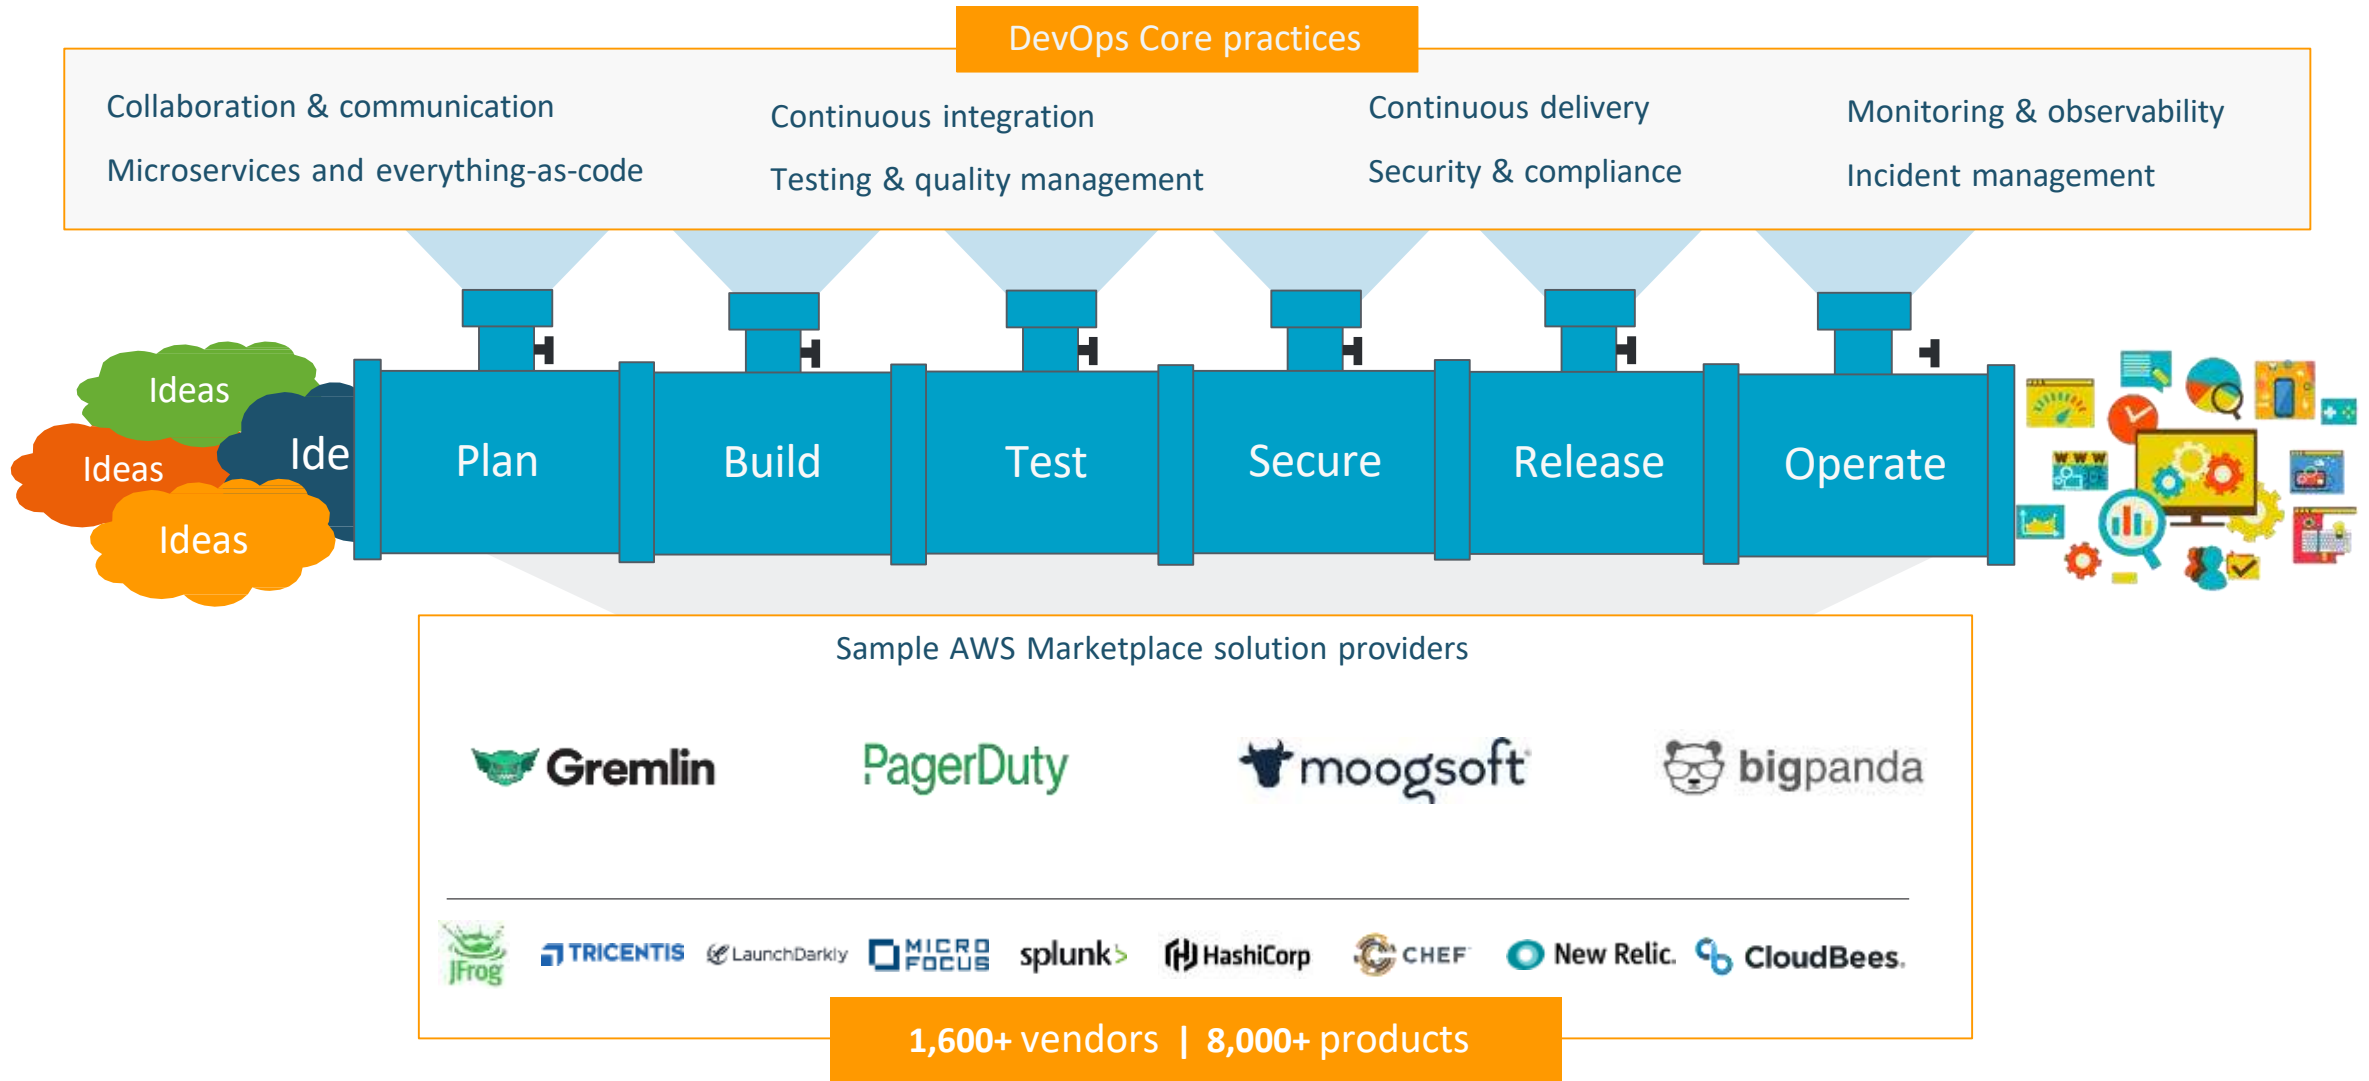
Why conduct a game day?

- Simulate a failure or event to test systems, processes, and team responses
- Should cover the areas of operations, security, reliability, performance, and cost
- Can be carried out with replicas of your production environment using AWS CloudFormation
- Should involve all personnel who normally operate a workload

Game day process



AWS Marketplace: Destination for third-party solutions to use with AWS





8,000+

listings



1,600+

ISVs



24

regions



290,000+

customers



1.5M+

subscriptions



And more coming soon!

How can you get started?

Find



A breadth of DevOps solutions:




Buy



Through flexible pricing options:

- Free trial
- Pay-as-you-go
- Budget alignment
- Bring Your Own License (BYOL)
- Private Offers
- Billing consolidation
- Enterprise Discount Program
- Private Marketplace

Deploy



With multiple deployment options:

- SaaS
- Amazon Machine Image (AMI)
- CloudFormation Template
- Containers
- Amazon EKS/ Amazon ECS
- AI / ML models
- AWS Data Exchange
- AWS Control Tower