


KSQL

Open-source streaming for Apache Kafka



KSQL and Kafka Streams in 3 minutes

In a nutshell



The streaming SQL engine
for Apache Kafka® to write
real-time applications in SQL



Apache Kafka® library to write
real-time applications and
microservices in Java and Scala

KSQL

is the


Streaming SQL Engine

for

Apache Kafka



Hello, Streaming World



KSQL



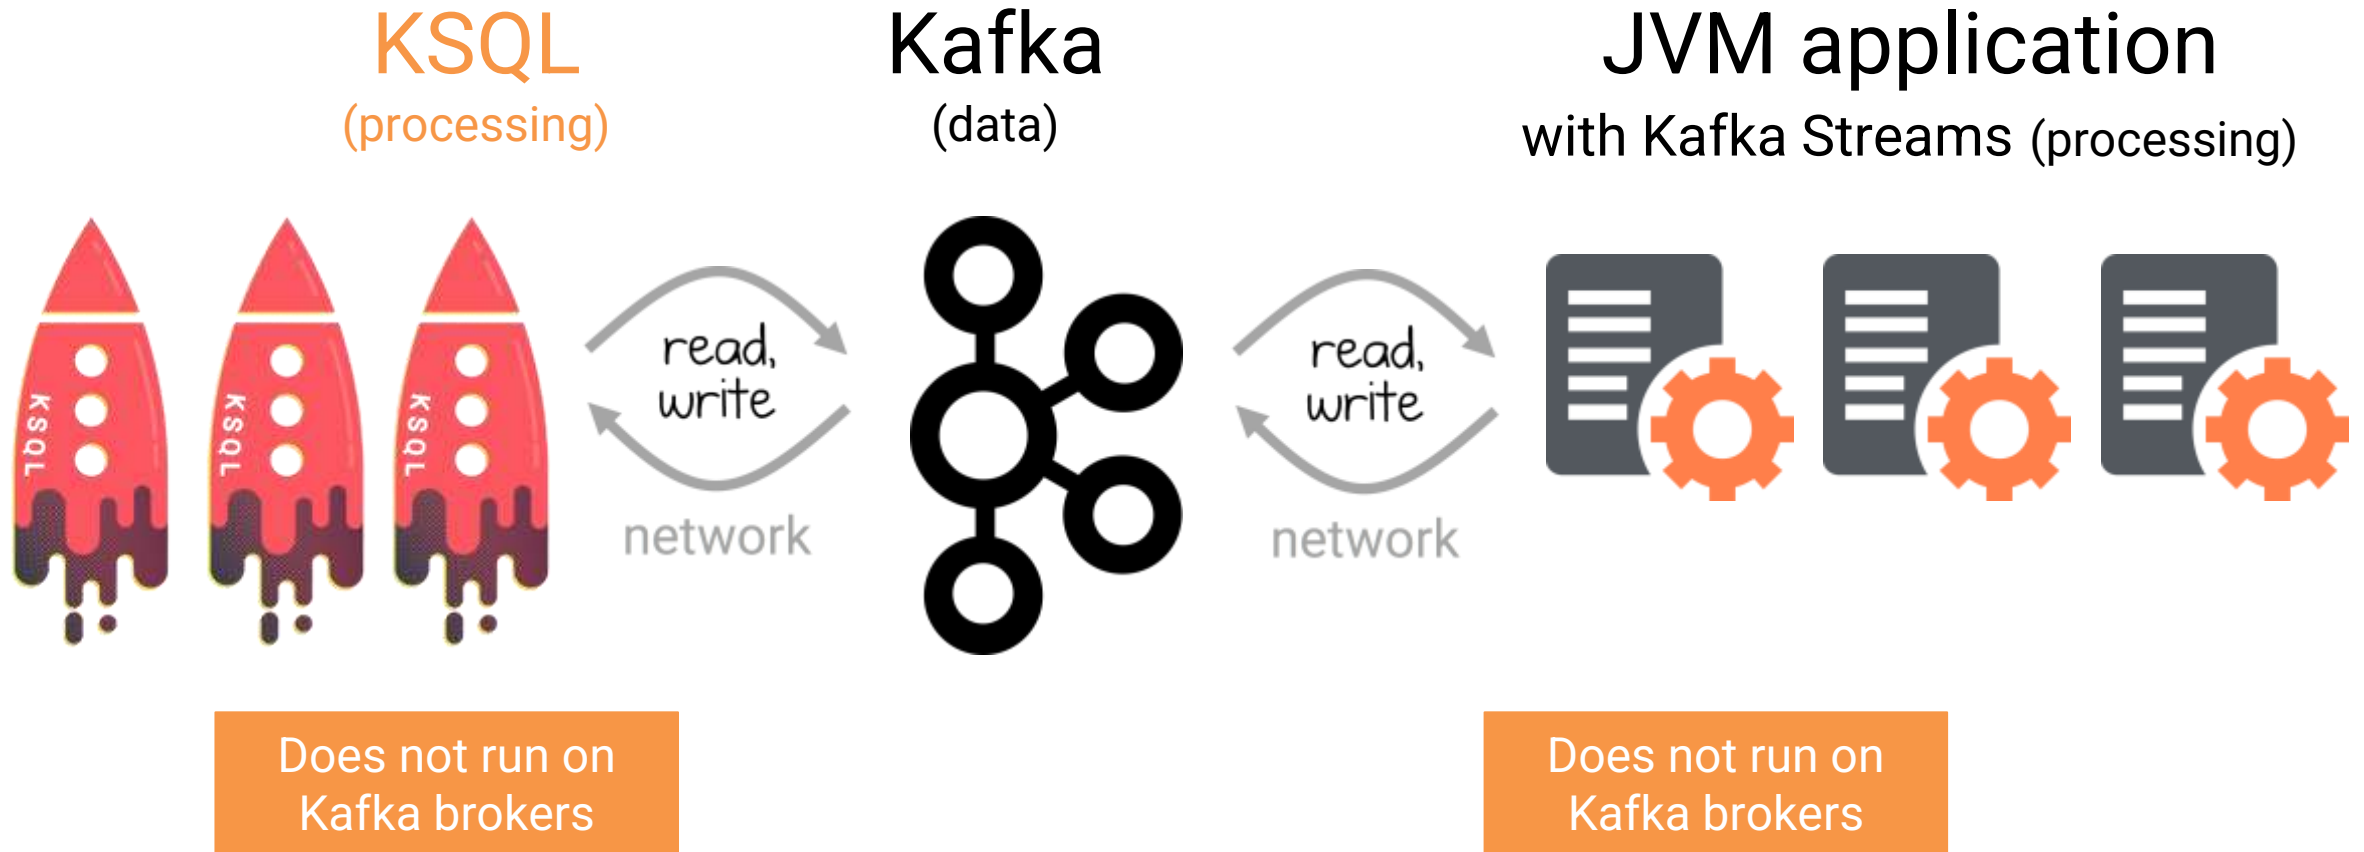
```
CREATE STREAM fraudulent_payments AS  
SELECT * FROM payments  
WHERE fraudProbability > 0.8;
```

You write *only* SQL. No Java, Python, or other boilerplate to wrap around it!

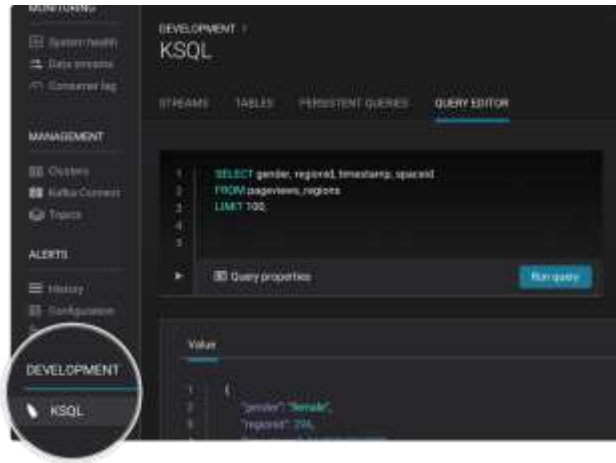
But you can create KSQL User Defined Functions in Java, if you want to.

```
object FraudFilteringApplication extends App {  
  
  val config = new java.util.Properties  
  config.put(StreamsConfig.APPLICATION_ID_CONFIG, "fraud-filtering-app")  
  config.put(StreamsConfig.BOOTSTRAP_SERVERS_CONFIG, "kafka-broker1:9092,kafka-broker2:9092")  
  
  val builder: StreamsBuilder = new StreamsBuilder()  
  val fraudulentPayments: KStream[String, Payment] = builder  
    .stream[String, Payment]("payments-kafka-topic")  
    .filter((_, payment) => payment.fraudProbability > 0.8)  
  
  val streams: KafkaStreams = new KafkaStreams(builder.build(), config)  
  streams.start()  
}
```

Interaction with Kafka



KSQL can be used interactively + programmatically



1 UI

ksql>

2 CLI

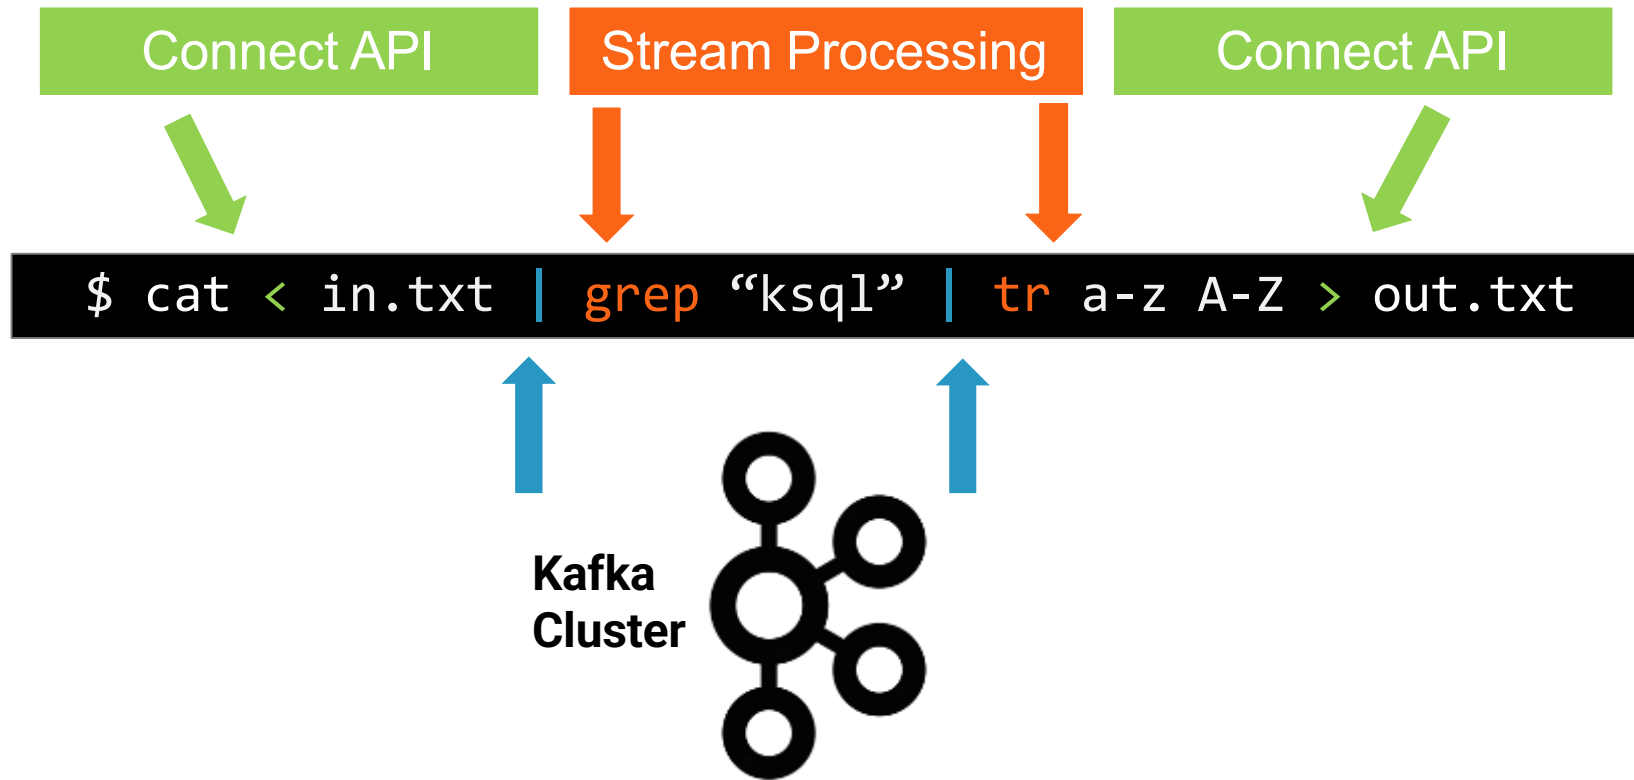
POST /query

3 REST

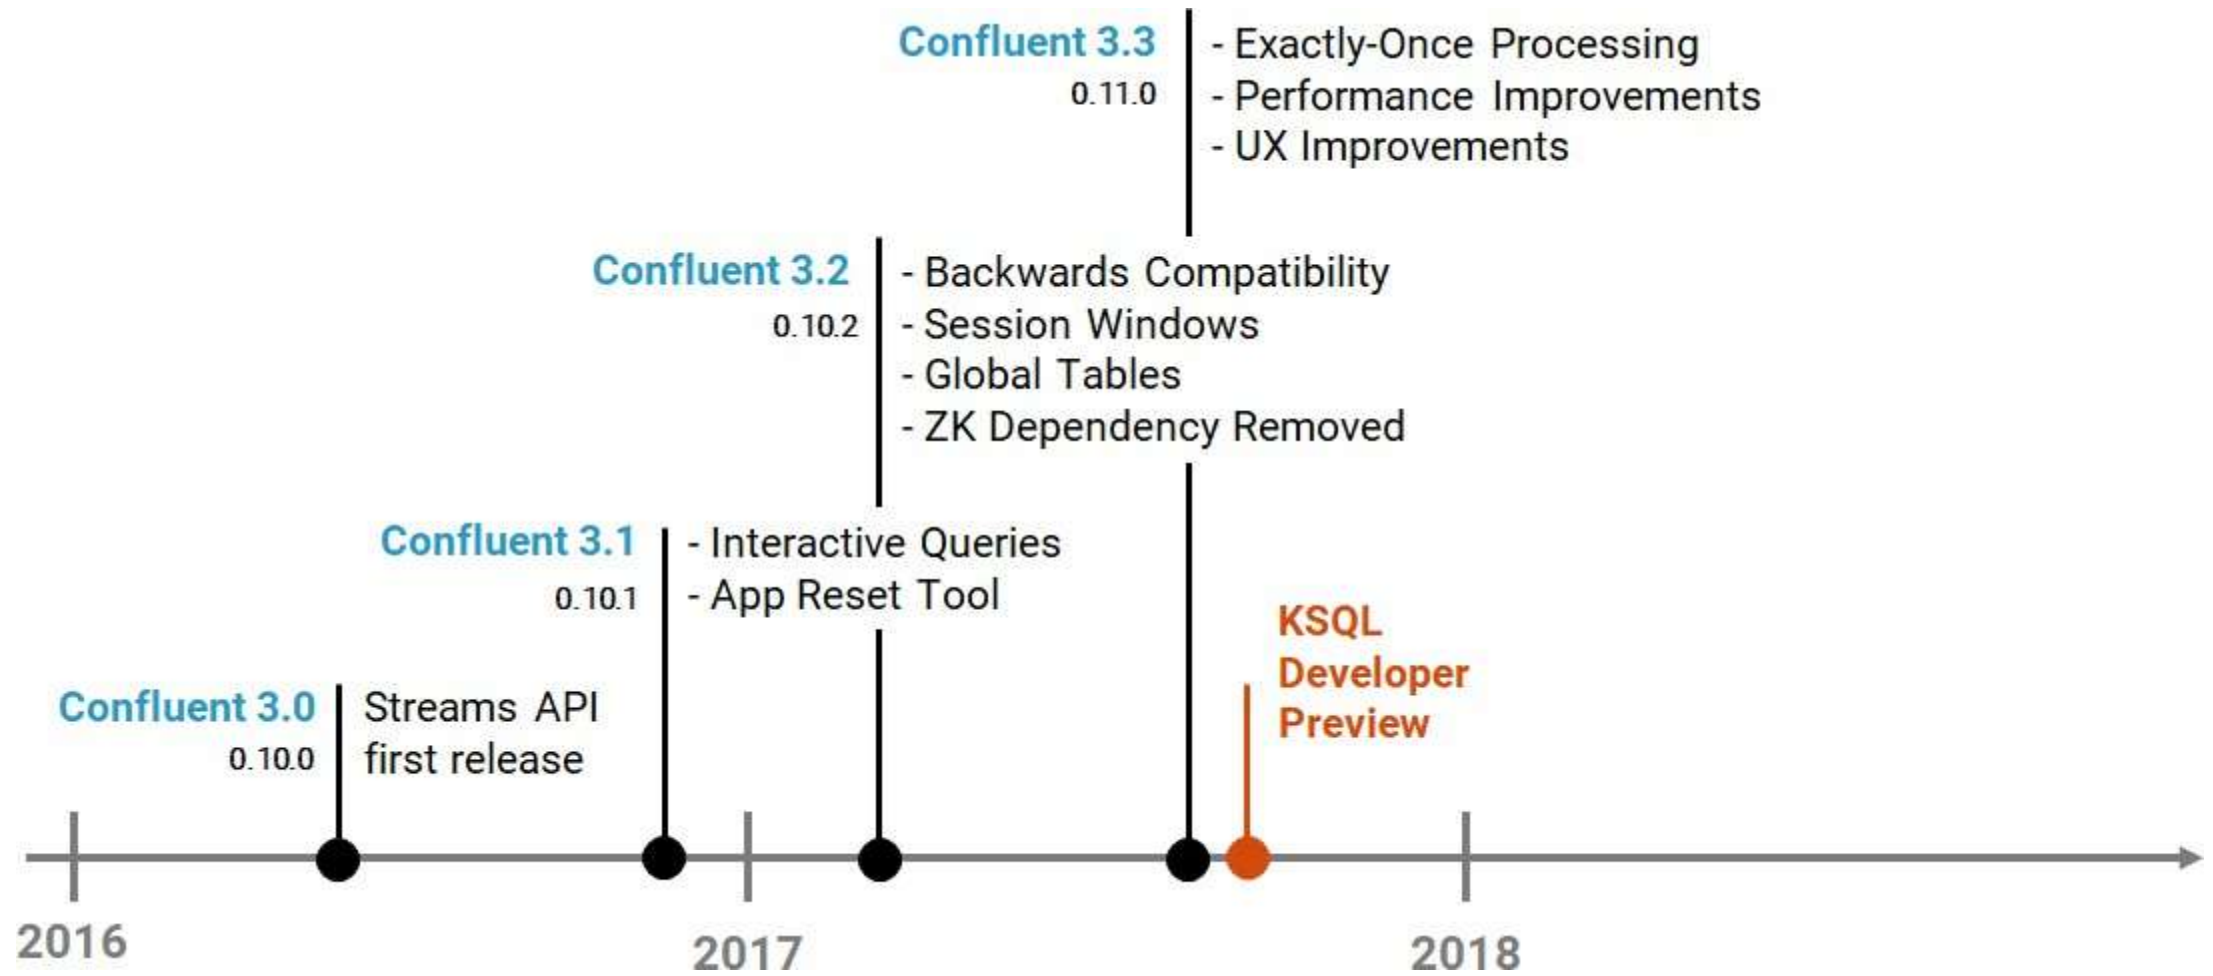


4 Headless

Stream Processing by Analogy



Kafka Stream Processing Evolution



Flexibility

The diagram features a large orange double-headed arrow spanning the width of the image. Below the arrow, three blue boxes are arranged horizontally, representing different Kafka ecosystem components. Each box is divided into two sections: the top section contains the component's name, and the bottom section lists its primary APIs. From left to right, the components are Consumer/Producer, Kafka Streams, and KSQL. The background of the entire image is a close-up, grayscale photograph of many macarons, which adds a visual texture to the technical content.

Simplicity

Consumer,
Producer

subscribe(),
poll(), send(),
flush()

Kafka
Streams

mapValues(),
filter(),
punctuate()

KSQL

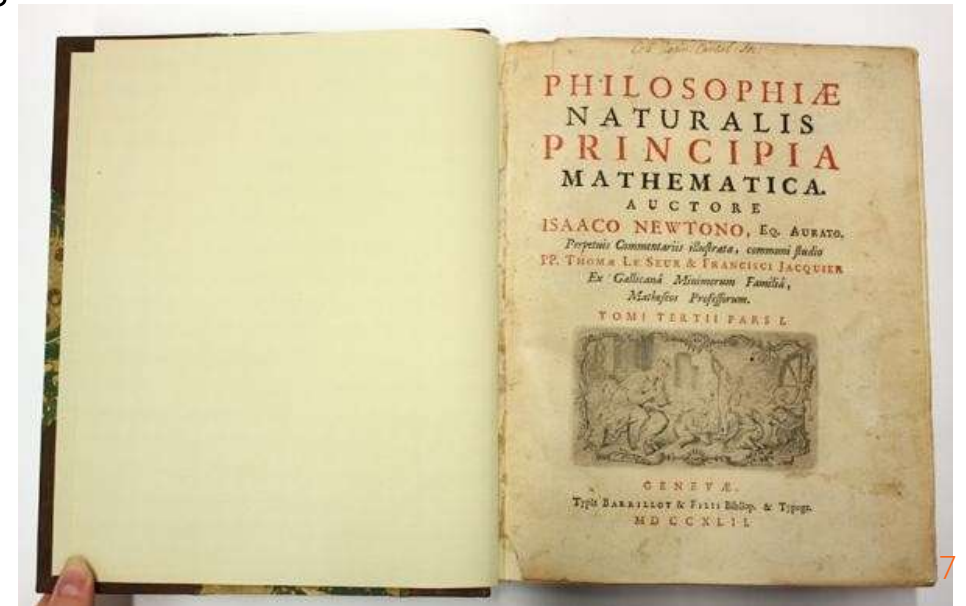
Select...from...
join...where...
group by..

On the Shoulders of (Streaming) Giants

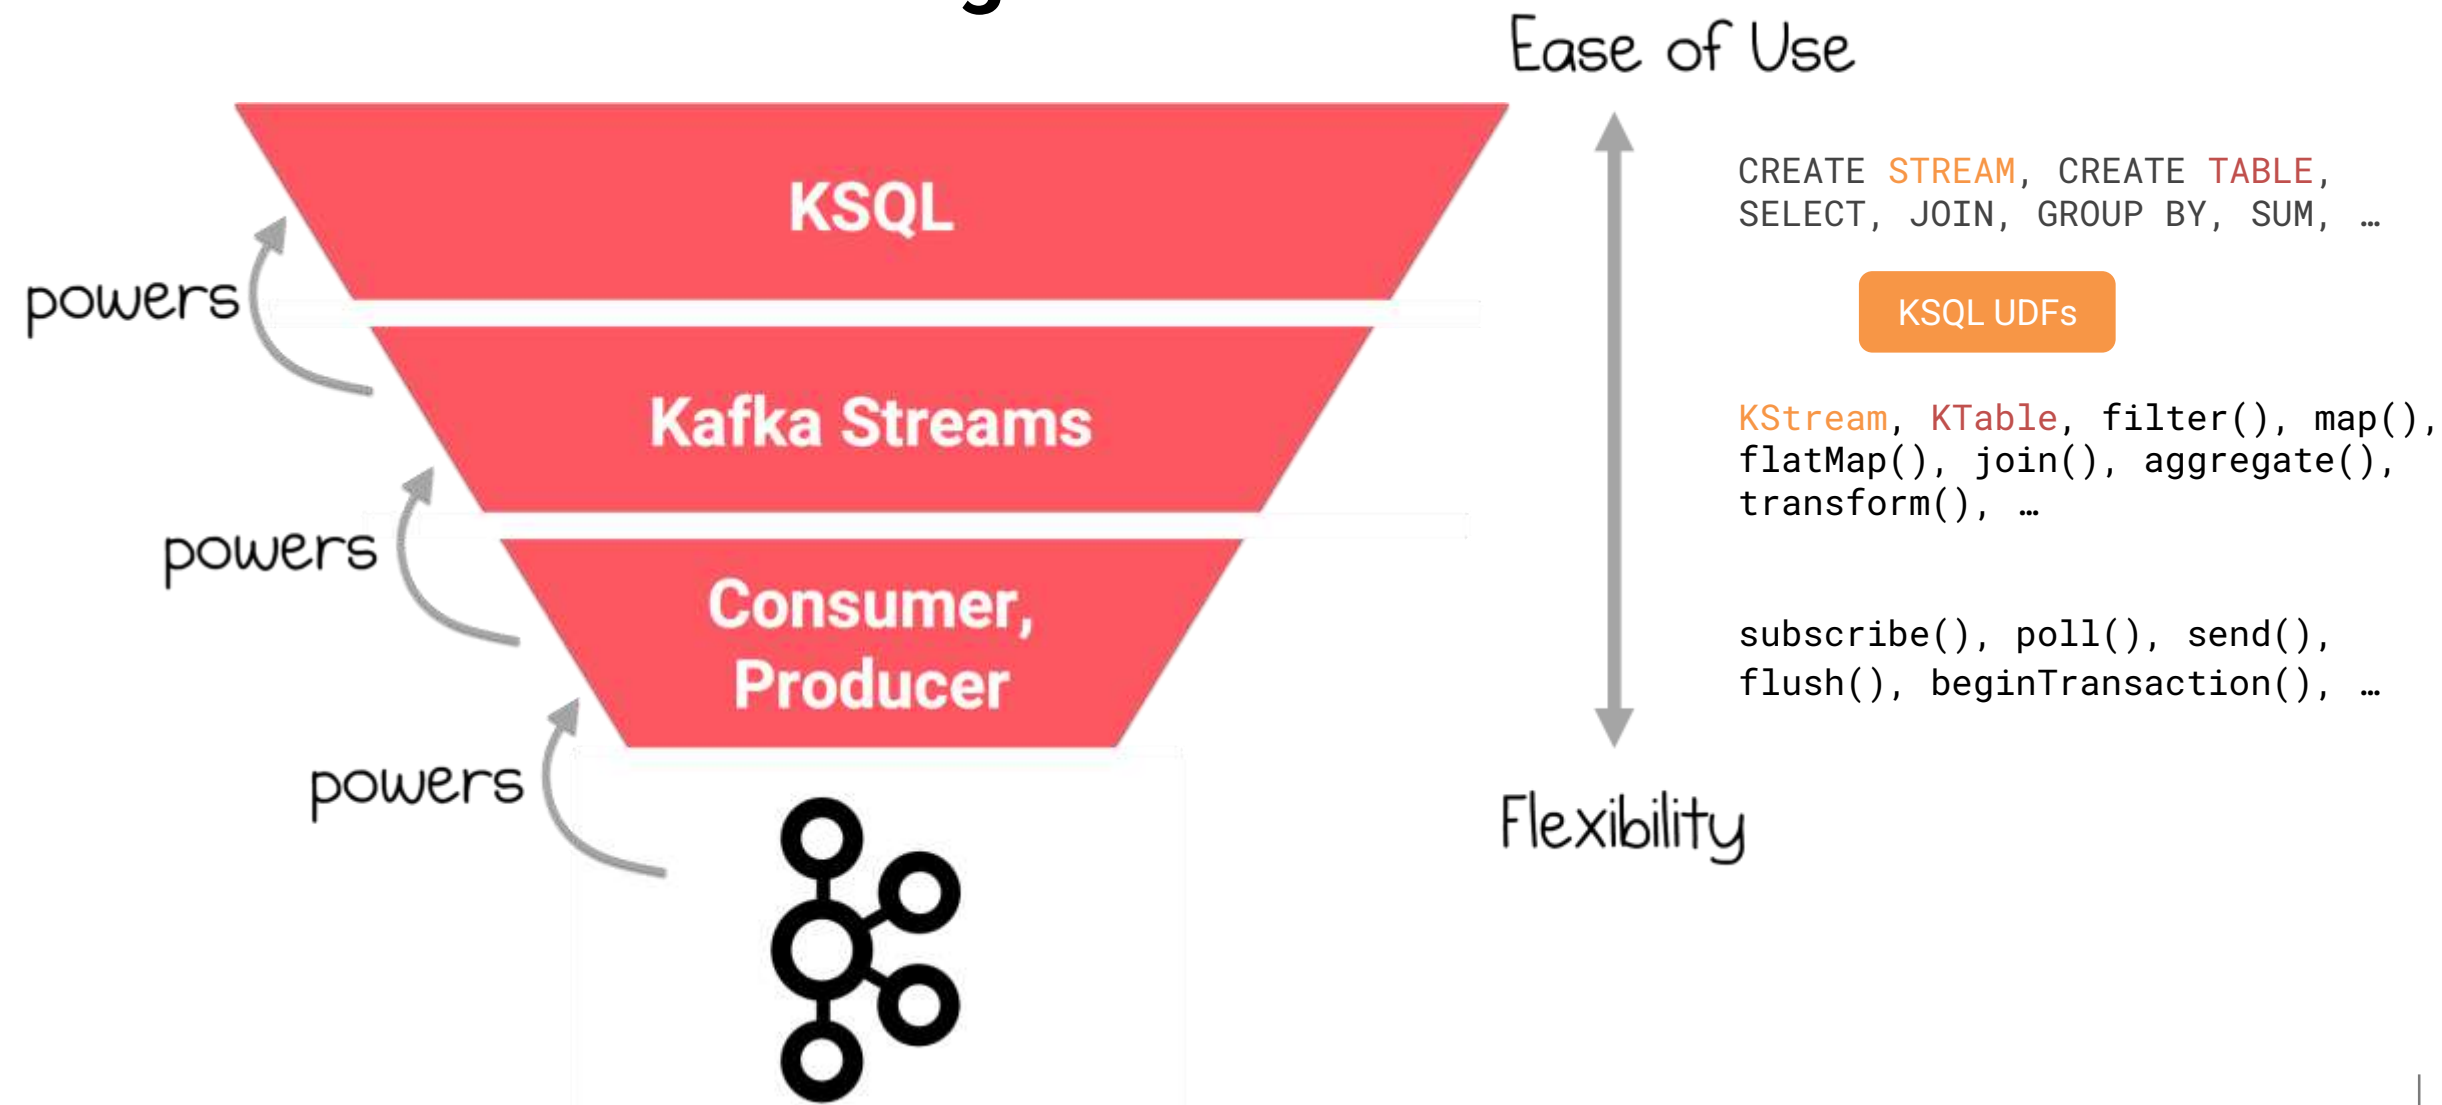


- Native, 100%-compatible Kafka integration
- Secure stream processing using Kafka's security features
- Elastic and highly scalable
- Fault-tolerant
- Stateful and stateless computations
- Interactive queries

- Time model
- Supports late-arriving and out-of-order data
- Windowing
- Millisecond processing latency, no micro-batching
- At-least-once and exactly-once processing guarantees



Shoulders of Streaming Giants



Example Use Cases

(focus on KSQL)

What is it for ?

- Streaming ETL
 - Kafka is popular for data pipelines.
 - KSQL enables easy transformations of data within the pipe

```
CREATE STREAM vip_actions AS  
SELECT userid, page, action FROM clickstream c  
LEFT JOIN users u ON c.userid = u.user_id  
WHERE u.level = 'Platinum';
```

What is it for ?

- Anomaly Detection
 - Identifying patterns or anomalies in real-time data, surfaced in milliseconds

```
CREATE TABLE possible_fraud AS  
SELECT card_number, count(*)  
FROM authorization_attempts  
WINDOW TUMBLING (SIZE 5 SECONDS)  
GROUP BY card_number  
HAVING count(*) > 3;
```


What is it for ?

- Real Time Monitoring
 - Log data monitoring, tracking and alerting
 - Sensor / IoT data

```
CREATE TABLE error_counts AS
SELECT error_code, count(*)
FROM monitoring_stream
WINDOW TUMBLING (SIZE 1 MINUTE)
WHERE type = 'ERROR'
GROUP BY error_code;
```


What is it for ?

- Simple Derivations of Existing Topics
 - One-liner to re-partition and/or re-key a topic for new uses

```
CREATE STREAM views_by_userid  
WITH (PARTITIONS=6,  
VALUE_FORMAT='JSON',  
TIMESTAMP='view_time') AS  
SELECT *  
FROM clickstream  
PARTITION BY user_id;
```

KSQL for Data Exploration

An easy way to inspect your data in Kafka

```
SHOW TOPICS;
```

```
PRINT 'my-topic' FROM BEGINNING;
```

```
SELECT page, user_id, status, bytes  
FROM clickstream  
WHERE user_agent LIKE 'Mozilla/5.0%';
```

KSQL for Data Transformation

Quickly make derivations of existing data in Kafka

```
CREATE STREAM clicks_by_user_id
  WITH (PARTITIONS=6,
        TIMESTAMP='view_time',
        VALUE_FORMAT='JSON') AS
  SELECT * FROM clickstream
  PARTITION BY user_id;
```

1 Change number of partitions

2 Convert data to JSON

3 Repartition the data

KSQL for Real-Time, Streaming ETL


Filter, cleanse, process data while it is in motion

```
CREATE STREAM clicks_from_vip_users AS
  SELECT user_id, u.country, page, action
  FROM clickstream c
  LEFT JOIN users u ON c.user_id = u.user_id
  WHERE u.level = 'Platinum';
```

1

Pick only VIP users

Example: CDC from DB via Kafka to Elastic



KSQL for Real-time Data Enrichment

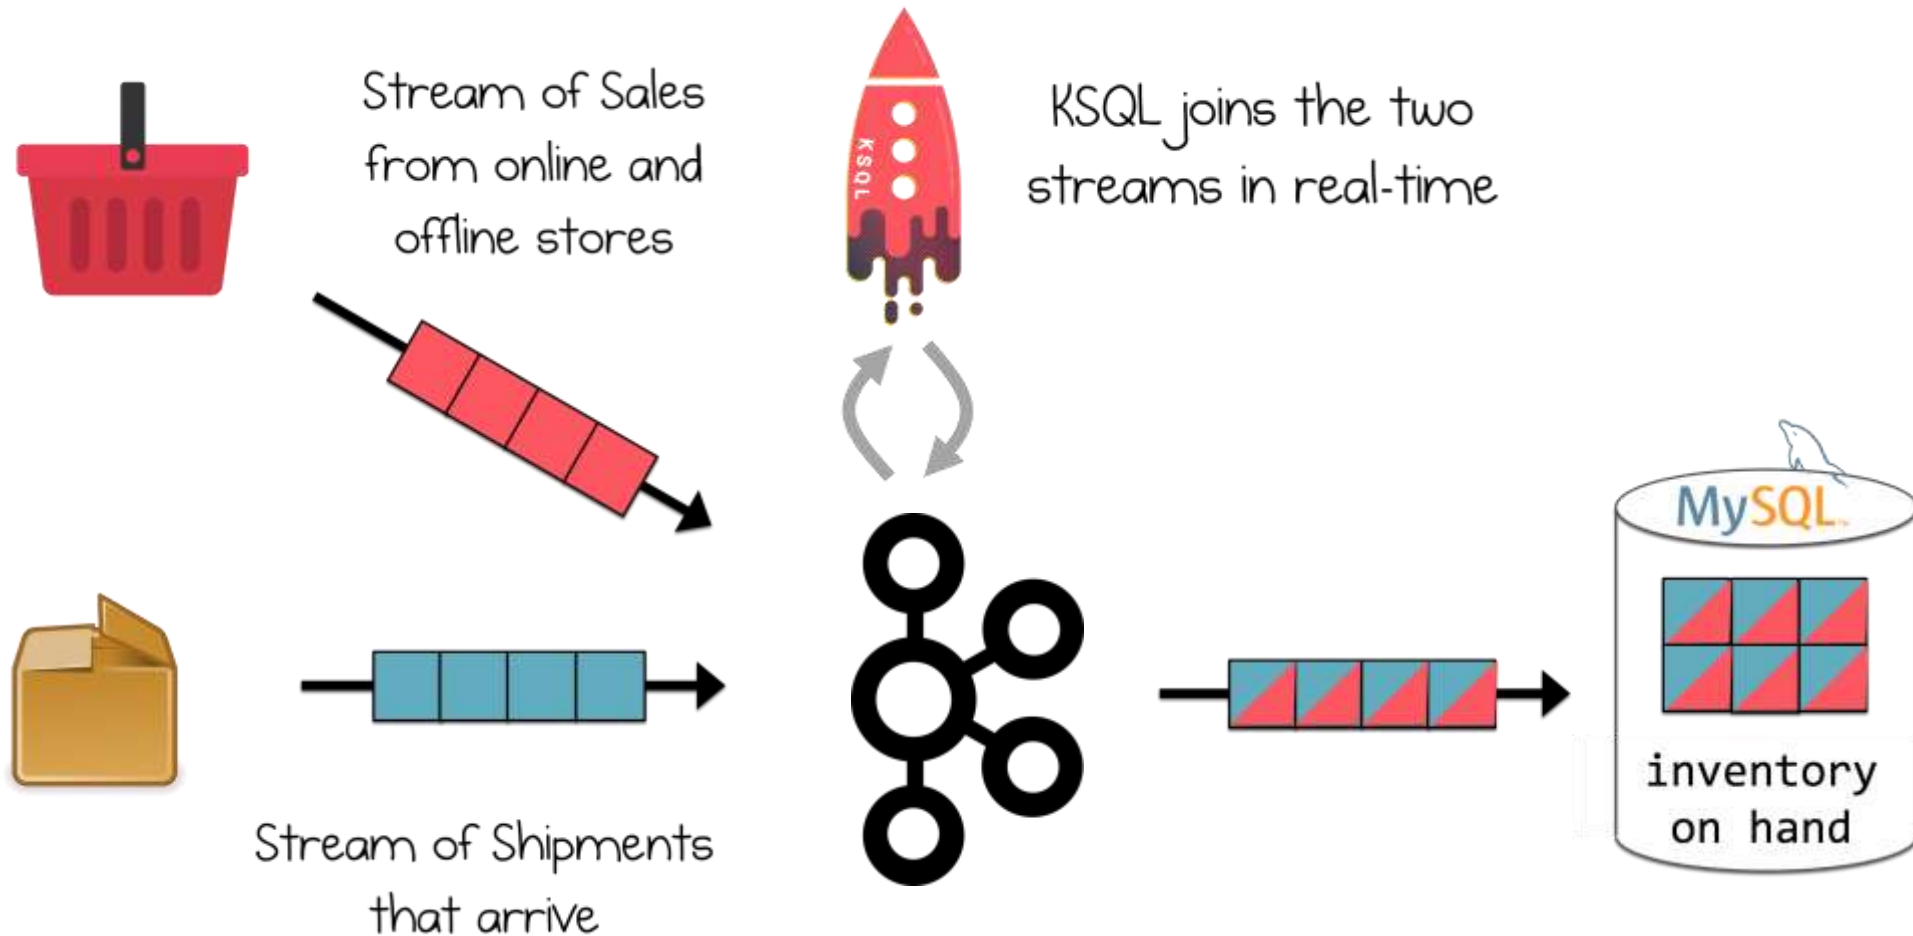
Join data from a variety of sources to see the full picture

```
CREATE STREAM enriched_payments AS
  SELECT payment_id, c.country, total
  FROM payments_stream p
  LEFT JOIN customers_table c
    ON p.user_id = c.user_id;
```

1

Stream-Table Join

Example: Retail



KSQL for Real-Time Monitoring


Derive insights from events (IoT, sensors, etc.) and turn them into actions

```
CREATE TABLE failing_vehicles AS
  SELECT vehicle, COUNT(*)
  FROM vehicle_monitoring_stream
  WINDOW TUMBLING (SIZE 1 MINUTE)
  WHERE event_type = 'ERROR'
  GROUP BY vehicle
  HAVING COUNT(*) >= 5;
```

1

Now we know to alert, and whom

Example: IoT, Automotive, Connected Cars



KSQL for Anomaly Detection

Aggregate data to identify patterns and anomalies in real-time

```
CREATE TABLE possible_fraud AS  
  SELECT card_number, COUNT(*)  
  FROM authorization_attempts  
  WINDOW TUMBLING (SIZE 30 SECONDS)  
  GROUP BY card_number  
  HAVING COUNT(*) > 3;
```

1

Aggregate data

2


... per 30-sec windows

Where is KSQL not such a great fit ?

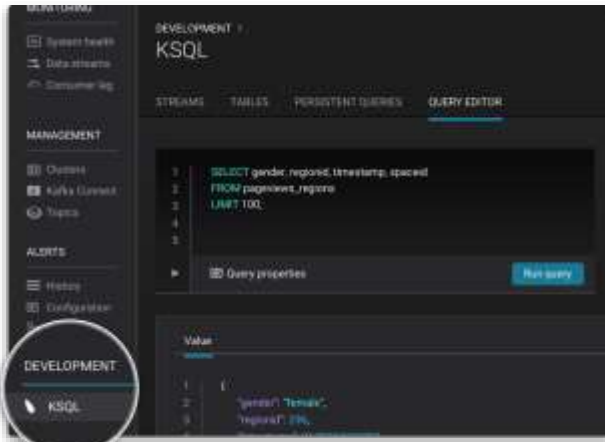
- Ad-hoc query
 - Limited span of time usually retained in Kafka
 - No indexes for random point lookups
- BI reports (Tableau etc.)
 - No indexes
 - No JDBC (most BI tools are not good with continuous results!)

Workflow Comparison

Typical developer interaction



KSQL



write KSQL
queries



view results
in real-time

write code in
Java or Scala

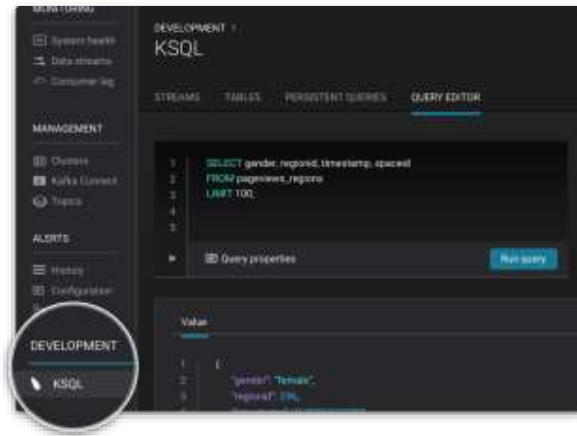


recompile,
then run/test
your app



KSQL: typical workflow from development to production

Interactive KSQL
for development



develop your application
and its queries

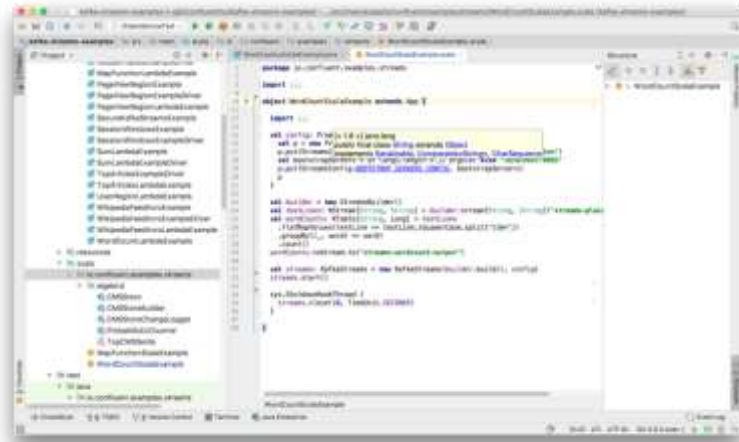
Headless KSQL
in production



deploy & run application

Kafka Streams: typical workflow from development to production

Local development and testing
with Java/Scala IDE



develop your application



build & package the
Java/Scala application

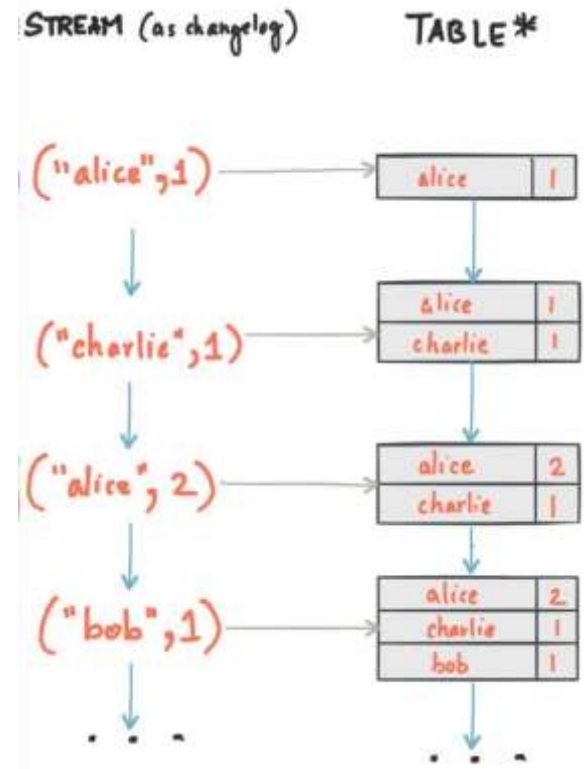
Production




deploy & run application

Streams & Tables

- **STREAM and TABLE as first-class citizens**
- Interpretations of topic content
- **STREAM** - data in motion
- **TABLE** - collected state of a stream
 - One record per key (per window)
 - Current values (compacted topic) ← **Not yet in KSQL**
 - Changelog
- **STREAM – TABLE Joins**




Running Example



Schema and Format

- A Kafka broker knows how to move <sets of bytes>
 - Technically a message is (ts, byte[], byte[])
- SQL-like queries require a richer structure
- Start with message (value) format
 - JSON
 - DELIMITED (comma-separated in this preview)
 - AVRO – requires that you supply a .avsc schema-file
- Pseudo-columns are automatically generated
 - ROWKEY, ROWTIME



Schema & Datatypes

- varchar / string
- boolean / bool
- integer / int
- bigint / long
- double
- array(*of_type*)
 - *of-type* must be primitive (no nested Array or Map yet)
- map(*key_type*, *value_type*)
 - *key-type* must be string, *value-type* must be primitive

Define a Stream

```
CREATE STREAM ratings (  
    rating_id long,  
    user_id int,  
    stars int,  
    route_id int,  
    rating_time long,  
    channel varchar,  
    message varchar)  
WITH (  
    value_format='JSON',  
    kafka_topic='ratings');
```

SELECTing from the Stream



Let's test our new stream definition by finding all the low-scoring ratings from our iPhone app

```
SELECT *  
FROM ratings  
WHERE stars <= 2  
AND lcase(channel) LIKE '%ios%'  
AND user_id > 0  
LIMIT 10;
```

SELECTing from the Stream



And set this to run as a continuous transformation, with results being saved into a new topic

```
CREATE STREAM poor_ratings AS  
SELECT *  
FROM ratings  
WHERE stars <= 2  
AND lcase(channel) LIKE '%ios%';
```


Bring in reference tables



Define reference tables

```
CREATE TABLE users (  
    uid int,  
    name varchar,  
    elite varchar)  
WITH (  
    key='uid',  
    value_format='JSON',  
    kafka_topic='mysql-users');
```


Joins for Enrichment

Enrich the 'poor_ratings' stream with data about each user, and derive a stream of low quality ratings posted only by our Platinum Elite users

```
CREATE STREAM vip_poor_ratings AS
SELECT uid, name, elite,
       stars, route_id, rating_time, message
FROM poor_ratings r
LEFT JOIN users u ON r.user_id = u.uid
WHERE u.elite = 'P';
```



Aggregates and Windowing



- COUNT, SUM, MIN, MAX
- Windowing - Not strictly ANSI SQL 😊
- Three window types supported:
 - TUMBLING
 - HOPPING (aka 'sliding')
 - SESSION

```
SELECT uid, name, count(*) as rating_count  
FROM vip_poor_ratings  
WINDOW TUMBLING(size 2 minutes)  
GROUP BY uid, name;
```

Continuous Aggregates

Save the results of our aggregation to a TABLE

```
CREATE TABLE sad_vips AS
SELECT uid, name, count(*) as rating_count
FROM vip_poor_ratings
WINDOW TUMBLING(size 1 minute)
GROUP BY uid, name
HAVING count(*) > 2;
```

Session Variables

- Just as in MySQL, ORCL etc. there are settings to control how your CLI behaves
- Set any property the Kafka Streams consumers/producers will understand
- Defaults can be set in the *ksql.properties* file
- To see a list of currently set or default variable values:
 - `ksql> show properties;`
- Useful examples:
 - `num.stream.threads=4`
 - `commit.interval.ms=1000`
 - `cache.max.bytes.buffering=2000000`
- TIP! - Your new best friend for testing and development is:
 - `ksql> set 'auto.offset.reset' = 'earliest';`

KSQL Components

- CLI
 - Designed to be familiar to users of MySQL, Postgres, etc
- Engine
 - Actually runs the Kafka Streams topologies
- REST Server
 - HTTP interface allows an Engine to receive instructions from the CLI

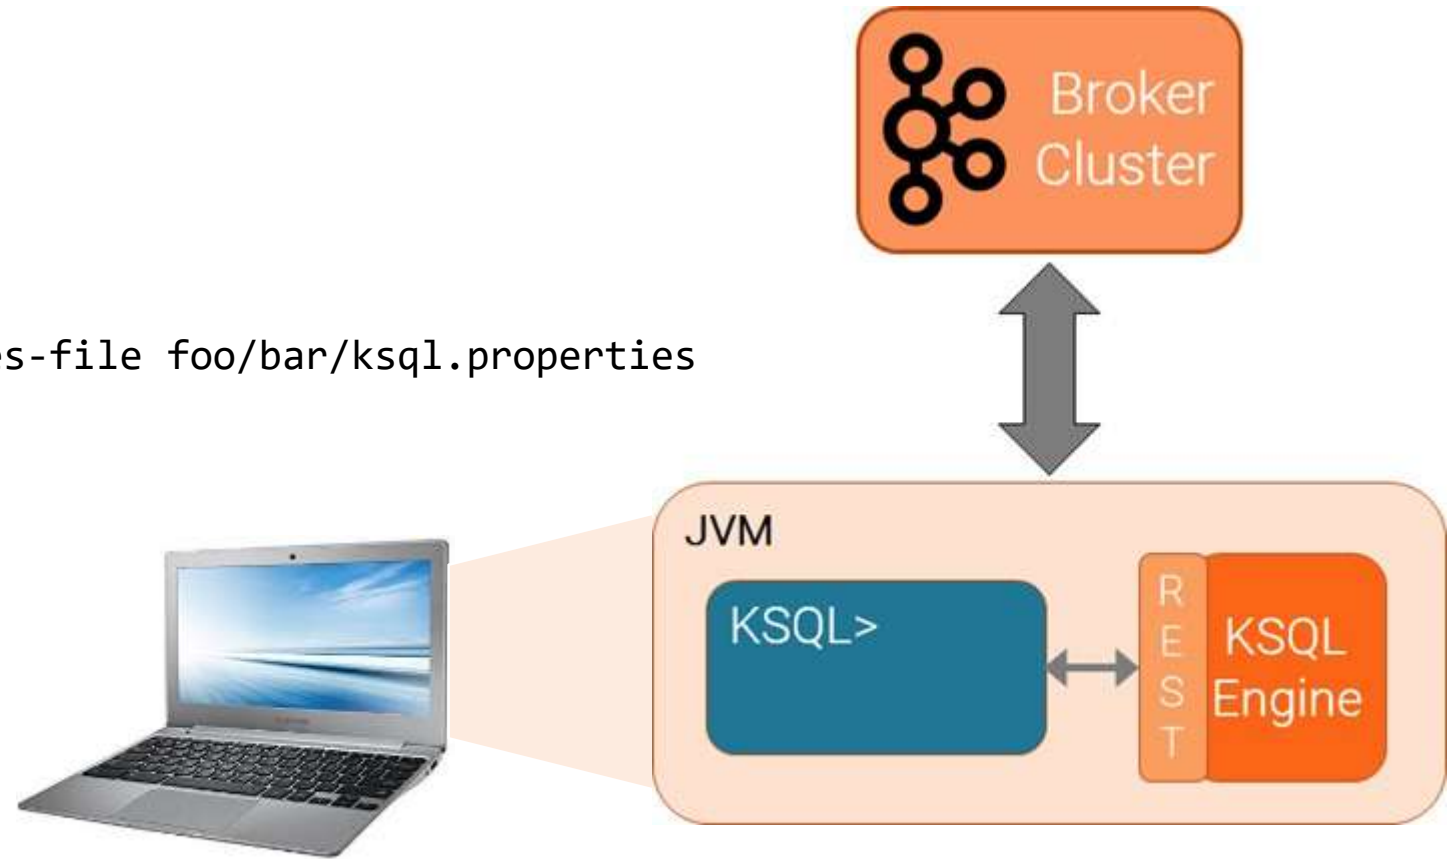


How to run KSQL - #1 Stand-alone aka 'local mode'

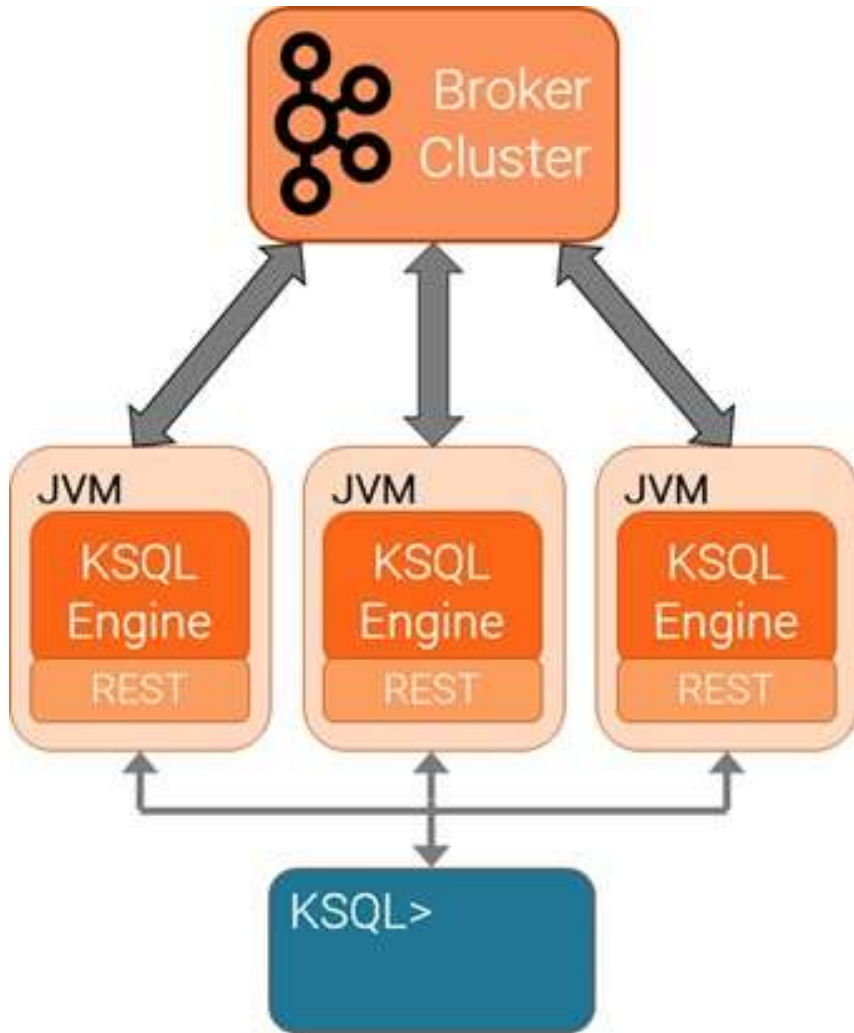
- Starts a CLI, an Engine, and a REST server all in the same JVM
- Ideal for laptop development
 - Start with default settings:

```
> bin/ksql-cli local
```
 - Or with customized settings:

```
> bin/ksql-cli local --properties-file foo/bar/ksql.properties
```




How to run KSQL - #2 Client-Server



- Start any number of Server nodes
 - `> bin/ksql-server-start`
- Start any number of CLIs and specify 'remote' server address
 - `> bin/ksql-cli remote http://myserver:8090`
- All running Engines share the processing load
 - Technically, instances of the same Kafka Streams Applications
 - Scale up/down without restart

How to run KSQL - #3 as an Application

- Ideal for streaming application deployment
 - Version control your queries and transformations as code
 - Deploy like any other java application
 - Avoid interactive changes to running apps from 'rogue' CLI users
- Start any number of Engine instances
 - Pass a file of KSQL statements to execute
 - > `bin/ksql-node foo/bar.sql`
- All running Engines share the processing load
 - Technically, instances of the same Kafka Streams Applications
 - Scale up/down without restart



Resources & Next Steps

Time to get involved !

- Try the Quickstart on Github
- Check out the code
- Play with the examples



The point of 'developer preview' is that we can change things for the better, together



<https://github.com/confluentinc/ksql>



<http://confluent.io/ksql>



<https://slackpass.io/confluentcommunity> #ksql